

Approximate Distributed Monitoring under Partial Synchrony: Balancing Speed with Accuracy

Author(s)

Affiliation(s)

Abstract. In distributed systems with processes that do not share a global clock, *partial synchrony* is achieved by clock synchronization that guarantees bounded clock skew among all applications. Existing solutions for distributed runtime verification under partial synchrony against temporal logic specifications are exact but suffer from significant computational overhead. In this paper, we propose an *approximate* distributed monitoring algorithm for Signal Temporal Logic (STL) that mitigates this issue by abstracting away potential interleaving behaviors. This conservative abstraction enables a significant speedup of the distributed monitors, albeit with a trade-off in accuracy. We address this trade-off with a methodology that combines our approximate monitor with its exact counterpart, resulting in enhanced monitoring efficiency without sacrificing precision. We validate our approach with multiple experiments, showcasing its effectiveness and efficacy on both a real-world application and synthetic examples.

You can leave notes using the command `\firstName{note}`.
Page limit: 16 + 4 (appendix)

1 Introduction

Distributed systems are networks of independent agents that work together to achieve a common objective. Distributed systems are everywhere around us and come in many different forms. For example, cloud computing uses distribution of resources and services over the internet to offer to their users a scalable infrastructure with transparent on-demand access to computing power and storage. Swarms of drones represent another family of distributed systems where individual drones collaborate to accomplish tasks like surveillance, search and rescue, or package delivery. While each drone operates independently, it also communicates and coordinates with others to successfully achieve their common objectives. The individual agents in a distributed system typically do not share a global clock. To coordinate actions across multiple agents, clock synchronization is often needed. While perfect clock synchronization is impractical due to network latency and node failures, algorithms such as the Network Time Protocol (NTP) allow agents

to maintain a *bounded skew* between the synchronized clocks. In that case, we say that a distributed system has *partial synchrony*.

Formal verification of distributed system is a notoriously hard problem, due to the combinatorial explosion of all possible interleavings in the behaviors collected from individual agents. *Runtime verification (RV)* provides a more pragmatic approach, in which a monitor observes a behavior of a distributed system and checks its correctness against a formal specification. The problem of distributed RV under partial synchrony assumption has been studied for Linear Temporal Logic (LTL) and Signal Temporal Logic (STL) specification languages. The proposed solutions use Satisfiability-Modulo-Theory (SMT) solving to provide sound and complete distributed monitoring procedures. Although distributed RV monitors consume only a single distributed behavior at a time, this behavior can nevertheless have an excessive number of possible interleavings. Hence, the exact distributed monitors from the literature can still suffer from significant computational overhead.

To mitigate this issue, we present an approach for *approximate* RV of STL specifications under partial synchrony. In essence, we abstract away potential interleavings in distributed behaviors in a conservative manner, resulting in an effective over-approximation of global behaviors. This abstraction simplifies the representation of distributed behaviors and the monitoring operations required to evaluate temporal specifications. There is an inevitable trade-off in approximate RV – gains in the monitoring speed-up may result in reduced accuracy. For some applications, reduced accuracy may not be acceptable. Therefore, we propose a methodology that combines our approximate monitors with their exact counterparts, with the aim to benefit from the enhanced monitoring efficiency without sacrificing precision. We implemented our approach in a prototype tool and performed thorough evaluations on both synthetic and real-world case studies. We first demonstrated that our approximate monitors achieve speed-ups of several orders of magnitudes compared to the exact SMT-based distributed RV solution. We empirically characterized the classes of specifications and behaviors for which our approximate monitoring approach achieves good precision. We finally showed that by combining exact and approximate distributed RV, there is still a significant efficiency gain on average without the sacrifice of the precision, even in cases where approximate monitors have low accuracy.

2 Preliminaries

We denote by $\mathbb{B} = \{\top, \perp\}$ the set of Booleans, \mathbb{R} the set of reals, $\mathbb{R}_{\geq 0}$ the set of nonnegative reals, and $\mathbb{R}_{> 0}$ the set of positive reals. An interval $I \subseteq \mathbb{R}$ of reals with the end points $a < b$ has length $|b - a|$.

Let Σ be a finite *alphabet*. We denote by Σ^* the set of finite words over Σ and by ϵ the empty word. For $u \in \Sigma^*$, we respectively write $\text{prefix}(u)$ and $\text{suffix}(u)$ for the sets of prefixes and suffixes of u . We also let $\text{infix}(u) = \{v \in \Sigma^* \mid \exists x, y \in \Sigma^* : u = xvy\}$. For a nonempty word $u \in \Sigma^*$ and $1 \leq i \leq |u|$, we denote by $u[i]$ the i th letter of u , by $u[..i]$ the prefix of u of length i , and by

79 $u[i..]$ the suffix of u of length $|u| - i + 1$. Given $u \in \Sigma^*$ and $\ell \geq 1$, we denote by
 80 u^ℓ the word obtained by concatenating u by itself $\ell - 1$ times. Moreover, given
 81 $L \subseteq \Sigma^*$, we define $\text{first}(L) = \{u[0] \mid u \in L\}$. For sets $L_1, L_2 \in \Sigma^*$ of words, we
 82 let $L_1 \cdot L_2 = \{u \cdot v \mid u \in L_1, v \in L_2\}$. For tuples (u_1, \dots, u_m) and (v_1, \dots, v_m) of
 83 words, we let $(u_1, \dots, u_m) \cdot (v_1, \dots, v_m) = (u_1v_1, \dots, u_mv_m)$.

84 We define the function $\text{destutter} : \Sigma^* \rightarrow \Sigma^*$ inductively as follows. For all
 85 $\sigma \in \Sigma \cup \{\epsilon\}$, let $\text{destutter}(\sigma) = \sigma$. For all $u \in \Sigma^*$ such that $u = \sigma_1\sigma_2v$ for
 86 some $\sigma_1, \sigma_2 \in \Sigma$ and $v \in \Sigma^*$, let (i) $\text{destutter}(u) = \text{destutter}(\sigma_2v)$ if $\sigma_1 = \sigma_2$,
 87 and (ii) $\text{destutter}(u) = \sigma_1 \cdot \text{destutter}(\sigma_2v)$ otherwise. By extension, for a set
 88 $L \subseteq \Sigma^*$ of finite words, we write $\text{destutter}(L) = \{\text{destutter}(u) \mid u \in L\}$. Given
 89 a tuple $(u_1, \dots, u_m) = (\sigma_{1,1}\sigma_{1,2}v_1, \dots, \sigma_{m,1}\sigma_{m,2}v_m)$ of finite words of the same
 90 length, we define $\text{destutter}(u_1, \dots, u_m)$ as expected: (i) $\text{destutter}(u_1, \dots, u_m) =$
 91 $\text{destutter}(\sigma_{1,2}v_1, \dots, \sigma_{m,2}v_m)$ if $\sigma_{i,1} = \sigma_{i,2}$ for all $1 \leq i \leq m$, and (ii) $\text{destutter}(u_1, \dots, u_m) =$
 92 $(\sigma_{1,1}, \dots, \sigma_{m,1}) \cdot \text{destutter}(\sigma_{1,2}v_1, \dots, \sigma_{m,2}v_m)$ otherwise.

93 Moreover, given an integer $k \geq 0$, we define $\text{stutter}_k : \Sigma^* \rightarrow \Sigma^*$ such
 94 that $\text{stutter}_k(u) = \{v \in \Sigma^* \mid |v| = k \wedge \text{destutter}(v) = \text{destutter}(u)\}$ if $k \geq$
 95 $|\text{destutter}(u)|$, and $\text{stutter}_k(u) = \emptyset$ otherwise.

96 **Signal Temporal Logic (STL)** Let $A, B \subset \mathbb{R}$. A function $f : A \rightarrow B$ is
 97 *right-continuous* iff $\lim_{a \rightarrow c^+} f(a) = f(c)$ for all $c \in A$, and *non-Zeno* iff for
 98 every bounded interval $I \subseteq A$ there are finitely many $a \in I$ such that f is not
 99 continuous at a . A *signal* is a right-continuous, non-Zeno, piecewise-constant
 100 function $x : [0, d) \rightarrow \mathbb{R}$ where $d \in \mathbb{R}_{>0}$ is the duration of x and $[0, d)$ is its
 101 temporal domain. Let $x : [0, d) \rightarrow \mathbb{R}$ be a signal. An *event* of x is a pair $(t, x(t))$
 102 where $t \in [0, d)$. An *edge* of x is an event $(t, x(t))$ such that $\lim_{s \rightarrow t^-} x(s) \neq$
 103 $\lim_{s \rightarrow t^+} x(s)$. In particular, an edge is *rising* if $\lim_{s \rightarrow t^-} x(s) < \lim_{s \rightarrow t^+} x(s)$, and
 104 it is *falling* otherwise. A signal $x : [0, d) \rightarrow \mathbb{R}$ can be represented finitely by its
 105 initial value and edges: if x has m edges, then $x = (t_0, v_0)(t_1, v_1) \dots (t_m, v_m)$
 106 such that $t_0 = 0$, $t_{i-1} < t_i$, and (t_i, v_i) is an edge of x for all $1 \leq i \leq m$.

107 Let AP be a set of atomic propositions. The syntax is given by the following
 108 grammar where $p \in \text{AP}$ and $I \subseteq \mathbb{R}_{\geq 0}$ is an interval.

$$\varphi := p \mid \neg\varphi \mid \varphi \wedge \varphi \mid \varphi \mathcal{U}_I \varphi$$

109 A *trace* $w = (x_1, \dots, x_n)$ is a finite vector of signals. We express atomic
 110 propositions as functions of trace values at a time point t , i.e., a proposition
 111 $p \in \text{AP}$ over a trace $w = (x_1, \dots, x_n)$ is defined as $f_p(x_1(t), \dots, x_n(t)) > 0$
 112 where $f_p : \mathbb{R}^n \rightarrow \mathbb{R}$ is a function. Given intervals $I, J \subseteq \mathbb{R}_{\geq 0}$, we define $I \oplus J =$
 113 $\{i + j \mid i \in I \wedge j \in J\}$, and we simply write t for the singleton set $\{t\}$.

114 Below we recall the finite-trace qualitative semantics of STL defined over \mathbb{B}
 115 $[1]$. Let $d \in \mathbb{R}_{>0}$ and $w = (x_1, \dots, x_n)$ with $x_i : [0, d) \rightarrow \mathbb{R}$ for all $1 \leq i \leq n$. Let
 116 φ_1, φ_2 be STL formulas and let $t \in [0, d)$.

$$\begin{aligned}
(w, t) \models p &\iff f_p(x_1(t), \dots, x_n(t)) > 0 \\
(w, t) \models \neg \varphi_1 &\iff \overline{(w, t) \models \varphi_1} \\
(w, t) \models \varphi_1 \wedge \varphi_2 &\iff (w, t) \models \varphi_1 \wedge (w, t) \models \varphi_2 \\
(w, t) \models \varphi_1 \mathcal{U}_I \varphi_2 &\iff \exists t' \in (t \oplus I) \cap [0, d) : \\
&\quad (w, t') \models \varphi_2 \wedge \forall t'' \in (t, t') : (w, t'') \models \varphi_1
\end{aligned}$$

117 We simply write $w \models \varphi$ for $(w, 0) \models \varphi$. We additionally use the following
118 standard abbreviations: **false** = $p \wedge \neg p$, **true** = $\neg \mathbf{false}$, $\varphi_1 \vee \varphi_2 = \neg(\neg \varphi_1 \wedge$
119 $\neg \varphi_2)$, $\Diamond_I \varphi = \mathbf{true} \mathcal{U}_I \varphi$, and $\Box_I \varphi = \neg \Diamond_I \neg \varphi$. Moreover, the untimed temporal
120 operators are defined through their timed counterparts on the interval $(0, \infty)$,
121 e.g., $\varphi_1 \mathcal{U} \varphi_2 = \varphi_1 \mathcal{U}_{(0, \infty)} \varphi_2$.

122 **Distributed Semantics of STL** We consider an asynchronous and loosely-
123 coupled message-passing system of $n \geq 2$ reliable agents producing a set of
124 signals x_1, \dots, x_n , where for some $d \in \mathbb{R}_{>0}$ we have $x_i : [0, d) \rightarrow \mathbb{R}$ for all
125 $1 \leq i \leq n$. The agents do not share memory or a global clock. Only to formalize
126 statements, we speak of a *hypothetical* global clock and denote its value by T .
127 For local time values, we use the lowercase letters t and s .

128 For a signal x_i , we denote by V_i the set of its events, by E_i^\uparrow the set of its
129 rising edges, and by E_i^\downarrow that of falling edges. Moreover, we let $E_i = E_i^\uparrow \cup E_i^\downarrow$. We
130 represent the local clock of the i th agent as an increasing and divergent function
131 $c_i : \mathbb{R}_{\geq 0} \rightarrow \mathbb{R}_{\geq 0}$ that maps a global time T to a local time $c_i(T)$.

132 We assume that the system is *partially synchronous*: the agents use a clock
133 synchronization algorithm that guarantees a bounded clock skew with respect
134 to the global clock, i.e., $|c_i(T) - c_j(T)| < \varepsilon$ for all $1 \leq i, j \leq N$ and $T \in \mathbb{R}_{\geq 0}$,
135 where $\varepsilon \in \mathbb{R}_{>0}$ is the maximum clock skew.

136 **Definition 1.** A distributed signal is a pair (S, \rightsquigarrow) , where $S = (x_1, \dots, x_n)$ is a
137 vector of signals and \rightsquigarrow is the happened-before relation between events in signals
138 extended with the partial synchrony assumption as follows.

- 139 – For every agent, the events of its signals are totally ordered, i.e., for all $1 \leq$
140 $i \leq n$ and all $(t, x_i(t)), (t', x_i(t')) \in V_i$, if $t < t'$ then $(t, x_i(t)) \rightsquigarrow (t', x_i(t'))$.
- 141 – Every pair of events whose timestamps are at least ε apart is totally ordered,
142 i.e., for all $1 \leq i, j \leq n$ and all $(t, x_i(t)) \in V_i$ and $(t', x_j(t')) \in V_j$, if $t + \varepsilon \leq t'$
143 then $(t, x_i(t)) \rightsquigarrow (t', x_j(t'))$.

144 *Example 2.* **TODO: distributed signal, happened-before relation**

145 **Definition 3.** Let (S, \rightsquigarrow) be a distributed signal of n signals, and $V = \bigcup_{i=1}^n V_i$
146 be the set of its events. A set $C \subseteq V$ is a consistent cut iff for every event in
147 C , all events that happened before it also belong to C , i.e., for all $e, e' \in V$, if
148 $e \in C$ and $e' \rightsquigarrow e$, then $e' \in C$.

149 We denote by $\mathbb{C}(T)$ the (infinite) set of consistent cuts at global time T .
 150 Given a consistent cut C , its *frontier* $\text{front}(C) \subseteq C$ is the set consisting of the
 151 last events in C of each signal, i.e., $\text{front}(C) = \bigcup_{i=1}^n \{(t, x_i(t)) \in V_i \cap C \mid \forall t' >$
 152 $t : (t', x_i(t')) \notin V_i \cap C\}$.

153 **Definition 4.** A consistent cut flow is a function $\text{ccf} : \mathbb{R}_{\geq 0} \rightarrow 2^V$ that maps a
 154 global clock value T to the frontier of a consistent cut at time T , i.e., $\text{ccf}(T) \in$
 155 $\{\text{front}(C) \mid C \in \mathbb{C}(T)\}$.

156 For all $T, T' \in \mathbb{R}_{\geq 0}$ and $1 \leq i \leq n$, if $T < T'$, then for every pair of
 157 events $(c_i(T), x_i(c_i(T))) \in \text{ccf}(T)$ and $(c_i(T'), x_i(c_i(T')))) \in \text{ccf}(T')$ we have
 158 $(c_i(T), x_i(c_i(T))) \rightsquigarrow (c_i(T'), x_i(c_i(T')))$. We denote by $\text{CCF}(S, \rightsquigarrow)$ the set of all
 159 consistent cut flows of the distributed signal (S, \rightsquigarrow) .

160 *Example 5. TODO: consistent cut, frontier, consistent cut flow*

161 Observe that a consistent cut flow of a distributed signal induces a vector
 162 of synchronous signals which can be evaluated using the standard semantics
 163 described in Section 2. Let (S, \rightsquigarrow) be a distributed signal of n signals x_1, \dots, x_n .
 164 A consistent cut flow $\text{ccf} \in \text{CCF}(S, \rightsquigarrow)$ yields a trace $w_{\text{ccf}} = (x'_1, \dots, x'_n)$ on a
 165 temporal domain $[0, D]$ where $D \in \mathbb{R}_{> 0}$ such that $(c_i(T), x_i(c_i(T))) \in \text{ccf}(T)$
 166 implies $x'_i(T) = x_i(c_i(T))$ for all $1 \leq i \leq n$ and $T \in [0, D]$. The set of traces of
 167 (S, \rightsquigarrow) is given by $\text{Tr}(S, \rightsquigarrow) = \{w_{\text{ccf}} \mid \text{ccf} \in \text{CCF}(S, \rightsquigarrow)\}$.

168 We define the satisfaction of an STL formula φ by a distributed signal (S, \rightsquigarrow)
 169 over a three-valued domain $\{\top, \perp, ?\}$. If the set of synchronous traces $\text{Tr}(S, \rightsquigarrow)$
 170 defined by a distributed signal (S, \rightsquigarrow) is contained in the set of traces allowed
 171 by the formula φ , then (S, \rightsquigarrow) satisfies φ . Similarly, if $\text{Tr}(S, \rightsquigarrow)$ has an empty
 172 intersection with the set of traces φ defines, then (S, \rightsquigarrow) violates φ . Otherwise,
 173 the evaluation is inconclusive since some traces satisfy the property and some
 174 violate it.

$$[(S, \rightsquigarrow) \models \varphi] = \begin{cases} \top & \text{if } \forall w \in \text{Tr}(S, \rightsquigarrow) : w \models \varphi \\ \perp & \text{if } \forall w \in \text{Tr}(S, \rightsquigarrow) : w \models \neg \varphi \\ ? & \text{otherwise} \end{cases}$$

175 3 Overapproximation of the STL Distributed Semantics

176 To address the computational overhead in exact distributed monitoring, we de-
 177 fine a new logic STL^+ whose syntax is the same as STL but semantics provide
 178 a sound approximation of the STL distributed semantics presented in Section 2.
 179 In essence, given a distributed signal (S, \rightsquigarrow) , STL^+ considers an overapproxima-
 180 tion $\text{Tr}^+(S, \rightsquigarrow)$ of the set $\text{Tr}(S, \rightsquigarrow)$ of synchronous traces. A signal (S, \rightsquigarrow) satisfies
 181 (resp. violates) an STL^+ formula φ iff all the traces in $\text{Tr}^+(S, \rightsquigarrow)$ belong to the
 182 language of φ (resp. $\neg \varphi$).

$$[(S, \rightsquigarrow) \models \varphi]_+ = \begin{cases} \top & \text{if } \forall w \in \text{Tr}^+(S, \rightsquigarrow) : w \models \varphi \\ \perp & \text{if } \forall w \in \text{Tr}^+(S, \rightsquigarrow) : w \models \neg \varphi \\ ? & \text{otherwise} \end{cases}$$

In Sections 4 and 5, we respectively define Tr^+ and present an algorithm to compute the semantics of STL^+ . We finally prove the correctness of our approach.

Theorem 6. *For every STL formula φ and every distributed signal (S, \rightsquigarrow) , if $[(S, \rightsquigarrow) \models \varphi]_+ = \top$ (resp. \perp) then $[(S, \rightsquigarrow) \models \varphi] = \top$ (resp. \perp).*

4 Overapproximation of Synchronous Traces

In this section, given a distributed signal (S, \rightsquigarrow) , we describe an overapproximation $\text{Tr}^+(S, \rightsquigarrow)$ of its set $\text{Tr}(S, \rightsquigarrow)$ of synchronous traces. First, we present the notion of *canonical segmentation*, a systematic way of partitioning the temporal domain of a given distributed signal to keep track of the partial asynchrony. Second, we introduce the notion of *value expressions*, sets of finite words representing how a signal behaves in a time interval. Finally, we define Tr^+ based on these notions, and show that it soundly approximates Tr .

Remark 7. We assume boolean signals in this section for convenience. The definitions and results presented here extend to real-valued signals because finite-length piecewise-constant signals will only use a finite number of values.

Canonical Segmentation Consider a boolean signal x with a rising edge at time $t > \varepsilon$. Due to clock skew, this edge occurs in the range $(t - \varepsilon, t + \varepsilon)$ from the monitor's point of view. This range is called an *uncertainty region* because in $(t - \varepsilon, t + \varepsilon)$ the monitor cannot tell the value of x precisely, but only that it changes from 0 to 1. Formally, given an edge $(t, x(t))$, we define $\theta_{\text{lo}}(x, t) = \max(0, t - \varepsilon)$ and $\theta_{\text{hi}}(x, t) = t + \varepsilon$ as the end points of the edge's uncertainty region.

Given a temporal domain $I = [0, d) \subset \mathbb{R}_{\geq 0}$, a *segmentation* of I is a partition of I into finitely many intervals I_1, \dots, I_k , called *segments*, of the form $I_j = [t_j, t_{j+1})$ such that $t_j < t_{j+1}$ for all $1 \leq j \leq k$. By extension, a segmentation of a collection of signals with the same temporal domain I is a segmentation of I .

Let (S, \rightsquigarrow) be a distributed signal of n signals. The *canonical segmentation* G_S of (S, \rightsquigarrow) is the segmentation of S where the end points of the segments coincide with the end points of its temporal domain and uncertainty regions. Formally, we define G_S as follows. For each signal x_i , let F_i be the set of end points of its uncertainty regions. Let $d' = \max(d, \max(\bigcup_{i=1}^n F_i))$, which corresponds to the duration of the distributed signal with respect to the monitor's clock. Let $F = \{0, d'\} \cup \bigcup_{i=1}^n F_i$ and let $(s_j)_{1 \leq j \leq |F|}$ be a nondecreasing sequence of clock values corresponding to the elements of F . Then, the canonical segmentation of (S, \rightsquigarrow) is $G_S = \{I_1, \dots, I_{|F|-1}\}$ where $I_j = [s_j, s_{j+1})$ for all $1 \leq j < |F|$.

Example 8. Let (S, \rightsquigarrow) be a distributed boolean signal with $S = (x_1, x_2)$ and $\varepsilon = 2$ over the temporal domain $[0, 8)$ as given in Figure 1. Both signals are initially 0. The signal x_1 has a rising edge at time 2 and a falling edge at time 5, while x_2 has a rising edge at time 3 and a falling edge at time 6. The uncertainty

regions of x_1 are $(0, 4)$ and $(3, 7)$, while those of x_2 are $(1, 5)$ and $(4, 8)$. Then,
 we have $F = \{0, 8\} \cup \{0, 1, 3, 4, 5, 7, 8\}$, and thus the canonical segmentation is
 $G_S = \{[0, 1), [1, 3), [3, 4), [4, 5), [5, 7), [7, 8)\}$.




Fig. 1. The signals x_1 (top, red, dashed) and x_2 (bottom, blue, solid) from Example 8. The edges are marked with solid balls and their uncertainty regions are given as semi-transparent boxes around the edges. The resulting canonical segmentation G_S is shown below the graphical representation of the signals.

Value Expressions Consider a boolean signal x with a rising edge with an uncertainty region of (t_1, t_2) . As discussed above, the monitor only knows that the value of x changes from 0 to 1 in this interval. We represent this knowledge as a finite word $v = 0 \cdot 1$. This representation is called a *value expression* and it encodes the uncertain behavior of an observed signal relative to the monitor. Formally, a value expression is an element of Σ^* where Σ is the finite alphabet of values the signal takes. Given a signal x and an edge $(t, x(t))$, the value expression corresponding to the uncertainty region $(\theta_{lo}(x, t), \theta_{hi}(x, t))$ is given by $v_{x,t} = v_- \cdot v_+$ where $v_- = \lim_{s \rightarrow t^-} x(s)$ and $v_+ = \lim_{s \rightarrow t^+} x(s)$. Let us remark that this definition is general because finite-length piecewise-constant real-valued signals will only have a finite number of values, making Σ finite.

Notice that (i) uncertainty regions may overlap, and (ii) the canonical segmentation may split an uncertainty region into multiple segments. Consider a signal x with a rising edge in $(1, 5)$ and a falling edge in $(4, 8)$. The corresponding value expressions are respectively $v_1 = 0 \cdot 1$ and $v_2 = 1 \cdot 0$. Notice that the behavior of x in the interval $[1, 4)$ can be expressed as $\text{prefix}(v_1)$, encoding whether

the rising edge has happened yet or not. Similarly, the behavior in $[4, 5)$ is given by $\text{suffix}(v_1) \cdot \text{prefix}(v_2)$, which captures whether the edges occur in this interval (thanks to prefixing and suffixing) and the fact that the rising edge happens before the falling edge (thanks to concatenation).

Formally, given a distributed signal (S, \rightsquigarrow) , we define a function $\gamma : S \times G_S \rightarrow 2^{\Sigma^*}$ that maps each signal and segment of the canonical segmentation to a set of value expressions, capturing the signal's potential behaviors in the given segment. Let x be a signal in S , and let R_1, \dots, R_m be its uncertainty regions where $R_i = (t_i, t'_i)$ and the corresponding value expression is v_i for all $1 \leq i \leq m$. Now, let $I \in G_S$ be a segment with $I = [s, s')$ and for each $1 \leq i \leq m$ define the set V_i of value expressions capturing how I relates with R_i as follows:

$$V_i = \begin{cases} \{v_i\} & \text{if } t_i = s \wedge s' = t'_i \\ \text{prefix}(v_i) & \text{if } t_i = s \wedge s' < t'_i \\ \text{suffix}(v_i) & \text{if } t_i > s \wedge s' = t'_i \\ \text{infix}(v_i) & \text{if } t_i > s \wedge s' < t'_i \\ \{\epsilon\} & \text{otherwise} \end{cases}$$

The last case happens only when $I \cap R_i$ is empty. We finally define γ as follows:

$$\gamma(x, I) = \text{destutter}(V_1 \cdot V_2 \cdot \dots \cdot V_m) \setminus \{\epsilon\}$$

Observe that $\gamma(x, I)$ contains all the potential behaviors of x in segment I by construction. However, it is potentially overapproximate. This is mainly because the sets V_1, \dots, V_m contain redundancy by definition and the concatenation does not guarantee that an edge is considered exactly once.

Example 9. Recall the distributed signal (S, \rightsquigarrow) in Example 8 and Figure 1. In Figure 2a, we show the value expressions corresponding to its uncertainty regions. For example, the falling edge of x_1 has an uncertainty region of $(3, 7)$, represented by the value expression $1 \cdot 0$. In Figure 2b, we give the function γ for (S, \rightsquigarrow) . For example, $\gamma(x_1, [3, 4))$ is obtained from $\text{suffix}(0 \cdot 1) \cdot \text{prefix}(1 \cdot 0)$ and $\gamma(x_2, [0, 1)) = \{0\}$.

Overapproximation of Tr Consider a distributed signal (S, \rightsquigarrow) of n signals, and let G_S be its canonical segmentation. We describe how the function γ defines a set $\text{Tr}^+(S, \rightsquigarrow)$ of synchronous traces that overapproximates the set $\text{Tr}(S, \rightsquigarrow)$.

Let $x \in S$ and x' be two signals with the same temporal domain, and let $I = [s, s')$ be a segment in G_S . Let $(t_1, x'(t_1)), \dots, (t_\ell, x'(t_\ell))$ be the edges of x' in segment I with $t_i < t_{i+1}$ for all $1 \leq i < \ell$. The signals x and x' are *consistent in I* iff the value expression $x'(s) \cdot x'(t_1) \cdot \dots \cdot x'(t_\ell)$ belongs to $\gamma(x, I)$. Moreover, x and x' are *consistent* iff they are consistent in I for all $I \in G_S$. Now, let $S = (x_1, \dots, x_n)$ and define $\text{Tr}^+(S, \rightsquigarrow)$ as follows:

$$\text{Tr}^+(S, \rightsquigarrow) = \{(x'_1, \dots, x'_n) \mid x_i \text{ and } x'_i \text{ are consistent for all } 1 \leq i \leq n\}$$




Fig. 2. (a) The uncertainty regions of the distributed signal in Example 8 and the corresponding value expressions. (b) The tabular representation of the function γ for the given distributed signal.

Example 10. Recall the distributed signal (S, \rightsquigarrow) in Example 8 whose γ function is given in Figure 2b. Consider the synchronous trace $w \in \text{Tr}(S, \rightsquigarrow)$ where the rising edges of both signals occur at time 3 and the falling edges at time 5. One can verify that $w \in \text{Tr}^+(S, \rightsquigarrow)$ since for each $i \in \{1, 2\}$ the value expression 1 is contained in $\gamma(x_i, [3, 4))$ and $\gamma(x_i, [4, 5))$ while 0 is contained in the remaining sets γ maps x_i to.

Now, consider a synchronous trace (x'_1, x'_2) where both signals are initially 0, have rising edges at time 2 and 3.5, and falling edges at time 3 and 5. Evidently, this trace does not belong to $\text{Tr}(S, \rightsquigarrow)$ since x'_1 and x'_2 have more edges than x_1 and x_2 . Nonetheless, it belongs to $\text{Tr}^+(S, \rightsquigarrow)$ since x'_1 and x'_2 are respectively consistent with x_1 and x_2 . To witness, notice that for each $i \in \{1, 2\}$ the value expression $0 \cdot 1$ is contained in $\gamma(x_i, [1, 3))$ and $\gamma(x_i, [3, 4))$, the expression 1 is contained in $\gamma(x_i, [4, 5))$, and 0 is contained in the remaining sets γ maps x_i to.

Finally, we prove Tr^+ overapproximates Tr .

Lemma 11. For every distributed signal (S, \rightsquigarrow) , we have $\text{Tr}(S, \rightsquigarrow) \subseteq \text{Tr}^+(S, \rightsquigarrow)$.

5 Monitoring Algorithm

In this section, given a distributed signal (S, \rightsquigarrow) , we describe an algorithm to compute $[(S, \rightsquigarrow) \models \varphi]_+$. The algorithm makes use of the function γ defined in Section 4 without explicitly computing $\text{Tr}^+(S, \rightsquigarrow)$. To achieve this, we first describe the notion of *asynchronous product* of value expressions to capture potential interleavings within segments. We continue with the evaluation of *untimed operators* and then *timed operators*. Finally, we conclude with putting all these

295 together to compute the *semantics* of STL^+ and discuss an efficient implemen-
 296 tation of the monitoring algorithm using *bit vectors* to represent and manipulate
 297 sets of boolean value expressions.

298 *Remark 12.* For the sake of convenience, we focus on boolean signals for the rest
 299 of the section. Note that asynchronous products and the algorithm to compute
 300 $[(S, \rightsquigarrow) \models \varphi]_+$ can be extended to value expressions over arbitrary finite alpha-
 301 bets, e.g., encoding real-valued signals. This allows us to express more complex
 302 properties where atomic propositions can be functions of real-valued signals.

303 **Asynchronous Products** Consider the value expressions $u_1 = 0 \cdot 1$ and
 304 $u_2 = 1 \cdot 0$ encoding the behaviors of two signals within a segment. Due to par-
 305 tial asynchrony, the behaviors within segments can be seen as completely asyn-
 306 chronous. To capture the potential interleavings of these behaviors, we consider
 307 how the values in u_1 and u_2 can align. In particular, there are three potential
 308 alignments: (i) the rising edge of u_1 happens before the falling edge of u_2 , (ii)
 309 the falling edge of u_2 happens before the rising edge of u_1 , and (iii) the two
 310 edges happen simultaneously. We respectively represent these with the tuples
 311 $(011, 110)$, $(001, 100)$, and $(01, 10)$ where the first component encodes u_1 and
 312 the second u_2 . Formally, given two value expressions u_1 and u_2 , we define their
 313 *asynchronous product* as follows:

$$u_1 \otimes u_2 = \{\text{destutter}(v_1, v_2) \mid v_i \in \text{stutter}_k(u_i), k = |u_1| + |u_2| - 1, i \in \{1, 2\}\}$$

314 Moreover, given two sets L_1 and L_2 of value expressions, we define the following:

$$L_1 \otimes L_2 = \{u_1 \otimes u_2 \mid u_1 \in L_1, u_2 \in L_2\}$$

315 Asynchronous products of value expressions allow us to lift value expressions
 316 to satisfaction signals of formulas.

317 *Example 13.* Recall the distributed signal (S, \rightsquigarrow) in Example 8 and its γ function
 318 given in Figure 2b. Suppose we want to compute the value expressions encoding
 319 the satisfaction of $x_1 \wedge x_2$ in the segment $[1, 3)$. We can achieve this by first com-
 320 puting the asynchronous product $\gamma(x_1, [3, 4)) \otimes \gamma(x_2, [3, 4))$, and then computing
 321 the bitwise conjunction of each pair in the set. For example, considering the ex-
 322 pression $0 \cdot 1 \cdot 0$ for x_1 and $0 \cdot 1$ for x_2 , the product contains the pair $(010, 011)$.
 323 Taking the bitwise conjunction of this pair gives us the expression $0 \cdot 1 \cdot 0$ as a
 324 potential behavior for the satisfaction of $x_1 \wedge x_2$ in this segment.

325 **Untimed Operations** As hinted in Example 13, to compute the semantics, we
 326 apply bitwise operations on value expressions and their asynchronous products
 327 to transform them into encodings of satisfaction signals of formulas. Consider
 328 the distributed signal (S, \rightsquigarrow) in Example 8 and suppose we want to compute
 329 $[(S, \rightsquigarrow) \models \Diamond(x_1 \wedge x_2)]_+$. To achieve this, we first compute for each segment in
 330 G_S the set of value expressions for the satisfaction of $x_1 \wedge x_2$, and then from these

331 compute that of $\Diamond(x_1 \wedge x_2)$. This compositional approach allows us to evaluate
 332 arbitrary STL⁺ formulas.

333 First, we define bitwise operations on boolean value expressions encoding
 334 atomic propositions. Then, we use these to evaluate (untimed) STL formulas
 335 over sets of value expressions.

336 Let u and v be boolean value expressions of length ℓ . We denote by $u \& v$ the
 337 bitwise-and operation, by $u \mid v$ the bitwise-or, and by $\sim u$ the bitwise-negation.
 338 In addition, we define the *bitwise strong until* operator as follows:

$$u \mathbf{U}^0 v = \left(\max_{i \leq j \leq \ell} \left(\min \left(v[j], \min_{i \leq k \leq j} u[k] \right) \right) \right)_{1 \leq i \leq \ell}$$

339 As usual, we derive *bitwise eventually* as $\mathbf{E}u = 1^\ell \mathbf{U}^0 u$, *bitwise always* as $\mathbf{A}u =$
 340 $\sim(\mathbf{E}\sim u)$, and *bitwise weak until* as $u \mathbf{U}^1 v = (u \mathbf{U}^0 v) \mid (\mathbf{A}u)$. The distinction be-
 341 tween \mathbf{U}^0 and \mathbf{U}^1 will be useful later when we evaluate a formula segment by
 342 segment. We remark that the definitions of these operators coincide with the
 343 robustness semantics of (discrete time) STL. Finally, note that the output of
 344 these operations is a value expression of length ℓ . For example, if $u = 010$, we
 345 have $\mathbf{E}u = 110$ and $\mathbf{A}u = 000$.

346 Let (S, \rightsquigarrow) be a distributed signal. Consider an atomic proposition $p \in \mathbf{AP}$
 347 encoded as $x_p \in S$ and let φ_1, φ_2 be two STL formulas. We define the evaluation
 348 of untimed formulas with respect to (S, \rightsquigarrow) and a segment $I \in G_S$ inductively:

$$\begin{aligned} \llbracket (S, \rightsquigarrow), I \models p \rrbracket &= \gamma(x_p, I) \\ \llbracket (S, \rightsquigarrow), I \models \neg \varphi_1 \rrbracket &= \{\sim u \mid u \in \llbracket (S, \rightsquigarrow), I \models \varphi_1 \rrbracket\} \\ \llbracket (S, \rightsquigarrow), I \models \varphi_1 \wedge \varphi_2 \rrbracket &= \{u_1 \& u_2 \mid (u_1, u_2) \in \llbracket (S, \rightsquigarrow), I \models \varphi_1 \rrbracket \otimes \llbracket (S, \rightsquigarrow), I \models \varphi_2 \rrbracket\} \\ \llbracket (S, \rightsquigarrow), I \models \varphi_1 \mathbf{U}^a \varphi_2 \rrbracket &= \{u_1 \mathbf{U}^a u_2 \mid (u_1, u_2) \in \llbracket (S, \rightsquigarrow), I \models \varphi_1 \rrbracket \otimes \llbracket (S, \rightsquigarrow), I \models \varphi_2 \rrbracket, \\ &\quad a \in \text{first}(\llbracket (S, \rightsquigarrow), I' \models \varphi_1 \mathbf{U} \varphi_2 \rrbracket)\} \end{aligned}$$

349 where I' is the segment that follows I in G_S , if it exists. For completeness, for
 350 every formula φ we define $\llbracket (S, \rightsquigarrow), I' \models \varphi \rrbracket = \{0\}$ when $I' \notin G_S$. When I is the
 351 first segment in G_S , we simply write $\llbracket (S, \rightsquigarrow) \models \varphi \rrbracket$. Similarly as above, we can
 352 use the standard derived operators to compute the corresponding sets of value
 353 expressions. Intuitively, for a given formula and a segment, the evaluation above
 354 produces a set of value expressions encoding the formula's satisfaction within
 355 the segment.

356 *Example 14.* Recall the distributed signal (S, \rightsquigarrow) in Example 8 and its γ function
 357 given in Figure 2b. Suppose we want to compute $\llbracket (S, \rightsquigarrow), [5, 7] \models \Diamond(x_1 \wedge x_2) \rrbracket$.
 358 First, we compute $\llbracket (S, \rightsquigarrow), [5, 7] \models x_1 \wedge x_2 \rrbracket$ by computing the bitwise conjunc-
 359 tion over the asynchronous product $\gamma(x_1, [5, 7]) \otimes \gamma(x_2, [5, 7])$ and destuttering.
 360 For example, since $010 \in \gamma(x_1, [5, 7])$ and $01 \in \gamma(x_2, [5, 7])$, the pair $(0010, 0111)$
 361 is in the product, whose conjunction gives us 010 after destuttering. Repeating
 362 this for the rest, we obtain $\llbracket (S, \rightsquigarrow), [5, 7] \models x_1 \wedge x_2 \rrbracket = \{0, 01, 010, 1, 10\}$. Fi-
 363 nally, we compute $\llbracket (S, \rightsquigarrow), [5, 7] \models \Diamond(x_1 \wedge x_2) \rrbracket$ by applying each expression in
 364 $\llbracket (S, \rightsquigarrow), [5, 7] \models x_1 \wedge x_2 \rrbracket$ the bitwise eventually operator and destuttering. The

365 resulting set $\{0, 1, 10\}$ encodes the satisfaction signal of $\Diamond(x_1 \wedge x_2)$ in $[5, 7)$.
 366 Note that we do not need to consider the evaluation of the next segment for the
 367 eventually operator since $\llbracket (S, \rightsquigarrow), [7, 8) \models x_1 \wedge x_2 \rrbracket = \{0\}$.

368 **Timed Operations** Handling timed operations requires a closer inspection as
 369 value expressions are untimed by definition. We address this issue by considering
 370 how a given evaluation interval relates with a given segmentation. For example,
 371 take a segmentation $G_S = \{[0, 4), [4, 6), [6, 10)\}$ and an evaluation interval $J =$
 372 $[0, 5)$. Suppose we are interested in how a signal $x \in S$ behaves with respect to
 373 J over the first segment $I = [0, 4)$. First, to see how J relates with G_S with
 374 respect to $I = [0, 4)$, we “slide” the interval J over $I \oplus J = [0, 9)$ and consider
 375 the different ways it intersects the segments in G_S . Initially, J covers the entire
 376 segment $[0, 4)$ and the beginning of $[4, 6)$, for which the potential behaviors of
 377 x are captured by the set $\gamma(x, [0, 4)) \cdot \text{prefix}(\gamma(x, [4, 6)))$. Now, if we slide the
 378 window and take $J' = [3, 7)$, the window covers the ending of $[0, 4)$, the entire
 379 $[4, 6)$, and the beginning of $[6, 10)$, for which the potential behaviors are captured
 380 by the set $\text{suffix}(\gamma(x, [0, 4))) \cdot \gamma(x, [4, 6)) \cdot \text{prefix}(\gamma(x, [6, 9)))$. We call these sets the
 381 *profiles* of J and J' with respect to (S, \rightsquigarrow) , x , and I .

382 Let (S, \rightsquigarrow) be a distributed signal, $I \in G_S$ be a segment, and φ be an STL
 383 formula. Let us introduce the notation we use in the definition below. First, we
 384 abbreviate the set $\llbracket (S, \rightsquigarrow), I \models \varphi \rrbracket$ of value expressions as $\tau_{\varphi, I}$. Second, given an
 385 interval K , we respectively denote by l_K and r_K its left and right end points.
 386 Third, recall that we denote by F the set of end points of G_S (see Section 4).
 387 Given an interval J , we define the *profile* of J with respect to (S, \rightsquigarrow) , φ , and I
 388 as follows.

$$\text{profile}((S, \rightsquigarrow), \varphi, I, J) = \begin{cases} \text{prefix}(\tau_{\varphi, I}) & \text{if } l_I = l_J \wedge r_I > r_J \\ \text{infix}(\tau_{\varphi, I}) & \text{if } l_I < l_J \wedge r_I > r_J \\ \tau_{\varphi, I} \cdot \kappa(\varphi, I, J) & \text{if } l_I = l_J \wedge r_I \leq r_J \wedge r_J \in F \setminus J \\ \tau_{\varphi, I} \cdot \kappa(\varphi, I, J) \cdot \text{first}(\tau_{\varphi, I'}) & \text{if } l_I = l_J \wedge r_I \leq r_J \wedge r_J \in F \cap J \\ \tau_{\varphi, I} \cdot \kappa(\varphi, I, J) \cdot \text{prefix}(\tau_{\varphi, I'}) & \text{if } l_I = l_J \wedge r_I \leq r_J \wedge r_J \notin F \\ \text{suffix}(\tau_{\varphi, I}) \cdot \kappa(\varphi, I, J) & \text{if } l_I < l_J < r_I \leq r_J \wedge r_J \in F \setminus J \\ \text{suffix}(\tau_{\varphi, I}) \cdot \kappa(\varphi, I, J) \cdot \text{first}(\tau_{\varphi, I'}) & \text{if } l_I < l_J < r_I \leq r_J \wedge r_J \in F \cap J \\ \text{suffix}(\tau_{\varphi, I}) \cdot \kappa(\varphi, I, J) \cdot \text{prefix}(\tau_{\varphi, I'}) & \text{if } l_I < l_J < r_I \leq r_J \wedge r_J \notin F \\ \{\epsilon\} & \text{otherwise} \end{cases}$$

389 where we assume J is trimmed to fit the temporal domain of S and $I' \in G_S$ is
 390 such that $r_J \in I'$. Moreover, $\kappa(\varphi, I, J)$ is the concatenation $\tau_{\varphi, I_1} \cdot \dots \cdot \tau_{\varphi, I_m}$ such
 391 that I, I_1, \dots, I_m, I' are consecutive segments in G_S . If I_1, \dots, I_m do not exist,
 392 we let $\kappa(\varphi, I, J) = \{\epsilon\}$. Note that the last case happens when $I \cap J$ is empty.
 393 We now formalize the intuitive approach of “sliding” J over the segmentation to
 394 obtain the various profiles it produces as follows.

$$\text{pfs}((S, \rightsquigarrow), \varphi, I, J) = \{\text{destutter}(\text{profile}((S, \rightsquigarrow), \varphi, I, J')) \mid J' \subseteq I \oplus J, J' \sim J\}$$

395 where $J' \sim J$ holds when $|J'| = |J|$ and J' contains an end point (left or
 396 right) iff J does so. Note that although infinitely many intervals J' satisfy the

conditions given above (due to denseness of time), the set defined by **pfs** is finite. We demonstrate this and the computation of **pfs** in Example 15 and Figure 3.

Example 15. Recall the distributed signal (S, \rightsquigarrow) in Example 8 and its γ function given in Figure 2b. We demonstrate the computation of $\mathbf{pfs}((S, \rightsquigarrow), x_1, [1, 3], [0, 1])$. Intuitively, sliding the interval $[0, 1]$ over the window $[1, 3] \oplus [0, 1]$ (as shown in Figure 3) gives us the following sets:

$$\begin{aligned} P_1 &= \text{destutter}(\text{prefix}(\gamma(x_1, [1, 3]))) = \{0, 01, 1\} \\ P_2 &= \text{destutter}(\text{infix}(\gamma(x_1, [1, 3]))) = \{0, 01, 1\} \\ P_3 &= \text{destutter}(\text{suffix}(\gamma(x_1, [1, 3]))) = \{0, 01, 1\} \\ P_4 &= \text{destutter}(\text{suffix}(\gamma(x_1, [1, 3])) \cdot \text{prefix}(\gamma(x_1, [3, 4]))) \\ &= \{0, 01, 010, 0101, 01010, 1, 10, 101, 1010\} \end{aligned}$$

Therefore, we obtain $\mathbf{pfs}((S, \rightsquigarrow), x_1, [1, 3], [0, 1]) = \{P_1, P_2, P_3, P_4\}$. This set overapproximates the potential behaviors of x_1 , for all $t \in [1, 3]$, in the interval $t \oplus [0, 1]$.




Fig. 3. The profiles of $J = [0, 1]$ with respect to $x_1 \in S$ of Example 8. The four representative intervals of each profile is shown with solid black lines below the tabular representation of γ for x_1 .

Let φ_1 and φ_2 be two STL formulas. Intuitively, once we have the profiles of a given interval J with respect to φ_1 and φ_2 , we can evaluate the corresponding untimed formulas on the product of these profiles and concatenate them. Formally, we handle the evaluation of timed formulas inductively as follows.

$$\llbracket (S, \rightsquigarrow), I \models \varphi_1 \mathcal{U}_J \varphi_2 \rrbracket = \text{destutter}(\{u_1 \mathcal{U}^0 u_2 \mid (u_1, u_2) \in P_1 \otimes Q_1\} \cdot \dots \cdot \{u_1 \mathcal{U}^0 u_2 \mid (u_1, u_2) \in P_k \otimes Q_k\})$$

where $\mathbf{pfs}((S, \rightsquigarrow), \varphi_1, I, J) = \{P_1, \dots, P_k\}$ and $\mathbf{pfs}((S, \rightsquigarrow), \varphi_2, I, J) = \{Q_1, \dots, Q_k\}$ such that the intervals producing P_i and Q_i respectively start before those producing P_{i+1} and Q_{i+1} for all $1 \leq i < k$.

Example 16. Let (S, \rightsquigarrow) be as in Example 8 and its γ function as given in Figure 2b. We demonstrate the evaluation of the timed formula $\Diamond_{[0,1]} x_1$ over the

segment $[1, 3)$. Recall from Example 15 the set $\text{pfs}((S, \rightsquigarrow), x_1, [1, 3), [0, 1)) = \{P_1, P_2, P_3, P_4\}$ of profiles. First, we apply the bitwise eventually operator to each value expression in each of these profiles separately: $\{Eu \mid u \in P_1\} = \{0, 1\}$, $\{Eu \mid u \in P_2\} = \{0, 1\}$, $\{Eu \mid u \in P_3\} = \{0, 1\}$, and $\{Eu \mid u \in P_4\} = \{0, 10, 1\}$. Then, we concatenate these sets and destutter to obtain the following:

$$\llbracket (S, \rightsquigarrow), [1, 3) \rrbracket \models \Diamond_{[0,1)x_1} = \{0, 01, 010, 0101, 01010, 1, 10, 101, 1010\}$$

Computing the Semantics of STL⁺ Putting it all together, given a distributed signal (S, \rightsquigarrow) and an STL⁺ formula φ , we can compute $\llbracket (S, \rightsquigarrow) \rrbracket \models \varphi_+$ thanks to the following theorem.

Theorem 17. *For every distributed signal (S, \rightsquigarrow) , we have $\llbracket (S, \rightsquigarrow) \rrbracket \models \varphi_+ = \top$ (resp. \perp , $?$) iff $\text{first}(\llbracket (S, \rightsquigarrow) \rrbracket \models \varphi) = \{1\}$ (resp. $\{0\}$, $\{0, 1\}$).*

Sets of Boolean Value Expressions as Bit Vectors Evidently, asynchronous products are expensive to compute. Our implementation of the algorithm we describe in this section relies on the following observation: Sets of boolean value expressions and their operations can be efficiently implemented through bit vectors. Intuitively, to represent such a set, we can encode each element using its first bit and its length since value expressions are boolean and always destuttered. Moreover, to evaluate untimed operations on such sets, we only need to know the maximal lengths of the four possible types of expressions ($0 \dots 0$, $0 \dots 1$, $1 \dots 0$, and $1 \dots 1$) and whether the set contains 0 or 1 (to handle some edge cases). This is because value expressions corresponding to same segments can be seen as completely asynchronous and the possible interleavings obtained from shorter expressions can be obtained from longer ones. This approach enables, for example, an algorithm for conjunction of sets of value expressions that runs in $O(|u| + |v|)$ time where u and v are the longest expressions in the two sets. Note that the same idea also applies to untimed temporal operators.

6 Experimental Evaluation

TODO

7 Conclusion

TODO

References

1. Maler, O., Nickovic, D.: Monitoring properties of analog and mixed-signal circuits. Int. J. Softw. Tools Technol. Transf. **15**(3), 247–268 (2013). <https://doi.org/10.1007/s10009-012-0247-9>

448 **Appendix**

449 *Proof (Proof of Theorem 6).* **TODO**

450 *Proof (Proof of Lemma 11).* **TODO**

451 *Proof (Proof of Theorem 17).* **TODO**