

Approximate Distributed Monitoring under Partial Asynchrony: Balancing Speed with Accuracy

Author(s)

Affiliation(s)

Abstract. In distributed systems with processes that do not share a global clock, *partial asynchrony* is achieved by clock synchronisation that guarantees bounded clock skew among all applications. Existing solutions for distributed runtime verification under partial asynchrony against temporal logic specifications are exact but suffer from significant computational overhead. In this paper, we propose an *approximate* distributed monitoring algorithm for Signal Temporal Logic (STL) that mitigates this issue by abstracting away potential interleaving behaviors. This conservative abstraction enables a significant speedup of the distributed monitors, albeit with a trade-off in accuracy. We address this trade-off with a methodology that combines our approximate monitor with its exact counterpart, resulting in enhanced monitoring efficiency without sacrificing precision. We validate our approach with multiple experiments, showcasing its effectiveness and efficacy on both a real-world application and synthetic examples.

You can leave notes using the command `\firstName{note}`.
Page limit: 16 + 4 (appendix)

1 Introduction

Distributed systems are networks of independent agents that work together to achieve a common objective. Distributed systems are everywhere around us and come in many different forms. For example, cloud computing uses distribution of resources and services over the internet to offer to their users a scalable infrastructure with transparent on-demand access to computing power and storage. Swarms of drones represent another family of distributed systems where individual drones collaborate to accomplish tasks like surveillance, search and rescue, or package delivery. While each drone operates independently, it also communicates and coordinates with others to successfully achieve their common objectives. The individual agents in a distributed system typically do not share a global clock. To coordinate actions across multiple agents, clock synchronisation is often needed. While perfect clock synchronisation is impractical due to network latency and node failures, algorithms such as the Network Time Protocol (NTP) allow agents

to maintain a *bounded skew* between the synchronised clocks. In that case, we say that a distributed system has *partial asynchrony*.

Formal verification of distributed system is a notoriously hard problem, due to the combinatorial explosion of all possible interleavings in the behaviors collected from individual agents. *Runtime verification (RV)* provides a more pragmatic approach, in which a monitor observes a behavior of a distributed system and checks its correctness against a formal specification. The problem of distributed RV under partial asynchrony assumption has been studied for Linear Temporal Logic (LTL) and Signal Temporal Logic (STL) specification languages. The proposed solutions use Satisfiability-Modulo-Theory (SMT) solving to provide sound and complete distributed monitoring procedures. Although distributed RV monitors consume only a single distributed behavior at a time, this behavior can nevertheless have an excessive number of possible interleavings. Hence, the exact distributed monitors from the literature can still suffer from significant computational overhead.

To mitigate this issue, we present an approach for *approximate* RV of STL specifications under partial asynchrony. In essence, we abstract away potential interleavings in distributed behaviors in a conservative manner, resulting in an effective over-approximation of global behaviors. This abstraction simplifies the representation of distributed behaviors and the monitoring operations required to evaluate temporal specifications. There is an inevitable trade-off in approximate RV – gains in the monitoring speed-up may result in reduced accuracy. For some applications, reduced accuracy may not be acceptable. Therefore, we propose a methodology that combines our approximate monitors with their exact counterparts, with the aim to benefit from the enhanced monitoring efficiency without sacrificing precision. We implemented our approach in a prototype tool and performed thorough evaluations on both synthetic and real-world case studies. We first demonstrated that our approximate monitors achieve speed-ups of several orders of magnitudes compared to the exact SMT-based distributed RV solution. We empirically characterized the classes of specifications and behaviors for which our approximate monitoring approach achieves good precision. We finally showed that by combining exact and approximate distributed RV, there is still a significant efficiency gain on average without the sacrifice of the precision, even in cases where approximate monitors have low accuracy.

2 Preliminaries

We define boolean domain $\mathbb{B} = \{\perp, \top\}$ as the set of boolean truth values, where $\perp < \top$ and they complement each other, i.e., $\overline{\perp} = \top$ and $\overline{\top} = \perp$. We denote by \mathbb{R} the set of reals, $\mathbb{R}_{\geq 0}$ the set of nonnegative reals, and $\mathbb{R}_{> 0}$ the set of positive reals. An interval $I \subseteq \mathbb{R}$ of reals with the end points $a < b$ has length $|b - a|$.

Let Σ be a finite *alphabet*. We denote by Σ^* the set of finite words over Σ and by ϵ the empty word. For $u \in \Sigma^*$, we respectively write $\text{prefix}(u)$ and $\text{suffix}(u)$ for the sets of prefixes and suffixes of u . We also let $\text{infix}(u) = \{v \in \Sigma^* \mid \exists x, y \in \Sigma^* : u = xvy\}$. For a nonempty word $u \in \Sigma^*$ and $1 \leq i \leq |u|$,

we denote by $u[i]$ the i th letter of u , by $u[..i]$ the prefix of u of length i , and by $u[i..]$ the suffix of u of length $|u| - i + 1$. Given $u \in \Sigma^*$ and $\ell \geq 1$, we denote by u^ℓ the word obtained by concatenating u by itself $\ell - 1$ times. Moreover, given $L \subseteq \Sigma^*$, we define $\text{first}(L) = \{u[0] \mid u \in L\}$. For sets $L_1, L_2 \in \Sigma^*$ of words, we let $L_1 \cdot L_2 = \{u \cdot v \mid u \in L_1, v \in L_2\}$. For tuples (u_1, \dots, u_m) and (v_1, \dots, v_m) of words, we let $(u_1, \dots, u_m) \cdot (v_1, \dots, v_m) = (u_1 v_1, \dots, u_m v_m)$.

We define the function $\text{destutter} : \Sigma^* \rightarrow \Sigma^*$ inductively as follows. For all $\sigma \in \Sigma \cup \{\epsilon\}$, let $\text{destutter}(\sigma) = \sigma$. For all $u \in \Sigma^*$ such that $u = \sigma_1 \sigma_2 v$ for some $\sigma_1, \sigma_2 \in \Sigma$ and $v \in \Sigma^*$, let (i) $\text{destutter}(u) = \text{destutter}(\sigma_2 v)$ if $\sigma_1 = \sigma_2$, and (ii) $\text{destutter}(u) = \sigma_1 \cdot \text{destutter}(\sigma_2 v)$ otherwise. By extension, for a set $L \subseteq \Sigma^*$ of finite words, we write $\text{destutter}(L) = \{\text{destutter}(u) \mid u \in L\}$. Given a tuple $(u_1, \dots, u_m) = (\sigma_{1,1} \sigma_{1,2} v_1, \dots, \sigma_{m,1} \sigma_{m,2} v_m)$ of finite words of the same length, we define $\text{destutter}(u_1, \dots, u_m)$ as expected: (i) $\text{destutter}(u_1, \dots, u_m) = \text{destutter}(\sigma_{1,2} v_1, \dots, \sigma_{m,2} v_m)$ if $\sigma_{i,1} = \sigma_{i,2}$ for all $1 \leq i \leq m$, and (ii) $\text{destutter}(u_1, \dots, u_m) = (\sigma_{1,1}, \dots, \sigma_{m,1}) \cdot \text{destutter}(\sigma_{1,2} v_1, \dots, \sigma_{m,2} v_m)$ otherwise.

Moreover, given an integer $k \geq 0$, we define $\text{stutter}_k : \Sigma^* \rightarrow \Sigma^*$ such that $\text{stutter}_k(u) = \{v \in \Sigma^* \mid |v| = k \wedge \text{destutter}(v) = \text{destutter}(u)\}$ if $k \geq |\text{destutter}(u)|$, and $\text{stutter}_k(u) = \emptyset$ otherwise.

Signal Temporal Logic (STL) Let $A, B \subset \mathbb{R}$. A function $f : A \rightarrow B$ is *right-continuous* iff $\lim_{a \rightarrow c^+} f(a) = f(c)$ for all $c \in A$, and *non-Zeno* iff for every bounded interval $I \subseteq A$ there are finitely many $a \in I$ such that f is not continuous at a .

A *signal* is a right-continuous, non-Zeno, piecewise-constant function $x : [0, d) \rightarrow \mathbb{R}$ where $d \in \mathbb{R}_{>0}$ is the duration of x and $[0, d)$ is its temporal domain. Let $x : [0, d) \rightarrow \mathbb{R}$ be a signal. An *event* of x is a pair $(t, x(t))$ where $t \in [0, d)$. An *edge* of x is an event $(t, x(t))$ such that $\lim_{s \rightarrow t^-} x(s) \neq \lim_{s \rightarrow t^+} x(s)$. In particular, an edge is *rising* if $\lim_{s \rightarrow t^-} x(s) < \lim_{s \rightarrow t^+} x(s)$, and it is *falling* otherwise. A signal $x : [0, d) \rightarrow \mathbb{R}$ can be represented finitely by its initial value and edges: if x has m edges, then $x = (t_0, v_0)(t_1, v_1) \dots (t_m, v_m)$ such that $t_0 = 0$, $t_{i-1} < t_i$, and (t_i, v_i) is an edge of x for all $1 \leq i \leq m$.

Let AP be a set of atomic propositions. The syntax is given by the following grammar where $p \in \text{AP}$ and $I \subseteq \mathbb{R}_{\geq 0}$ is an interval.

$$\varphi := p \mid \neg \varphi \mid \varphi \wedge \varphi \mid \varphi \mathcal{U}_I \varphi$$

A *trace* $w = (x_1, \dots, x_n)$ is a finite vector of signals. We express atomic propositions as functions of trace values at a time point t , i.e., a proposition $p \in \text{AP}$ over a trace $w = (x_1, \dots, x_n)$ is defined as $f_p(x_1(t), \dots, x_n(t)) > 0$ where $f_p : \mathbb{R}^n \rightarrow \mathbb{R}$ is a function. Given intervals $I, J \subseteq \mathbb{R}_{\geq 0}$, we define $I \oplus J = \{i + j \mid i \in I \wedge j \in J\}$, and we simply write t for the singleton set $\{t\}$.

Below we recall the finite-trace qualitative semantics of STL defined over \mathbb{B} [1]. Let $d \in \mathbb{R}_{>0}$ and $w = (x_1, \dots, x_n)$ with $x_i : [0, d) \rightarrow \mathbb{R}$ for all $1 \leq i \leq n$. Let φ_1, φ_2 be STL formulas and let $t \in [0, d)$.

$$\begin{aligned}
[w, t \models p]_{\text{STL}} &\iff f_p(x_1(t), \dots, x_n(t)) > 0 \\
[w, t \models \neg\varphi_1]_{\text{STL}} &\iff \overline{[w, t \models \varphi_1]_{\text{STL}}} \\
[w, t \models \varphi_1 \wedge \varphi_2]_{\text{STL}} &\iff [w, t \models \varphi_1]_{\text{STL}} \wedge [w, t \models \varphi_2]_{\text{STL}} \\
[w, t \models \varphi_1 \mathcal{U}_I \varphi_2]_{\text{STL}} &\iff \exists t' \in (t \oplus I) \cap [0, d) : \\
&\quad [w, t' \models \varphi_2]_{\text{STL}} \wedge \forall t'' \in (t, t') : [w, t'' \models \varphi_1]_{\text{STL}}
\end{aligned}$$

We simply write $[w \models \varphi]$ for $[w, 0 \models \varphi]$. We additionally use the following standard abbreviations: **false** = $p \wedge \neg p$, **true** = $\neg \text{false}$, $\varphi_1 \vee \varphi_2 = \neg(\neg\varphi_1 \wedge \neg\varphi_2)$, $\Diamond_I \varphi = \text{true} \mathcal{U}_I \varphi$, and $\Box_I \varphi = \neg \Diamond_I \neg \varphi$. Moreover, the untimed temporal operators are defined through their timed counterparts on the interval $(0, \infty)$, e.g., $\varphi_1 \mathcal{U} \varphi_2 = \varphi_1 \mathcal{U}_{(0, \infty)} \varphi_2$.

Distributed Semantics of STL We consider an asynchronous and loosely-coupled message-passing system of $n \geq 2$ reliable agents producing a set of signals x_1, \dots, x_n , where for some $d \in \mathbb{R}_{>0}$ we have $x_i : [0, d) \rightarrow \mathbb{R}$ for all $1 \leq i \leq n$. The agents do not share memory or a global clock. Only to formalize statements, we speak of a *hypothetical* global clock and denote its value by T . For local time values, we use the lowercase letters t and s .

For a signal x_i , we denote by V_i the set of its events, by E_i^\uparrow the set of its rising edges, and by E_i^\downarrow that of falling edges. Moreover, we let $E_i = E_i^\uparrow \cup E_i^\downarrow$. We represent the local clock of the i th agent as an increasing and divergent function $c_i : \mathbb{R}_{\geq 0} \rightarrow \mathbb{R}_{\geq 0}$ that maps a global time T to a local time $c_i(T)$. We denote by c_i^{-1} the inverse of the local clock function c_i .

We assume that the system is *partially synchronous*: the agents use a clock synchronization algorithm that guarantees a bounded clock skew with respect to the global clock, i.e., $|c_i(T) - c_j(T)| < \varepsilon$ for all $1 \leq i, j \leq N$ and $T \in \mathbb{R}_{\geq 0}$, where $\varepsilon \in \mathbb{R}_{>0}$ is the maximum clock skew.

Definition 1. A distributed signal is a pair (S, \rightsquigarrow) , where $S = (x_1, \dots, x_n)$ is a vector of signals and \rightsquigarrow is the happened-before relation between events in signals extended with the partial synchrony assumption as follows.

- For every agent, the events of its signals are totally ordered, i.e., for all $1 \leq i \leq n$ and all $(t, x_i(t)), (t', x_i(t')) \in V_i$, if $t < t'$ then $(t, x_i(t)) \rightsquigarrow (t', x_i(t'))$.
- Every pair of events whose timestamps are at least ε apart is totally ordered, i.e., for all $1 \leq i, j \leq n$ and all $(t, x_i(t)) \in V_i$ and $(t', x_j(t')) \in V_j$, if $t + \varepsilon \leq t'$ then $(t, x_i(t)) \rightsquigarrow (t', x_j(t'))$.

Example 2. **TODO: distributed signal, happened-before relation**

Definition 3. Let (S, \rightsquigarrow) be a distributed signal of n signals, and $V = \bigcup_{i=1}^n V_i$ be the set of its events. A set $C \subseteq V$ is a consistent cut iff for every event in C , all events that happened before it also belong to C , i.e., for all $e, e' \in V$, if $e \in C$ and $e' \rightsquigarrow e$, then $e' \in C$.

We denote by $\mathbb{C}(T)$ the (infinite) set of consistent cuts at global time T . Given a consistent cut C , its *frontier* $\text{front}(C) \subseteq C$ is the set consisting of the last events in C of each signal, i.e., $\text{front}(C) = \bigcup_{i=1}^n \{(t, x_i(t)) \in V_i \cap C \mid \forall t' > t : (t', x_i(t')) \notin V_i \cap C\}$.

Definition 4. A consistent cut flow is a function $\text{ccf} : \mathbb{R}_{\geq 0} \rightarrow 2^V$ that maps a global clock value T to the frontier of a consistent cut at time T , i.e., $\text{ccf}(T) \in \{\text{front}(C) \mid C \in \mathbb{C}(T)\}$.

For all $T, T' \in \mathbb{R}_{\geq 0}$ and $1 \leq i \leq n$, if $T < T'$, then for every pair of events $(c_i(T), x_i(c_i(T))) \in \text{ccf}(T)$ and $(c_i(T'), x_i(c_i(T')))) \in \text{ccf}(T')$ we have $(c_i(T), x_i(c_i(T))) \rightsquigarrow (c_i(T'), x_i(c_i(T')))$. We denote by $\text{CCF}(S, \rightsquigarrow)$ the set of all consistent cut flows of the distributed signal (S, \rightsquigarrow) .

Example 5. TODO: consistent cut, frontier, consistent cut flow

Observe that a consistent cut flow of a distributed signal induces a vector of synchronous signals which can be evaluated using the standard semantics described in Section 2. Let (S, \rightsquigarrow) be a distributed signal of n signals x_1, \dots, x_n . A consistent cut flow $\text{ccf} \in \text{CCF}(S, \rightsquigarrow)$ yields a trace $w_{\text{ccf}} = (x'_1, \dots, x'_n)$ on a temporal domain $[0, D]$ where $D \in \mathbb{R}_{> 0}$ such that $(c_i(T), x_i(c_i(T))) \in \text{ccf}(T)$ implies $x'_i(T) = x_i(c_i(T))$ for all $1 \leq i \leq n$ and $T \in [0, D]$. The set of traces of (S, \rightsquigarrow) is given by $\text{Tr}(S, \rightsquigarrow) = \{w_{\text{ccf}} \mid \text{ccf} \in \text{CCF}(S, \rightsquigarrow)\}$.

We define the satisfaction of an STL formula φ by a distributed signal (S, \rightsquigarrow) over a three-valued domain $\{\top, \perp, ?\}$. If the set of synchronous traces $\text{Tr}(S, \rightsquigarrow)$ defined by a distributed signal (S, \rightsquigarrow) is contained in the set of traces allowed by the formula φ , then (S, \rightsquigarrow) satisfies φ . Similarly, if $\text{Tr}(S, \rightsquigarrow)$ has an empty intersection with the set of traces φ defines, then (S, \rightsquigarrow) violates φ . Otherwise, the evaluation is inconclusive since some traces satisfy the property and some violate it.

$$[(S, \rightsquigarrow) \models \varphi]_{\text{STL}} = \begin{cases} \top & \text{if } \forall w \in \text{Tr}(S, \rightsquigarrow) : [w \models \varphi]_{\text{STL}} \\ \perp & \text{if } \forall w \in \text{Tr}(S, \rightsquigarrow) : [w \models \neg \varphi]_{\text{STL}} \\ ? & \text{otherwise} \end{cases}$$

3 Overapproximation of the STL Distributed Semantics

To address the computational overhead in exact distributed monitoring, we define a new logic STL^+ whose syntax is the same as STL but semantics provide a sound approximation of the STL distributed semantics presented in Section 2. In essence, given a distributed signal (S, \rightsquigarrow) , STL^+ considers an overapproximation $\text{Tr}^+(S, \rightsquigarrow)$ of the set $\text{Tr}(S, \rightsquigarrow)$ of synchronous traces. A signal (S, \rightsquigarrow) satisfies (resp. violates) an STL^+ formula φ iff all the traces in $\text{Tr}^+(S, \rightsquigarrow)$ belong to the language of φ (resp. $\neg \varphi$).

$$[(S, \rightsquigarrow) \models \varphi]_{\text{STL}^+} = \begin{cases} \top & \text{if } \forall w \in \text{Tr}^+(S, \rightsquigarrow) : [w \models \varphi]_{\text{STL}} \\ \perp & \text{if } \forall w \in \text{Tr}^+(S, \rightsquigarrow) : [w \models \neg \varphi]_{\text{STL}} \\ ? & \text{otherwise} \end{cases}$$

In Sections 4 and 5, we respectively define Tr^+ and present an algorithm to compute the semantics of STL^+ . We finally prove the correctness of our approach.

Theorem 6. *For every STL formula φ and every distributed signal (S, \rightsquigarrow) , if $[(S, \rightsquigarrow) \models \varphi]_{\text{STL}^+} = \top$ (resp. \perp) then $[(S, \rightsquigarrow) \models \varphi]_{\text{STL}} = \top$ (resp. \perp).*

4 Overapproximation of Synchronous Traces

In this section, given a distributed signal (S, \rightsquigarrow) , we describe an overapproximation $\text{Tr}^+(S, \rightsquigarrow)$ of its set $\text{Tr}(S, \rightsquigarrow)$ of synchronous traces. First, we present the notion of *canonical segmentation*, a systematic way of partitioning the temporal domain of a given distributed signal to keep track of the partial asynchrony. Second, we introduce the notion of *value expressions*, sets of finite words representing how a signal behaves in a time interval. Finally, we define Tr^+ based on these notions, and show that it soundly approximates Tr .

Remark 7. We assume boolean signals in this section for convenience. The definitions and results presented here extend to real-valued signals.

Canonical Segmentation Consider a boolean signal x with a rising edge at time $t > \varepsilon$. Due to clock skew, this edge occurs in the range $(t - \varepsilon, t + \varepsilon)$ from the monitor's point of view. This range is called an *uncertainty region* because in $(t - \varepsilon, t + \varepsilon)$ the monitor cannot tell the value of x precisely, but only that it changes from 0 to 1. Formally, given an edge $(t, x(t))$, we define $\theta_{\text{lo}}(x, t) = \max(0, t - \varepsilon)$ and $\theta_{\text{hi}}(x, t) = t + \varepsilon$ as the end points of the edge's uncertainty region.

Given a temporal domain $I = [0, d) \subset \mathbb{R}_{\geq 0}$, a *segmentation* of I is a partition of I into finitely many intervals I_1, \dots, I_k , called *segments*, of the form $I_j = [t_j, t_{j+1})$ such that $t_j < t_{j+1}$ for all $1 \leq j \leq k$. By extension, a segmentation of a collection of signals with the same temporal domain I is a segmentation of I .

Let (S, \rightsquigarrow) be a distributed signal of n signals. The *canonical segmentation* G_S of (S, \rightsquigarrow) is the segmentation of S where the end points of the segments coincide with the end points of its temporal domain and uncertainty regions. Formally, we define G_S as follows. For each signal x_i , let F_i be the set of end points of its uncertainty regions. Let $d' = \max(d, \max(\bigcup_{i=1}^n F_i))$, which corresponds to the duration of the distributed signal with respect to the monitor's clock. Let $F = \{0, d'\} \cup \bigcup_{i=1}^n F_i$ and let $(s_j)_{1 \leq j \leq |F|}$ be a nondecreasing sequence of clock values corresponding to the elements of F . Then, the canonical segmentation of (S, \rightsquigarrow) is $G_S = \{I_1, \dots, I_{|F|-1}\}$ where $I_j = [s_j, s_{j+1})$ for all $1 \leq j < |F|$.

Example 8. Let (S, \rightsquigarrow) be a distributed boolean signal with $S = (x_1, x_2)$ and $\varepsilon = 2$ over the temporal domain $[0, 8)$ as given in Figure 1. Both signals are initially 0. The signal x_1 has a rising edge at time 2 and a falling edge at time 5, while x_2 has a rising edge at time 3 and a falling edge at time 6. The uncertainty regions of x_1 are $(0, 4)$ and $(3, 7)$, while those of x_2 are $(1, 5)$ and $(4, 8)$. Then,

we have $F = \{0, 8\} \cup \{0, 1, 3, 4, 5, 7, 8\}$, and thus the canonical segmentation is $G_S = \{[0, 1), [1, 3), [3, 4), [4, 5), [5, 7), [7, 8)\}$.




Fig. 1. The signals x_1 (top, red, dashed) and x_2 (bottom, blue, solid) from Example 8. The edges are marked with solid balls and their uncertainty regions are given as semi-transparent boxes around the edges. The resulting canonical segmentation G_S is shown below the graphical representation of the signals.

Value Expressions Consider a boolean signal x with a rising edge with an uncertainty region of (t_1, t_2) . As discussed above, the monitor only knows that the value of x changes from 0 to 1 in this interval. We represent this knowledge as a finite word $v = 0 \cdot 1$. This representation is called a *value expression* and it encodes the uncertain behavior of an observed signal relative to the monitor. Formally, a value expression is an element of Σ^* where Σ is the finite alphabet of values the signal takes. Given a signal x and an edge $(t, x(t))$, the value expression corresponding to the uncertainty region $(\theta_{lo}(x, t), \theta_{hi}(x, t))$ is given by $v_{x,t} = v_- \cdot v_+$ where $v_- = \lim_{s \rightarrow t^-} x(s)$ and $v_+ = \lim_{s \rightarrow t^+} x(s)$.

Notice that (i) uncertainty regions may overlap, and (ii) the canonical segmentation may split an uncertainty region into multiple segments. Consider a signal x with a rising edge in $(1, 5)$ and a falling edge in $(4, 8)$. The corresponding value expressions are respectively $v_1 = 0 \cdot 1$ and $v_2 = 1 \cdot 0$. Notice that the behavior of x in the interval $[1, 4)$ can be expressed as $\text{prefix}(v_1)$, encoding whether the rising edge has happened yet or not. Similarly, the behavior in $[4, 5)$ is given by $\text{suffix}(v_1) \cdot \text{prefix}(v_2)$, which captures whether the edges occur in this interval

(thanks to prefixing and suffixing) and the fact that the rising edge happens before the falling edge (thanks to concatenation).

Formally, given a distributed signal (S, \rightsquigarrow) , we define a function $\gamma : S \times G_S \rightarrow 2^{\Sigma^*}$ that maps each signal and segment of the canonical segmentation to a set of value expressions, capturing the signal's potential behaviors in the given segment. Let x be a signal in S , and let R_1, \dots, R_m be its uncertainty regions where $R_i = (t_i, t'_i)$ and the corresponding value expression is v_i for all $1 \leq i \leq m$. Now, let $I \in G_S$ be a segment with $I = [s, s']$ and for each $1 \leq i \leq m$ define the set V_i of value expressions capturing how I relates with R_i as follows:

$$V_i = \begin{cases} \{v_i\} & \text{if } t_i = s \wedge s' = t'_i \\ \text{prefix}(v_i) & \text{if } t_i = s \wedge s' < t'_i \\ \text{suffix}(v_i) & \text{if } t_i > s \wedge s' = t'_i \\ \text{infix}(v_i) & \text{if } t_i > s \wedge s' < t'_i \\ \{\epsilon\} & \text{otherwise} \end{cases}$$

Note that the last case happens only when $I \cap R_i$ is empty. We finally define $\gamma(x, I) = \text{destutter}(V_1 \cdot V_2 \cdot \dots \cdot V_m) \setminus \{\epsilon\}$. Observe that $\gamma(x, I)$ contains all the potential behaviors of x in segment I by construction. However, it is potentially overapproximate. This is mainly because the sets V_1, \dots, V_m contain redundancy by definition and the concatenation does not guarantee that an edge is considered exactly once.

Example 9. Recall the distributed signal (S, \rightsquigarrow) in Example 8 and Figure 1. In Figure 2a, we show the value expressions corresponding to its uncertainty regions. For example, the falling edge of x_1 has an uncertainty region of $(3, 7)$, represented by the value expression $1 \cdot 0$. In Figure 2b, we give the function γ for (S, \rightsquigarrow) . For example, $\gamma(x_1, [3, 4])$ is obtained from $\text{suffix}(0 \cdot 1) \cdot \text{prefix}(1 \cdot 0)$ and $\gamma(x_2, [0, 1]) = \{0\}$.

Overapproximation of Tr Consider a distributed signal (S, \rightsquigarrow) of n signals, and let G_S be its canonical segmentation. We describe how the function γ defines a set $\text{Tr}^+(S, \rightsquigarrow)$ of synchronous traces that overapproximates the set $\text{Tr}(S, \rightsquigarrow)$.

Let $x \in S$ and x' be two signals with the same temporal domain, and let $I = [s, s']$ be a segment in G_S . Let $(t_1, x'(t_1)), \dots, (t_\ell, x'(t_\ell))$ be the edges of x' in segment I with $t_i < t_{i+1}$ for all $1 \leq i < \ell$. The signals x and x' are *consistent in I* iff the value expression $x'(s) \cdot x'(t_1) \cdot \dots \cdot x'(t_\ell)$ belongs to $\gamma(x, I)$. Moreover, x and x' are *consistent* iff they are consistent in I for all $I \in G_S$. Now, let $S = (x_1, \dots, x_n)$ and define $\text{Tr}^+(S, \rightsquigarrow)$ as follows:

$$\text{Tr}^+(S, \rightsquigarrow) = \{(x'_1, \dots, x'_n) \mid x_i \text{ and } x'_i \text{ are consistent for all } 1 \leq i \leq n\}$$

Example 10. Recall the distributed signal (S, \rightsquigarrow) in Example 8 whose γ function is given in Figure 2b. Consider the synchronous trace $w \in \text{Tr}(S, \rightsquigarrow)$ where the rising edges of both signals occur at time 3 and the falling edges at time 5. One can verify that $w \in \text{Tr}^+(S, \rightsquigarrow)$ since for each $i \in \{1, 2\}$ the value expression 1 is




Fig. 2. (a) The uncertainty regions of the distributed signal in Example 8 and the corresponding value expressions. (b) The tabular representation of the function γ for the given distributed signal.

contained in $\gamma(x_i, [3, 4))$ and $\gamma(x_i, [4, 5))$ while 0 is contained in the remaining sets γ maps x_i to.

Now, consider a synchronous trace (x'_1, x'_2) where both signals are initially 0, have rising edges at time 2 and 3.5, and falling edges at time 3 and 5. Evidently, this trace does not belong to $\text{Tr}(S, \rightsquigarrow)$ since x'_1 and x'_2 have more edges than x_1 and x_2 . Nonetheless, it belongs to $\text{Tr}^+(S, \rightsquigarrow)$ since x'_1 and x'_2 are respectively consistent with x_1 and x_2 . To witness, notice that for each $i \in \{1, 2\}$ the value expression $0 \cdot 1$ is contained in $\gamma(x_i, [1, 3))$ and $\gamma(x_i, [3, 4))$, the expression 1 is contained in $\gamma(x_i, [4, 5))$, and 0 is contained in the remaining sets γ maps x_i to.

Finally, we show the correctness of the canonical overapproximation.

Lemma 11. *For every distributed signal (S, \rightsquigarrow) , we have $\text{Tr}(S, \rightsquigarrow) \subseteq \text{Tr}^+(S, \rightsquigarrow)$.*

5 Monitoring Algorithm

TODO

6 Experimental Evaluation

TODO

7 Conclusion

TODO

References

1. Maler, O., Nickovic, D.: Monitoring properties of analog and mixed-signal circuits. *Int. J. Softw. Tools Technol. Transf.* **15**(3), 247–268 (2013). <https://doi.org/10.1007/s10009-012-0247-9>