

Approximate Distributed Monitoring under Partial Synchrony: Balancing Speed with Accuracy

Author(s)

Affiliation(s)

Abstract. In distributed systems with processes that do not share a global clock, *partial synchrony* is achieved by clock synchronization that guarantees bounded clock skew among all applications. Existing solutions for distributed runtime verification under partial synchrony against temporal logic specifications are exact but suffer from significant computational overhead. In this paper, we propose an *approximate* distributed monitoring algorithm for Signal Temporal Logic (STL) that mitigates this issue by abstracting away potential interleaving behaviors. This conservative abstraction enables a significant speedup of the distributed monitors, albeit with a trade-off in accuracy. We address this trade-off with a methodology that combines our approximate monitor with its exact counterpart, resulting in enhanced monitoring efficiency without sacrificing precision. We validate our approach with multiple experiments, showcasing its effectiveness and efficacy on both a real-world application and synthetic examples.

1 Introduction

Distributed systems are networks of independent agents that work together to achieve a common objective. Distributed systems are everywhere around us and come in many different forms. For example, cloud computing uses distribution of resources and services over the internet to offer to their users a scalable infrastructure with transparent on-demand access to computing power and storage. Swarms of drones represent another family of distributed systems where individual drones collaborate to accomplish tasks like surveillance, search and rescue, or package delivery. While each drone operates independently, it also communicates and coordinates with others to successfully achieve their common objectives. The individual agents in a distributed system typically do not share a global clock. To coordinate actions across multiple agents, clock synchronization is often needed. While perfect clock synchronization is impractical due to network latency and node failures, algorithms such as the Network Time Protocol (NTP) allow agents to maintain a *bounded skew* between the synchronized clocks. In that case, we say that a distributed system has *partial synchrony*.

Formal verification of distributed system is a notoriously hard problem, due to the combinatorial explosion of all possible interleavings in the behaviors col-

lected from individual agents. *Runtime verification (RV)* provides a more pragmatic approach, in which a monitor observes a behavior of a distributed system and checks its correctness against a formal specification. The problem of distributed RV under partial synchrony assumption has been studied for Linear Temporal Logic (LTL) and Signal Temporal Logic (STL) specification languages. The proposed solutions use Satisfiability-Modulo-Theory (SMT) solving to provide sound and complete distributed monitoring procedures. Although distributed RV monitors consume only a single distributed behavior at a time, this behavior can nevertheless have an excessive number of possible interleavings. Hence, the exact distributed monitors from the literature can still suffer from significant computational overhead.

To mitigate this issue, we present an approach for *approximate* RV of STL specifications under partial synchrony. In essence, we abstract away potential interleavings in distributed behaviors in a conservative manner, resulting in an effective over-approximation of global behaviors. This abstraction simplifies the representation of distributed behaviors and the monitoring operations required to evaluate temporal specifications. There is an inevitable trade-off in approximate RV – gains in the monitoring speed-up may result in reduced accuracy. For some applications, reduced accuracy may not be acceptable. Therefore, we propose a methodology that combines our approximate monitors with their exact counterparts, with the aim to benefit from the enhanced monitoring efficiency without sacrificing precision. We implemented our approach in a prototype tool and performed thorough evaluations on both synthetic and real-world case studies. We first demonstrated that our approximate monitors achieve speed-ups of several orders of magnitudes compared to the exact SMT-based distributed RV solution. We empirically characterized the classes of specifications and behaviors for which our approximate monitoring approach achieves good precision. We finally showed that by combining exact and approximate distributed RV, there is still a significant efficiency gain on average without the sacrifice of the precision, even in cases where approximate monitors have low accuracy.

2 Preliminaries

We denote by $\mathbb{B} = \{\top, \perp\}$ the set of Booleans, \mathbb{R} the set of reals, $\mathbb{R}_{\geq 0}$ the set of nonnegative reals, and $\mathbb{R}_{> 0}$ the set of positive reals. An interval $I \subseteq \mathbb{R}$ of reals with the end points $a < b$ has length $|b - a|$.

Let Σ be a finite *alphabet*. We denote by Σ^* the set of finite words over Σ and by ϵ the empty word. For $u \in \Sigma^*$, we respectively write $\text{prefix}(u)$ and $\text{suffix}(u)$ for the sets of prefixes and suffixes of u . We also let $\text{infix}(u) = \{v \in \Sigma^* \mid \exists x, y \in \Sigma^* : u = xvy\}$. For a nonempty word $u \in \Sigma^*$ and $1 \leq i \leq |u|$, we denote by $u[i]$ the i th letter of u , by $u[..i]$ the prefix of u of length i , and by $u[i..]$ the suffix of u of length $|u| - i + 1$. Given $u \in \Sigma^*$ and $\ell \geq 1$, we denote by u^ℓ the word obtained by concatenating u by itself $\ell - 1$ times. Moreover, given $L \subseteq \Sigma^*$, we define $\text{first}(L) = \{u[0] \mid u \in L\}$. For sets $L_1, L_2 \subseteq \Sigma^*$ of words, we

81 let $L_1 \cdot L_2 = \{u \cdot v \mid u \in L_1, v \in L_2\}$. For tuples (u_1, \dots, u_m) and (v_1, \dots, v_m) of
 82 words, we let $(u_1, \dots, u_m) \cdot (v_1, \dots, v_m) = (u_1v_1, \dots, u_mv_m)$.

83 We define the function $\text{destutter} : \Sigma^* \rightarrow \Sigma^*$ inductively. For all $\sigma \in \Sigma \cup \{\epsilon\}$,
 84 let $\text{destutter}(\sigma) = \sigma$. For all $u \in \Sigma^*$ such that $u = \sigma_1\sigma_2v$ for some $\sigma_1, \sigma_2 \in \Sigma$
 85 and $v \in \Sigma^*$, we define it as follows.

$$\text{destutter}(u) = \begin{cases} \text{destutter}(\sigma_2v) & \text{if } \sigma_1 = \sigma_2 \\ \sigma_1 \cdot \text{destutter}(\sigma_2v) & \text{otherwise} \end{cases}$$

86 For a set $L \subseteq \Sigma^*$ of finite words, we define $\text{destutter}(L) = \{\text{destutter}(u) \mid u \in L\}$.
 87 To extend destutter to tuples of words, for all $\sigma \in \Sigma \cup \{\epsilon\}$ let $\text{destutter}(\sigma, \dots, \sigma) =$
 88 (σ, \dots, σ) . Given a tuple $(u_1, \dots, u_m) = (\sigma_{1,1}\sigma_{1,2}v_1, \dots, \sigma_{m,1}\sigma_{m,2}v_m)$ of finite
 89 words of the same length, we define $\text{destutter}(u_1, \dots, u_m)$ as expected:

$$\text{destutter}(u_1, \dots, u_m) = \begin{cases} \text{destutter}(\sigma_{1,2}v_1, \dots, \sigma_{m,2}v_m) & \text{if } \sigma_{i,1} = \sigma_{i,2} \text{ for all } 1 \leq i \leq m \\ (\sigma_{1,1}, \dots, \sigma_{m,1}) \cdot \text{destutter}(\sigma_{1,2}v_1, \dots, \sigma_{m,2}v_m) & \text{otherwise} \end{cases}$$

90 Moreover, given an integer $k \geq 0$, we define $\text{stutter}_k : \Sigma^* \rightarrow \Sigma^*$ such
 91 that $\text{stutter}_k(u) = \{v \in \Sigma^* \mid |v| = k \wedge \text{destutter}(v) = \text{destutter}(u)\}$ if $k \geq$
 92 $|\text{destutter}(u)|$, and $\text{stutter}_k(u) = \emptyset$ otherwise.

93 **Signal Temporal Logic (STL) [1].** Let $A, B \subset \mathbb{R}$. A function $f : A \rightarrow B$
 94 is *right-continuous* iff $\lim_{a \rightarrow c^+} f(a) = f(c)$ for all $c \in A$, and *non-Zeno* iff for
 95 every bounded interval $I \subseteq A$ there are finitely many $a \in I$ such that f is not
 96 continuous at a . A *signal* is a right-continuous, non-Zeno, piecewise-constant
 97 function $x : [0, d) \rightarrow \mathbb{R}$ where $d \in \mathbb{R}_{>0}$ is the duration of x and $[0, d)$ is its
 98 temporal domain. Let $x : [0, d) \rightarrow \mathbb{R}$ be a signal. An *event* of x is a pair $(t, x(t))$
 99 where $t \in [0, d)$. An *edge* of x is an event $(t, x(t))$ such that $\lim_{s \rightarrow t^-} x(s) \neq$
 100 $\lim_{s \rightarrow t^+} x(s)$. In particular, an edge is *rising* if $\lim_{s \rightarrow t^-} x(s) < \lim_{s \rightarrow t^+} x(s)$, and
 101 it is *falling* otherwise. A signal $x : [0, d) \rightarrow \mathbb{R}$ can be represented finitely by its
 102 initial value and edges: if x has m edges, then $x = (t_0, v_0)(t_1, v_1) \dots (t_m, v_m)$
 103 such that $t_0 = 0$, $t_{i-1} < t_i$, and (t_i, v_i) is an edge of x for all $1 \leq i \leq m$.

104 Let AP be a set of *atomic propositions*. The syntax is given by the following
 105 grammar where $p \in \text{AP}$ and $I \subseteq \mathbb{R}_{\geq 0}$ is an interval.

$$\varphi := p \mid \neg\varphi \mid \varphi \wedge \varphi \mid \varphi \mathcal{U}_I \varphi$$

106 A *trace* $w = (x_1, \dots, x_n)$ is a finite vector of signals. We express atomic
 107 propositions as functions of trace values at a time point t , i.e., a proposition
 108 $p \in \text{AP}$ over a trace $w = (x_1, \dots, x_n)$ is defined as $f_p(x_1(t), \dots, x_n(t)) > 0$
 109 where $f_p : \mathbb{R}^n \rightarrow \mathbb{R}$ is a function. Given intervals $I, J \subseteq \mathbb{R}_{\geq 0}$, we define $I \oplus J =$
 110 $\{i + j \mid i \in I \wedge j \in J\}$, and we simply write t for the singleton set $\{t\}$.

111 Below, we recall the finite-trace qualitative semantics of STL defined over \mathbb{B}
 112 [1]. Let $d \in \mathbb{R}_{>0}$ and $w = (x_1, \dots, x_n)$ with $x_i : [0, d) \rightarrow \mathbb{R}$ for all $1 \leq i \leq n$. Let
 113 φ_1, φ_2 be STL formulas and let $t \in [0, d)$.

$$\begin{aligned}
(w, t) \models p &\iff f_p(x_1(t), \dots, x_n(t)) > 0 \\
(w, t) \models \neg \varphi_1 &\iff \overline{(w, t) \models \varphi_1} \\
(w, t) \models \varphi_1 \wedge \varphi_2 &\iff (w, t) \models \varphi_1 \wedge (w, t) \models \varphi_2 \\
(w, t) \models \varphi_1 \mathcal{U}_I \varphi_2 &\iff \exists t' \in (t \oplus I) \cap [0, d) : \\
&\quad (w, t') \models \varphi_2 \wedge \forall t'' \in (t, t') : (w, t'') \models \varphi_1
\end{aligned}$$

114 We simply write $w \models \varphi$ for $(w, 0) \models \varphi$. We additionally use the following
115 standard abbreviations: **false** = $p \wedge \neg p$, **true** = $\neg \mathbf{false}$, $\varphi_1 \vee \varphi_2 = \neg(\neg \varphi_1 \wedge$
116 $\neg \varphi_2)$, $\Diamond_I \varphi = \mathbf{true} \mathcal{U}_I \varphi$, and $\Box_I \varphi = \neg \Diamond_I \neg \varphi$. Moreover, the untimed temporal
117 operators are defined through their timed counterparts on the interval $[0, \infty)$,
118 e.g., $\varphi_1 \mathcal{U} \varphi_2 = \varphi_1 \mathcal{U}_{[0, \infty)} \varphi_2$.

119 **Distributed Semantics of STL [2].** We consider an asynchronous and loosely-
120 coupled message-passing system of $n \geq 2$ reliable agents producing a set of
121 signals x_1, \dots, x_n , where for some $d \in \mathbb{R}_{>0}$ we have $x_i : [0, d) \rightarrow \mathbb{R}$ for all
122 $1 \leq i \leq n$. The agents do not share memory or a global clock. Only to formalize
123 statements, we speak of a *hypothetical* global clock and denote its value by T .
124 For local time values, we use the lowercase letters t and s .

125 For a signal x_i , we denote by V_i the set of its events, by E_i^\uparrow the set of its
126 rising edges, and by E_i^\downarrow that of falling edges. Moreover, we let $E_i = E_i^\uparrow \cup E_i^\downarrow$. We
127 represent the local clock of the i th agent as an increasing and divergent function
128 $c_i : \mathbb{R}_{\geq 0} \rightarrow \mathbb{R}_{\geq 0}$ that maps a global time T to a local time $c_i(T)$.

129 We assume that the system is *partially synchronous*: the agents use a clock
130 synchronization algorithm that guarantees a bounded clock skew with respect
131 to the global clock, i.e., $|c_i(T) - c_j(T)| < \varepsilon$ for all $1 \leq i, j \leq N$ and $T \in \mathbb{R}_{\geq 0}$,
132 where $\varepsilon \in \mathbb{R}_{>0}$ is the maximum clock skew.

133 **Definition 1.** A distributed signal is a pair (S, \rightsquigarrow) , where $S = (x_1, \dots, x_n)$ is a
134 vector of signals and \rightsquigarrow is the happened-before relation between events in signals
135 extended with the partial synchrony assumption as follows.

- 136 – For every agent, the events of its signals are totally ordered, i.e., for all $1 \leq$
137 $i \leq n$ and all $(t, x_i(t)), (t', x_i(t')) \in V_i$, if $t < t'$ then $(t, x_i(t)) \rightsquigarrow (t', x_i(t'))$.
- 138 – Every pair of events whose timestamps are at least ε apart is totally ordered,
139 i.e., for all $1 \leq i, j \leq n$ and all $(t, x_i(t)) \in V_i$ and $(t', x_j(t')) \in V_j$, if $t + \varepsilon \leq t'$
140 then $(t, x_i(t)) \rightsquigarrow (t', x_j(t'))$.

141 *Example 2.* **TODO: distributed signal, happened-before relation**

142 **Definition 3.** Let (S, \rightsquigarrow) be a distributed signal of n signals, and $V = \bigcup_{i=1}^n V_i$
143 be the set of its events. A set $C \subseteq V$ is a consistent cut iff for every event in
144 C , all events that happened before it also belong to C , i.e., for all $e, e' \in V$, if
145 $e \in C$ and $e' \rightsquigarrow e$, then $e' \in C$.

We denote by $\mathbb{C}(T)$ the (infinite) set of consistent cuts at global time T . Given a consistent cut C , its *frontier* $\text{front}(C) \subseteq C$ is the set consisting of the last events in C of each signal, i.e., $\text{front}(C) = \bigcup_{i=1}^n \{(t, x_i(t)) \in V_i \cap C \mid \forall t' > t : (t', x_i(t')) \notin V_i \cap C\}$.

Definition 4. A consistent cut flow is a function $\text{ccf} : \mathbb{R}_{\geq 0} \rightarrow 2^V$ that maps a global clock value T to the frontier of a consistent cut at time T , i.e., $\text{ccf}(T) \in \{\text{front}(C) \mid C \in \mathbb{C}(T)\}$.

For all $T, T' \in \mathbb{R}_{\geq 0}$ and $1 \leq i \leq n$, if $T < T'$, then for every pair of events $(c_i(T), x_i(c_i(T))) \in \text{ccf}(T)$ and $(c_i(T'), x_i(c_i(T')))) \in \text{ccf}(T')$ we have $(c_i(T), x_i(c_i(T))) \rightsquigarrow (c_i(T'), x_i(c_i(T')))$. We denote by $\text{CCF}(S, \rightsquigarrow)$ the set of all consistent cut flows of the distributed signal (S, \rightsquigarrow) .

Example 5. **TODO: consistent cut, frontier, consistent cut flow**

Observe that a consistent cut flow of a distributed signal induces a vector of synchronous signals which can be evaluated using the standard semantics described in Section 2. Let (S, \rightsquigarrow) be a distributed signal of n signals x_1, \dots, x_n . A consistent cut flow $\text{ccf} \in \text{CCF}(S, \rightsquigarrow)$ yields a trace $w_{\text{ccf}} = (x'_1, \dots, x'_n)$ on the temporal domain $[0, d]$ such that $(c_i(T), x_i(c_i(T))) \in \text{ccf}(T)$ implies $x'_i(T) = x_i(c_i(T))$ for all $1 \leq i \leq n$ and $T \in [0, d]$. The set of traces of (S, \rightsquigarrow) is given by $\text{Tr}(S, \rightsquigarrow) = \{w_{\text{ccf}} \mid \text{ccf} \in \text{CCF}(S, \rightsquigarrow)\}$.

We define the satisfaction of an STL formula φ by a distributed signal (S, \rightsquigarrow) over a three-valued domain $\{\top, \perp, ?\}$. If the set of synchronous traces $\text{Tr}(S, \rightsquigarrow)$ defined by a distributed signal (S, \rightsquigarrow) is contained in the set of traces allowed by the formula φ , then (S, \rightsquigarrow) satisfies φ . Similarly, if $\text{Tr}(S, \rightsquigarrow)$ has an empty intersection with the set of traces φ defines, then (S, \rightsquigarrow) violates φ . Otherwise, the evaluation is inconclusive since some traces satisfy the property and some violate it. Notice that we quantify universally over traces for both satisfaction and violation.

$$[(S, \rightsquigarrow) \models \varphi] = \begin{cases} \top & \text{if } \forall w \in \text{Tr}(S, \rightsquigarrow) : w \models \varphi \\ \perp & \text{if } \forall w \in \text{Tr}(S, \rightsquigarrow) : w \models \neg \varphi \\ ? & \text{otherwise} \end{cases}$$

3 Overapproximation of the STL Distributed Semantics

To address the computational overhead in exact distributed monitoring, we define STL^+ , a variant of STL whose syntax is the same as STL but semantics provide a sound approximation of the STL distributed semantics. STL^+ is better adapted to distributed monitoring as it overapproximates the set of traces. In particular, given a distributed signal (S, \rightsquigarrow) , STL^+ considers an approximation $\text{Tr}^+(S, \rightsquigarrow)$ of the set $\text{Tr}(S, \rightsquigarrow)$ of synchronous traces where $\text{Tr}(S, \rightsquigarrow) \subseteq \text{Tr}^+(S, \rightsquigarrow)$. A signal (S, \rightsquigarrow) satisfies (resp. violates) an STL^+ formula φ iff all the traces in $\text{Tr}^+(S, \rightsquigarrow)$ belong to the language of φ (resp. $\neg \varphi$).

$$[(S, \rightsquigarrow) \models \varphi]_+ = \begin{cases} \top & \text{if } \forall w \in \text{Tr}^+(S, \rightsquigarrow) : w \models \varphi \\ \perp & \text{if } \forall w \in \text{Tr}^+(S, \rightsquigarrow) : w \models \neg \varphi \\ ? & \text{otherwise} \end{cases}$$

In Sections 4 and 5, we respectively define Tr^+ and present an algorithm to compute the semantics of STL^+ . We finally prove the correctness of our approach.

Theorem 6. *For every STL formula φ and every distributed signal (S, \rightsquigarrow) , if $[(S, \rightsquigarrow) \models \varphi]_+ = \top$ (resp. \perp) then $[(S, \rightsquigarrow) \models \varphi] = \top$ (resp. \perp).*

Notice that both the distributed semantics of STL and the semantics of STL^+ quantify universally over the set of traces for the verdicts \top and \perp . Therefore, the Theorem 6 holds for the verdicts \top and \perp , but not for $?$.

4 Overapproximation of Synchronous Traces

In this section, given a distributed signal (S, \rightsquigarrow) , we describe an overapproximation $\text{Tr}^+(S, \rightsquigarrow)$ of its set $\text{Tr}(S, \rightsquigarrow)$ of synchronous traces. First, we present the notion of *canonical segmentation*, a systematic way of partitioning the temporal domain of a given distributed signal to keep track of the partial asynchrony. Second, we introduce the notion of *value expressions*, sets of finite words representing how a signal behaves in a time interval. Finally, we define Tr^+ based on these notions, and show that it soundly approximates Tr .

Remark 7. We assume Boolean signals in this section for convenience (all results hold for non-Boolean signals). The definitions and results presented here extend to real-valued signals because finite-length piecewise-constant signals will only use a finite number of values.

Canonical Segmentation Consider a Boolean signal x with a rising edge at time $t > \varepsilon$. Due to clock skew, this edge occurs in the range $(t - \varepsilon, t + \varepsilon)$ from the monitor's point of view. This range is called an *uncertainty region* because in $(t - \varepsilon, t + \varepsilon)$ the monitor cannot tell the value of x precisely, but only that it changes from 0 to 1. Formally, given an edge $(t, x(t))$, we define $\theta_{\text{lo}}(x, t) = \max(0, t - \varepsilon)$ and $\theta_{\text{hi}}(x, t) = \min(d, t + \varepsilon)$ as the end points of the edge's uncertainty region.

Given a temporal domain $I = [0, d) \subset \mathbb{R}_{\geq 0}$, a *segmentation* of I is a partition of I into finitely many intervals I_1, \dots, I_k , called *segments*, of the form $I_j = [t_j, t_{j+1})$ such that $t_j < t_{j+1}$ for all $1 \leq j \leq k$. By extension, a segmentation of a collection of signals with the same temporal domain I is a segmentation of I .

Let (S, \rightsquigarrow) be a distributed signal of n signals. The *canonical segmentation* G_S of (S, \rightsquigarrow) is the segmentation of S where the end points of the segments coincide with the end points of its temporal domain and uncertainty regions. Formally, we define G_S as follows. For each signal x_i , where $1 \leq i \leq n$, let F_i be

the set of end points of its uncertainty regions. Let $F = \{0, d\} \cup \bigcup_{i=1}^n F_i$ and let $(s_j)_{1 \leq j \leq |F|}$ be a nondecreasing sequence of clock values corresponding to the elements of F . Then, the canonical segmentation of (S, \rightsquigarrow) is $G_S = \{I_1, \dots, I_{|F|-1}\}$ where $I_j = [s_j, s_{j+1})$ for all $1 \leq j < |F|$.

Example 8. Let (S, \rightsquigarrow) be a distributed Boolean signal with $S = (x_1, x_2)$ and $\varepsilon = 2$ over the temporal domain $[0, 8)$ as given in Figure 1. Both signals are initially 0. The signal x_1 has a rising edge at time 2 and a falling edge at time 5, while x_2 has a rising edge at time 3 and a falling edge at time 6. The uncertainty regions of x_1 are $(0, 4)$ and $(3, 7)$, while those of x_2 are $(1, 5)$ and $(4, 8)$. Then, we have $F = \{0, 8\} \cup \{0, 1, 3, 4, 5, 7, 8\}$, and thus the canonical segmentation is $G_S = \{[0, 1), [1, 3), [3, 4), [4, 5), [5, 7), [7, 8)\}$.




Fig. 1. The signals x_1 (top, red, dashed) and x_2 (bottom, blue, solid) from Example 8. The edges are marked with solid balls and their uncertainty regions are given as semi-transparent boxes around the edges. The resulting canonical segmentation G_S is shown below the graphical representation of the signals.

Value Expressions Consider a boolean signal x with a rising edge with an uncertainty region of (t_1, t_2) . As discussed above, the monitor only knows that the value of x changes from 0 to 1 in this interval. We represent this knowledge as a finite word $v = 01$ over the alphabet $\Sigma = \{0, 1\}$. This representation is called a *value expression* and it encodes the uncertain behavior of an observed signal relative to the monitor. Formally, a value expression is an element of Σ^* where Σ is the finite alphabet of values the signal takes. Given a signal

235 x and an edge $(t, x(t))$, the value expression corresponding to the uncertainty
 236 region $(\theta_{\text{lo}}(x, t), \theta_{\text{hi}}(x, t))$ is given by $v_{x,t} = v_- \cdot v_+$ where $v_- = \lim_{s \rightarrow t^-} x(s)$
 237 and $v_+ = \lim_{s \rightarrow t^+} x(s)$. When drop the concatenation symbol \cdot when the letters
 238 are clear from the context. Let us remark that this definition is general because
 239 finite-length piecewise-constant real-valued signals will only have a finite number
 240 of values, making Σ finite.

241 Notice that (i) uncertainty regions may overlap, and (ii) the canonical seg-
 242 mentation may split an uncertainty region into multiple segments. Consider a
 243 signal x with a rising edge in $(1, 5)$ and a falling edge in $(4, 8)$. The corresponding
 244 value expressions are respectively $v_1 = 01$ and $v_2 = 10$. Notice that the behavior
 245 of x in the interval $[1, 4)$ can be expressed as $\text{prefix}(v_1)$, encoding whether the
 246 rising edge has happened yet or not. Similarly, the behavior in $[4, 5)$ is given
 247 by $\text{suffix}(v_1) \cdot \text{prefix}(v_2)$, which captures whether the edges occur in this interval
 248 (thanks to prefixing and suffixing) and the fact that the rising edge happens
 249 before the falling edge (thanks to concatenation).

250 Formally, given a distributed signal (S, \rightsquigarrow) , we define a function $\gamma : S \times$
 251 $G_S \rightarrow 2^{\Sigma^*}$ that maps each signal and segment of the canonical segmentation to
 252 a set of value expressions, capturing the signal's potential behaviors in the given
 253 segment. Let x be a signal in S , and let R_1, \dots, R_m be its uncertainty regions
 254 where $R_i = (t_i, t'_i)$ and the corresponding value expression is v_i for all $1 \leq i \leq m$.
 255 Now, let $I \in G_S$ be a segment with $I = [s, s')$ and for each $1 \leq i \leq m$ define the
 256 set V_i of value expressions capturing how I relates with R_i as follows:

$$V_i = \begin{cases} \{v_i\} & \text{if } t_i = s \wedge s' = t'_i \\ \text{prefix}(v_i) & \text{if } t_i = s \wedge s' < t'_i \\ \text{suffix}(v_i) & \text{if } t_i > s \wedge s' = t'_i \\ \text{infix}(v_i) & \text{if } t_i > s \wedge s' < t'_i \\ \{\epsilon\} & \text{otherwise} \end{cases} \quad (1)$$

257 The last case happens only when $I \cap R_i$ is empty. We finally define γ as follows:

$$\gamma(x, I) = \text{destutter}(V_1 \cdot V_2 \cdot \dots \cdot V_m) \setminus \{\epsilon\}$$

258 Observe that $\gamma(x, I)$ contains all the potential behaviors of x in segment I by
 259 construction. However, it is potentially overapproximate. This is mainly because
 260 the sets V_1, \dots, V_m contain redundancy by definition and the concatenation does
 261 not guarantee that an edge is considered exactly once.

262 *Example 9.* Recall the distributed signal (S, \rightsquigarrow) in Example 8 and Figure 1.
 263 In Figure 2a, we show the value expressions corresponding to its uncertainty
 264 regions. For example, the falling edge of x_1 has an uncertainty region of $(3, 7)$,
 265 represented by the value expression 10. In Figure 2b, we give the function γ for
 266 (S, \rightsquigarrow) . For example, $\gamma(x_1, [3, 4)) = (\text{suffix}(01) \cdot \text{prefix}(10)) \setminus \{\epsilon\} = \{01, 010, 1, 10\}$,
 267 and $\gamma(x_2, [0, 1)) = \{0\}$.

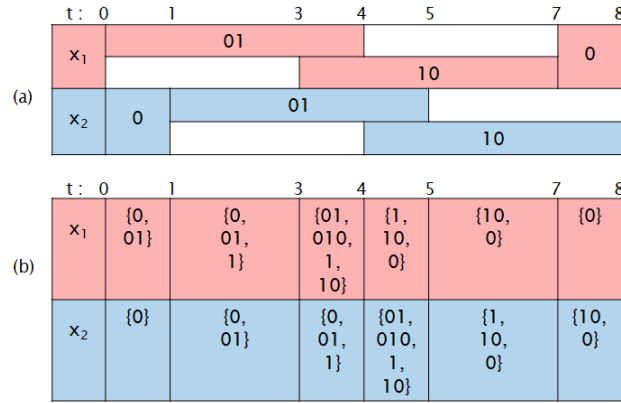


Fig. 2. (a) The uncertainty regions of the distributed signal in Example 8 and the corresponding value expressions. (b) The tabular representation of the function γ for the given distributed signal.

Overapproximation of Tr. Consider a distributed signal (S, \rightsquigarrow) of n signals, and let G_S be its canonical segmentation. We describe how the function γ defines a set $\text{Tr}^+(S, \rightsquigarrow)$ of synchronous traces that overapproximates the set $\text{Tr}(S, \rightsquigarrow)$. Let $x \in S$ and x' be two signals with the same temporal domain, and let $I = [s, s']$ be a segment in G_S . Let $(t_1, x'(t_1)), \dots, (t_\ell, x'(t_\ell))$ be the edges of x' in segment I with $t_i < t_{i+1}$ for all $1 \leq i < \ell$. The signal x' is *I-consistent with x* iff the value expression $x'(s) \cdot x'(t_1) \cdot \dots \cdot x'(t_\ell)$ belongs to $\gamma(x, I)$. Moreover, x' is *consistent with x* iff it is *I-consistent with x* for all $I \in G_S$.

Now, let $S = (x_1, \dots, x_n)$ and define $\text{Tr}^+(S, \rightsquigarrow)$ as follows:

$$\text{Tr}^+(S, \rightsquigarrow) = \{(x'_1, \dots, x'_n) \mid x'_i \text{ is consistent with } x_i \text{ for all } 1 \leq i \leq n\}$$

Example 10. Recall the distributed signal (S, \rightsquigarrow) in Example 8 whose γ function is given in Figure 2b. Consider the synchronous trace $w \in \text{Tr}(S, \rightsquigarrow)$ where the rising edges of both signals occur at time 3 and the falling edges at time 5. A signal such as w would be included in $\text{Tr}^+(S, \rightsquigarrow)$ since for each $i \in \{1, 2\}$ the value expression 1 is contained in $\gamma(x_i, [3, 4))$ and $\gamma(x_i, [4, 5))$ while 0 is contained in the remaining sets γ maps x_i to. Now, consider a synchronous trace (x'_1, x'_2) where both signals are initially 0, have rising edges at time 2 and 3.5, and falling edges at time 3 and 5. Evidently, this trace does not belong to $\text{Tr}(S, \rightsquigarrow)$ since x'_1 and x'_2 have more edges than x_1 and x_2 . Nonetheless, it belongs to $\text{Tr}^+(S, \rightsquigarrow)$ since x'_1 and x'_2 are respectively consistent with x_1 and x_2 . To witness, notice that for each $i \in \{1, 2\}$ the value expression 01 is contained in $\gamma(x_i, [1, 3))$ and $\gamma(x_i, [3, 4))$, the expression 1 is contained in $\gamma(x_i, [4, 5))$, and 0 is contained in the remaining sets γ maps x_i to.

Finally, we prove that Tr^+ overapproximates Tr .

Lemma 11. For every distributed signal (S, \rightsquigarrow) , we have $\text{Tr}(S, \rightsquigarrow) \subseteq \text{Tr}^+(S, \rightsquigarrow)$.

5 Monitoring Algorithm

In this section, given a distributed signal (S, \rightsquigarrow) , we describe an algorithm to compute $[(S, \rightsquigarrow) \models \varphi]_+$. The algorithm makes use of the function γ defined in Section 4 without explicitly computing $\text{Tr}^+(S, \rightsquigarrow)$. To achieve this, we first describe the notion of *asynchronous product* of value expressions to capture potential interleavings within segments. We continue with the evaluation of *untimed operators* and then *timed operators*. Finally, we conclude with putting all these together to compute the *semantics* of STL^+ and discuss an efficient implementation of the monitoring algorithm using *bit vectors* to represent and manipulate sets of Boolean value expressions.

Remark 12. For the sake of convenience, we focus on Boolean signals for the rest of the section. Note that asynchronous products and the algorithm to compute $[(S, \rightsquigarrow) \models \varphi]_+$ can be extended to value expressions over arbitrary finite alphabets, e.g., encoding finite-length piecewise-constant real-valued signals. This allows us to express more complex properties where atomic propositions can be functions of real-valued signals.

Asynchronous Products Consider the value expressions $u_1 = 0 \cdot 1$ and $u_2 = 1 \cdot 0$ encoding the behaviors of two signals within a segment. The behaviors within a segment are seen as completely asynchronous. To capture the potential interleavings of these behaviors, we consider how the values in u_1 and u_2 can align. In particular, there are three potential alignments: (i) the rising edge of u_1 happens before the falling edge of u_2 , (ii) the falling edge of u_2 happens before the rising edge of u_1 , and (iii) the two edges happen simultaneously. We respectively represent these with the tuples $(011, 110)$, $(001, 100)$, and $(01, 10)$ where the first component encodes u_1 and the second u_2 . Formally, given two value expressions u_1 and u_2 , we define their *asynchronous product* as follows:

$$u_1 \otimes u_2 = \{\text{destutter}(v_1, v_2) \mid v_i \in \text{stutter}_k(u_i), k = |u_1| + |u_2| - 1, i \in \{1, 2\}\}$$

Moreover, given two sets L_1 and L_2 of value expressions, we define the following:

$$L_1 \otimes L_2 = \{u_1 \otimes u_2 \mid u_1 \in L_1, u_2 \in L_2\}$$

Asynchronous products of value expressions allow us to lift value expressions to satisfaction signals of formulas.

Example 13. Recall the distributed signal (S, \rightsquigarrow) in Example 8 and its γ function given in Figure 2b. Suppose we want to compute the value expressions encoding the satisfaction of $x_1 \wedge x_2$ in the segment $[1, 3)$. We can achieve this by first computing the asynchronous product $\gamma(x_1, [3, 4)) \otimes \gamma(x_2, [3, 4))$, and then computing the bitwise conjunction of each pair in the set. For example, considering the expression $0 \cdot 1 \cdot 0$ for x_1 and $0 \cdot 1$ for x_2 , the product contains the pair $(010, 011)$. Taking the bitwise conjunction of this pair gives us the expression $0 \cdot 1 \cdot 0$ as a potential behavior for the satisfaction of $x_1 \wedge x_2$ in this segment.

Untimed Operations As hinted in Example 13, to compute the semantics, we apply bitwise operations on value expressions and their asynchronous products to transform them into encodings of satisfaction signals of formulas. Consider the distributed signal (S, \rightsquigarrow) in Example 8 and suppose we want to compute $[(S, \rightsquigarrow) \models \Diamond(x_1 \wedge x_2)]_+$. To achieve this, we first compute for each segment in G_S the set of value expressions for the satisfaction of $x_1 \wedge x_2$, and then from these compute that of $\Diamond(x_1 \wedge x_2)$. This compositional approach allows us to evaluate arbitrary STL⁺ formulas.

First, we define bitwise operations on boolean value expressions encoding atomic propositions. Then, we use these to evaluate (untimed) STL formulas over sets of value expressions.

Let u and v be boolean value expressions of length ℓ . We denote by $u \& v$ the bitwise-and operation, by $u \mid v$ the bitwise-or, and by $\sim u$ the bitwise-negation. In addition, we define the *bitwise strong until* operator as follows:

$$u \mathbf{U}^0 v = \left(\max_{i \leq j \leq \ell} \left(\min \left(v[j], \min_{i \leq k \leq j} u[k] \right) \right) \right)_{1 \leq i \leq \ell}$$

As usual, we derive *bitwise eventually* as $\mathbf{E}u = 1^\ell \mathbf{U}^0 u$, *bitwise always* as $\mathbf{A}u = \sim(\mathbf{E}\sim u)$, and *bitwise weak until* as $u \mathbf{U}^1 v = (u \mathbf{U}^0 v) \mid (\mathbf{A}u)$. The distinction between \mathbf{U}^0 and \mathbf{U}^1 will be useful later when we evaluate a formula segment by segment. We remark that the definitions of these operators coincide with the robustness semantics of (discrete time) STL. Finally, note that the output of these operations is a value expression of length ℓ . For example, if $u = 010$, we have $\mathbf{E}u = 110$ and $\mathbf{A}u = 000$.

Let (S, \rightsquigarrow) be a distributed signal. Consider an atomic proposition $p \in \mathbf{AP}$ encoded as $x_p \in S$ and let φ_1, φ_2 be two STL formulas. We define the evaluation of untimed formulas with respect to (S, \rightsquigarrow) and a segment $I \in G_S$ inductively:

$$\begin{aligned} \llbracket (S, \rightsquigarrow), I \models p \rrbracket &= \gamma(x_p, I) \\ \llbracket (S, \rightsquigarrow), I \models \neg \varphi_1 \rrbracket &= \{\sim u \mid u \in \llbracket (S, \rightsquigarrow), I \models \varphi_1 \rrbracket\} \\ \llbracket (S, \rightsquigarrow), I \models \varphi_1 \wedge \varphi_2 \rrbracket &= \{u_1 \& u_2 \mid (u_1, u_2) \in \llbracket (S, \rightsquigarrow), I \models \varphi_1 \rrbracket \otimes \llbracket (S, \rightsquigarrow), I \models \varphi_2 \rrbracket\} \\ \llbracket (S, \rightsquigarrow), I \models \varphi_1 \mathbf{U}^a \varphi_2 \rrbracket &= \{u_1 \mathbf{U}^a u_2 \mid (u_1, u_2) \in \llbracket (S, \rightsquigarrow), I \models \varphi_1 \rrbracket \otimes \llbracket (S, \rightsquigarrow), I \models \varphi_2 \rrbracket, \\ &\quad a \in \mathbf{first}(\llbracket (S, \rightsquigarrow), I' \models \varphi_1 \mathbf{U} \varphi_2 \rrbracket)\} \end{aligned}$$

where I' is the segment that follows I in G_S , if it exists. For completeness, for every formula φ we define $\llbracket (S, \rightsquigarrow), I' \models \varphi \rrbracket = \{0\}$ when $I' \notin G_S$. When I is the first segment in G_S , we simply write $\llbracket (S, \rightsquigarrow) \models \varphi \rrbracket$. Similarly as above, we can use the standard derived operators to compute the corresponding sets of value expressions. Intuitively, for a given formula and a segment, the evaluation above produces a set of value expressions encoding the formula's satisfaction within the segment.

Example 14. Recall the distributed signal (S, \rightsquigarrow) in Example 8 and its γ function given in Figure 2b. Suppose we want to compute $\llbracket (S, \rightsquigarrow), [5, 7) \models \Diamond(x_1 \wedge x_2) \rrbracket$.

First, we compute $\llbracket (S, \rightsquigarrow), [5, 7) \models x_1 \wedge x_2 \rrbracket$ by computing the bitwise conjunction over the asynchronous product $\gamma(x_1, [5, 7)) \otimes \gamma(x_2, [5, 7))$ and destuttering. For example, since $010 \in \gamma(x_1, [5, 7))$ and $01 \in \gamma(x_2, [5, 7))$, the pair $(0010, 0111)$ is in the product, whose conjunction gives us 010 after destuttering. Repeating this for the rest, we obtain $\llbracket (S, \rightsquigarrow), [5, 7) \models x_1 \wedge x_2 \rrbracket = \{0, 01, 010, 1, 10\}$. Finally, we compute $\llbracket (S, \rightsquigarrow), [5, 7) \models \Diamond(x_1 \wedge x_2) \rrbracket$ by applying each expression in $\llbracket (S, \rightsquigarrow), [5, 7) \models x_1 \wedge x_2 \rrbracket$ the bitwise eventually operator and destuttering. The resulting set $\{0, 1, 10\}$ encodes the satisfaction signal of $\Diamond(x_1 \wedge x_2)$ in $[5, 7)$. Note that we do not need to consider the evaluation of the next segment for the eventually operator since $\llbracket (S, \rightsquigarrow), [7, 8) \models x_1 \wedge x_2 \rrbracket = \{0\}$.

Timed Operations Handling timed operations requires a closer inspection as value expressions are untimed by definition. We address this issue by considering how a given evaluation interval relates with a given segmentation. For example, take a segmentation $G_S = \{[0, 4), [4, 6), [6, 10)\}$ and an evaluation interval $J = [0, 5)$. Suppose we are interested in how a signal $x \in S$ behaves with respect to J over the first segment $I = [0, 4)$. First, to see how J relates with G_S with respect to $I = [0, 4)$, we “slide” the interval J over $I \oplus J = [0, 9)$ and consider the different ways it intersects the segments in G_S . Initially, J covers the entire segment $[0, 4)$ and the beginning of $[4, 6)$, for which the potential behaviors of x are captured by the set $\gamma(x, [0, 4)) \cdot \text{prefix}(\gamma(x, [4, 6)))$. Now, if we slide the window and take $J' = [3, 7)$, the window covers the ending of $[0, 4)$, the entire $[4, 6)$, and the beginning of $[6, 10)$, for which the potential behaviors are captured by the set $\text{suffix}(\gamma(x, [0, 4))) \cdot \gamma(x, [4, 6)) \cdot \text{prefix}(\gamma(x, [6, 9)))$. We call these sets the *profiles* of J and J' with respect to (S, \rightsquigarrow) , x , and I .

Let (S, \rightsquigarrow) be a distributed signal, $I \in G_S$ be a segment, and φ be an STL formula. Let us introduce the notation we use in the definition below. First, we abbreviate the set $\llbracket (S, \rightsquigarrow), I \models \varphi \rrbracket$ of value expressions as $\tau_{\varphi, I}$. Second, given an interval K , we respectively denote by l_K and r_K its left and right end points. Third, recall that we denote by F the set of end points of G_S (see Section 4). Given an interval J , we define the *profile* of J with respect to (S, \rightsquigarrow) , φ , and I as follows.

$$\text{profile}((S, \rightsquigarrow), \varphi, I, J) = \begin{cases} \text{prefix}(\tau_{\varphi, I}) & \text{if } l_I = l_J \wedge r_I > r_J \\ \text{infix}(\tau_{\varphi, I}) & \text{if } l_I < l_J \wedge r_I > r_J \\ \tau_{\varphi, I} \cdot \kappa(\varphi, I, J) & \text{if } l_I = l_J \wedge r_I \leq r_J \wedge r_J \in F \setminus J \\ \tau_{\varphi, I} \cdot \kappa(\varphi, I, J) \cdot \text{first}(\tau_{\varphi, I'}) & \text{if } l_I = l_J \wedge r_I \leq r_J \wedge r_J \in F \cap J \\ \tau_{\varphi, I} \cdot \kappa(\varphi, I, J) \cdot \text{prefix}(\tau_{\varphi, I'}) & \text{if } l_I = l_J \wedge r_I \leq r_J \wedge r_J \notin F \\ \text{suffix}(\tau_{\varphi, I}) \cdot \kappa(\varphi, I, J) & \text{if } l_I < l_J < r_I \leq r_J \wedge r_J \in F \setminus J \\ \text{suffix}(\tau_{\varphi, I}) \cdot \kappa(\varphi, I, J) \cdot \text{first}(\tau_{\varphi, I'}) & \text{if } l_I < l_J < r_I \leq r_J \wedge r_J \in F \cap J \\ \text{suffix}(\tau_{\varphi, I}) \cdot \kappa(\varphi, I, J) \cdot \text{prefix}(\tau_{\varphi, I'}) & \text{if } l_I < l_J < r_I \leq r_J \wedge r_J \notin F \\ \{\epsilon\} & \text{otherwise} \end{cases}$$

where we assume J is trimmed to fit the temporal domain of S and $I' \in G_S$ is such that $r_J \in I'$. Moreover, $\kappa(\varphi, I, J)$ is the concatenation $\tau_{\varphi, I_1} \cdot \dots \cdot \tau_{\varphi, I_m}$ such that I, I_1, \dots, I_m, I' are consecutive segments in G_S . If I_1, \dots, I_m do not exist, we let $\kappa(\varphi, I, J) = \{\epsilon\}$. Note that the last case happens when $I \cap J$ is empty.

397 We now formalize the intuitive approach of “sliding” J over the segmentation to
 398 obtain the various profiles it produces as follows.

$$\mathbf{pfs}((S, \rightsquigarrow), \varphi, I, J) = \{\text{destutter}(\text{profile}((S, \rightsquigarrow), \varphi, I, J')) \mid J' \subseteq I \oplus J, J' \sim J\}$$

399 where $J' \sim J$ holds when $|J'| = |J|$ and J' contains an end point (left or
 400 right) iff J does so. Note that although infinitely many intervals J' satisfy the
 401 conditions given above (due to denseness of time), the set defined by \mathbf{pfs} is finite.
 402 We demonstrate this and the computation of \mathbf{pfs} in Example 15 and Figure 3.

403 *Example 15.* Recall the distributed signal (S, \rightsquigarrow) in Example 8 and its γ function
 404 given in Figure 2b. We demonstrate the computation of $\mathbf{pfs}((S, \rightsquigarrow), x_1, [1, 3], [0, 1])$.
 405 Intuitively, sliding the interval $[0, 1)$ over the window $[1, 3] \oplus [0, 1)$ (as shown in
 406 Figure 3) gives us the following sets:

$$\begin{aligned} P_1 &= \text{destutter}(\text{prefix}(\gamma(x_1, [1, 3]))) = \{0, 01, 1\} \\ P_2 &= \text{destutter}(\text{infix}(\gamma(x_1, [1, 3]))) = \{0, 01, 1\} \\ P_3 &= \text{destutter}(\text{suffix}(\gamma(x_1, [1, 3]))) = \{0, 01, 1\} \\ P_4 &= \text{destutter}(\text{suffix}(\gamma(x_1, [1, 3])) \cdot \text{prefix}(\gamma(x_1, [3, 4]))) \\ &= \{0, 01, 010, 0101, 01010, 1, 10, 101, 1010\} \end{aligned}$$

407 Therefore, we obtain $\mathbf{pfs}((S, \rightsquigarrow), x_1, [1, 3], [0, 1]) = \{P_1, P_2, P_3, P_4\}$. This set over-
 408 approximates the potential behaviors of x_1 , for all $t \in [1, 3)$, in the interval
 409 $t \oplus [0, 1)$.




Fig. 3. The profiles of $J = [0, 1)$ with respect to $x_1 \in S$ of Example 8. The four representative intervals of each profile is shown with solid black lines below the tabular representation of γ for x_1 .

410 Let φ_1 and φ_2 be two STL formulas. Intuitively, once we have the profiles of
 411 a given interval J with respect to φ_1 and φ_2 , we can evaluate the correspond-
 412 ing untimed formulas on the product of these profiles and concatenate them.
 413 Formally, we handle the evaluation of timed formulas inductively as follows.

$$\llbracket (S, \rightsquigarrow), I \models \varphi_1 \mathcal{U}_J \varphi_2 \rrbracket = \text{destutter}(\{u_1 \mathcal{U}^0 u_2 \mid (u_1, u_2) \in P_1 \otimes Q_1\} \cdot \dots \cdot \{u_1 \mathcal{U}^0 u_2 \mid (u_1, u_2) \in P_k \otimes Q_k\})$$

where $\text{pfs}((S, \rightsquigarrow), \varphi_1, I, J) = \{P_1, \dots, P_k\}$ and $\text{pfs}((S, \rightsquigarrow), \varphi_2, I, J) = \{Q_1, \dots, Q_k\}$ such that the intervals producing P_i and Q_i respectively start before those producing P_{i+1} and Q_{i+1} for all $1 \leq i < k$.

Example 16. Let (S, \rightsquigarrow) be as in Example 8 and its γ function as given in Figure 2b. We demonstrate the evaluation of the timed formula $\Diamond_{[0,1)} x_1$ over the segment $[1, 3)$. Recall from Example 15 the set $\text{pfs}((S, \rightsquigarrow), x_1, [1, 3), [0, 1)) = \{P_1, P_2, P_3, P_4\}$ of profiles. First, we apply the bitwise eventually operator to each value expression in each of these profiles separately: $\{Eu \mid u \in P_1\} = \{0, 1\}$, $\{Eu \mid u \in P_2\} = \{0, 1\}$, $\{Eu \mid u \in P_3\} = \{0, 1\}$, and $\{Eu \mid u \in P_4\} = \{0, 10, 1\}$. Then, we concatenate these sets and destutter to obtain the following:

$$\llbracket (S, \rightsquigarrow), [1, 3) \models \Diamond_{[0,1)} x_1 \rrbracket = \{0, 01, 010, 0101, 01010, 1, 10, 101, 1010\}$$

Computing the Semantics of STL⁺ Putting it all together, given a distributed signal (S, \rightsquigarrow) and an STL⁺ formula φ , we can compute $\llbracket (S, \rightsquigarrow) \models \varphi \rrbracket_+$ thanks to the following theorem.

Theorem 17. *For every distributed signal (S, \rightsquigarrow) , we have $\llbracket (S, \rightsquigarrow) \models \varphi \rrbracket_+ = \top$ (resp. \perp , $?$) iff $\text{first}(\llbracket (S, \rightsquigarrow) \models \varphi \rrbracket) = \{1\}$ (resp. $\{0\}$, $\{0, 1\}$).*

Sets of Boolean Value Expressions as Bit Vectors Evidently, asynchronous products are expensive to compute. Our implementation of the algorithm we describe in this section relies on the following observation: Sets of boolean value expressions and their operations can be efficiently implemented through bit vectors. Intuitively, to represent such a set, we can encode each element using its first bit and its length since value expressions are boolean and always destuttered. Moreover, to evaluate untimed operations on such sets, we only need to know the maximal lengths of the four possible types of expressions ($0 \dots 0$, $0 \dots 1$, $1 \dots 0$, and $1 \dots 1$) and whether the set contains 0 or 1 (to handle some edge cases). This is because value expressions corresponding to same segments can be seen as completely asynchronous and the possible interleavings obtained from shorter expressions can be obtained from longer ones. This approach enables, for example, an algorithm for conjunction of sets of value expressions that runs in $O(|u| + |v|)$ time where u and v are the longest expressions in the two sets. Note that the same idea also applies to untimed temporal operators.

6 Experimental Evaluation

This section presents the research questions and the experiments that we conducted and analysed to validate our approximate distributed monitoring approach.

448 6.1 Research Questions

449 We seek answers to the following research questions (RQs):

450 **RQ1.** *What is the tradeoff between the efficiency and the accuracy of approximate*
 451 *distributed monitors?* The approximate distributed monitoring comes with a
 452 price in terms of the loss of accuracy. We want to understand what is the tradeoff
 453 between the potential speed-ups that an approximate distributed monitor can
 454 potentially achieve when compared to its exact counterpart and the consequent
 455 loss in accuracy due to the approximations. We would also like to identify the
 456 classes of signals and properties for which this tradeoff is effective.

457 **RQ2.** *Can the combination of approximate and exact distributed monitors in-*
 458 *crease efficiency while preserving accuracy?* We are interested in evaluating
 459 whether a smart combined use of approximate and exact distributed monitors
 460 can still bring improvements in monitoring efficiency while guaranteeing the ac-
 461 curacy of the monitoring verdicts.

462 6.2 Experimental Setup

463 **Distributed Monitors** In our study, we compare our approximate distributed
 464 monitoring (ADM) approach to an exact distributed monitoring approach. For
 465 the exact monitoring, we take a variant of the distributed monitoring procedure
 466 from [?] that allows to evaluate STL specifications over distributed traces using
 467 SMT-solving. Originally, that procedure assumes that input signals are polynomial
 468 continuous functions. We adapt the SMT-based approach to consider input
 469 signals as piecewise-constant signals to make a consistent comparison with ADM.
 470 We note that the passage from the polynomial continuous to piecewise-constant
 471 input signals reduces the efficiency of the SMT-based monitors. We also observe
 472 that the SMT-based monitors from [?] can split the input trace into multiple
 473 segments and evaluate the specification incrementally, segment-by-segment, al-
 474 lowing early termination of the monitor in some cases. Since the focus of this
 475 paper is purely on the offline monitoring, we also use the exact monitors without
 476 their incremental mode. We use the abbreviation EDM to denote the variant of
 477 the exact SMT-based distributed monitors used in this study.

478 **Experimental Subjects** To answer our research questions, we also use a *ran-*
 479 *dom generator* (RG) of distributed traces, in which we can control the trace
 480 duration d (in terms of the number of events) and the bound on the clock skew
 481 ϵ . We also consider two real-world, publicly accessible applications, a *heat pump*
 482 (HP) and a *swarm of drones* (SD). In addition, to have more quantitative and
 483 statistically relevant results,

484 *Heat pump (HP) models...*

485 *Swarm of drones (SD) models...*

486 Specifications

Subject	Spec ID	STL formula	Comments
RG	φ_1	$\Box(p \wedge q)$	untimed
	φ_2	$\Diamond(p \vee q)$	
	φ_3	$\Box(p \vee q)$	
	φ_4	$\Diamond(p \wedge q)$	
	φ_5	$p\mathcal{U}q$	
HP	φ_{HP}	$\Diamond(p \vee q)$	arithmetic + numeric predicates
SD	φ_{SD}	$\Diamond(p \vee q)$	

Table 1. STL specifications used in the experiments.

487 **Computing Platform** To conduct the experiments, we used...

488 7 Conclusion

489 **TODO**

490 References

- 491 1. Maler, O., Nickovic, D.: Monitoring properties of analog and mixed-signal circuits.
492 Int. J. Softw. Tools Technol. Transf. **15**(3), 247–268 (2013). <https://doi.org/10.1007/s10009-012-0247-9>
- 493 2. Momtaz, A., Abbas, H., Bonakdarpour, B.: Monitoring signal temporal logic in
494 distributed cyber-physical systems. In: Mitra, S., Venkatasubramanian, N., Dubey,
495 A., Feng, L., Ghasemi, M., Sprinkle, J. (eds.) Proceedings of the ACM/IEEE 14th
496 International Conference on Cyber-Physical Systems, ICCPS 2023, (with CPS-IoT
497 Week 2023), San Antonio, TX, USA, May 9-12, 2023. pp. 154–165. ACM (2023).
498 <https://doi.org/10.1145/3576841.3585937>
- 499

500 Appendix

501 Proof of Theorem 6.

502 *Proof.* Let φ be an STL formula and (S, \rightsquigarrow) be a distributed signal. Assume
 503 $[(S, \rightsquigarrow) \models \varphi]_+ = \top$. We want to show that $[(S, \rightsquigarrow) \models \varphi] = \top$. Expanding
 504 the definition of $[(S, \rightsquigarrow) \models \varphi]_+ = \top$, we have $w \models \varphi$ for all $w \in \text{Tr}^+(S, \rightsquigarrow)$.
 505 By Lemma 11, we have $\text{Tr}(S, \rightsquigarrow) \subseteq \text{Tr}^+(S, \rightsquigarrow)$. Then, it holds that $w \models \varphi$
 506 for all $w \in \text{Tr}(S, \rightsquigarrow)$. Therefore, $[(S, \rightsquigarrow) \models \varphi] = \top$ by definition. The case of
 507 $[(S, \rightsquigarrow) \models \varphi]_+ = \perp$ follows from the same arguments.

508 Proof of Lemma 11.

509 *Proof.* Let (S, \rightsquigarrow) be a distributed signal where $S = (x_1, \dots, x_n)$. Let $w =$
 510 $(y_1, \dots, y_n) \in \text{Tr}(S, \rightsquigarrow)$ be a trace. We want to show that $w \in \text{Tr}^+(S, \rightsquigarrow)$. First,
 511 let us recall the definition of Tr^+ .

$$\text{Tr}^+(S, \rightsquigarrow) = \{(x'_1, \dots, x'_n) \mid x'_i \text{ is consistent with } x_i \text{ for all } 1 \leq i \leq n\}$$

512 Let $1 \leq i \leq n$ be arbitrary. To show that y_i is consistent with x_i , we need to
 513 show that y_i is I -consistent with x_i for all $I \in G_S$. Let $I = [t_0, s)$ be an arbitrary
 514 segment in G_S , let $(t_1, y_i(t_1)), \dots, (t_\ell, y_i(t_\ell))$ be the edges of y_i in segment I with
 515 $t_j < t_{j+1}$ for all $1 \leq j < \ell$. To show that y_i is I -consistent with x_i , we need to
 516 show that the expression $y_i(t_0) \cdot y_i(t_1) \cdot \dots \cdot y_i(t_\ell)$ belongs to $\gamma(x_i, I)$. We sketch
 517 the proof idea below.

518 Note that w can be seen as a trace obtained through an ε -retiming of S (see
 519 [2, Section 4.2]). Then, the timestamp t of any edge of x_i is mapped to some
 520 clock value in the range $(\theta_{\text{lo}}(t), \theta_{\text{hi}}(t))$. In particular, $|t - c_i^{-1}(t)| < \varepsilon$ for all
 521 $t \in \{t_0, t_1, \dots, t_\ell\}$, where $c_i^{-1}(t)$ is the local clock value of x_i that is mapped to
 522 t .

523 Since y_i has ℓ edges in I , it holds that x_i has at least ℓ edges in $(t_0 - \varepsilon, s +$
 524 $\varepsilon)$. Since I is a segment in G_S , there are ℓ of these that are consecutive such
 525 that the intersection of their uncertainty regions contain (t_0, s) , i.e., $(t_0, s) \subseteq$
 526 $\bigcap_{1 \leq j \leq \ell} (\theta_{\text{lo}}(t'_j), \theta_{\text{hi}}(t'_j))$ where $t'_j = c_i^{-1}(t_j)$ is the corresponding timestamp in x_i
 527 for all $0 \leq j \leq \ell$. In particular, note that $y_i(t_j) = x_i(t'_j)$ for all $0 \leq j \leq \ell$.

528 Now, notice that, by definition, $\gamma(x_i, I)$ takes into account every edge of x_i
 529 whose uncertainty region has a nonempty intersection with I , and preserves their
 530 order. Let V_j be the set of value expressions capturing how I relates with the
 531 uncertainty intervals of the edge $(t'_j, x_i(t'_j))$ for all $1 \leq j \leq \ell$ (as defined in
 532 Equation (1)). Then, $\text{destutter}(\{x_i(t'_0)\} \cdot V_1 \cdot \dots \cdot V_\ell) \subseteq \gamma(x_i, I)$. One can verify
 533 that for all $1 \leq j \leq \ell$, either $x_i(t'_j)$ or $x_i(t'_{j-1}) \cdot x_i(t'_j)$ belongs to V_j . This allows
 534 us to choose a value expression v_j from each V_j such that $\text{destutter}(\{x_i(t'_0)\} \cdot v_1 \cdot$
 535 $\dots \cdot v_\ell) = x_i(t'_0) \cdot x_i(t'_1) \cdot \dots \cdot x_i(t'_\ell)$, which concludes the proof.

536 Note that if there are more edges of x_i with a timestamp smaller than t'_0 or
 537 larger than t'_ℓ whose uncertainty intervals intersect with I , then the correspond-
 538 ing set of value expressions is obtained either by prefixing or suffixing. In either
 539 case, we can choose ϵ from these sets for concatenation with the remaining edges'
 540 value expressions and obtain the desired result.

541 **Proof of Theorem 17. TODO: Fix**

542 Let $w = (x_1, \dots, x_n)$ be a (synchronous) trace and $E_w = \{t_1, \dots, t_m\}$ be the
 543 set of timestamps for its edges with $t_j < t_{j+1}$ for all $1 \leq j < m$. We define
 544 $\pi(w) = (v_1, \dots, v_n)$ where $v_i = x_i(0) \cdot x_i(t_1) \cdot \dots \cdot x_i(t_m)$ for all $1 \leq i \leq n$. Let
 545 (S, \rightsquigarrow) be a distributed signal. We define

$$\pi(\text{Tr}^+(S, \rightsquigarrow)) = \{\pi(w) \mid w \in \text{Tr}^+(S, \rightsquigarrow)\}.$$

546 Note that for each $u \in \pi(\text{Tr}^+(S, \rightsquigarrow))$ we have $u = \text{destutter}(u)$.

547 Let $k \in \mathbb{N}$ and recall the function $\text{stutter}_k : \Sigma^* \rightarrow \Sigma^*$. For $L \subseteq \Sigma^*$, we define
 548 $\text{stutter}_k(L) = \bigcup_{u \in L} \text{stutter}_k(u)$.

549 Let (S, \rightsquigarrow) be a distributed signal such that $S = (x_1, \dots, x_n)$ and $G_S =$
 550 $\{I_1, \dots, I_k\}$ where the segments with lower indices start (and end) before those
 551 with higher indices. Moreover, we define

$$\Omega(S, \rightsquigarrow) = \{\text{destutter}(u) \mid u \in \Gamma(S, \rightsquigarrow)\}$$

552 where

$$\Gamma(S, \rightsquigarrow) = (\text{stutter}_3(\gamma(x_1, I_1)) \dots \text{stutter}_3(\gamma(x_1, I_k)), \dots, \text{stutter}_3(\gamma(x_n, I_1)) \dots \text{stutter}_3(\gamma(x_n, I_k))).$$

553 Note that stutter_3 only serves to capture the interleavings caused by the asyn-
 554 chronicity within the segments.

555 **Lemma 18.** *For every distributed signal (S, \rightsquigarrow) , we have $\pi(\text{Tr}^+(S, \rightsquigarrow)) = \Omega(S, \rightsquigarrow)$.*

556 *Proof.* We briefly sketch each direction. The key observation is that both sets of
 557 value expressions are essentially defined through γ .

558 Let $u \in \pi(\text{Tr}^+(S, \rightsquigarrow))$. By definition, u is obtained from a trace (x'_1, \dots, x'_n)
 559 such that, for each $1 \leq i \leq n$, the signal x'_i is I -consistent with x_i for all $I \in G_S$.
 560 Namely, for each $1 \leq i \leq n$ and $I \in G_S$, the value expression obtained by
 561 concatenating the initial value of x'_i in I and the values of its edges in I belongs
 562 to $\gamma(x_i, I)$. Then, one can verify that $u \in \Omega(S, \rightsquigarrow)$.

563 Let $u \in \Omega(S, \rightsquigarrow)$. Recall that u is obtained by choosing, for each $1 \leq i \leq$
 564 n and $I \in G_S$, a value expression $v_{i,I} \in \gamma(x_i, I)$. By definition, there exists
 565 $(x'_1, \dots, x'_n) \in \text{Tr}^+(S, \rightsquigarrow)$ such that, for each $1 \leq i \leq n$ and $I \in G_S$, the behavior
 566 of x'_i in the interval I is given by the value expression $v_{i,I}$. Then, one can verify
 567 that $u \in \pi(\text{Tr}^+(S, \rightsquigarrow))$.

568 We now proceed to prove Theorem 17.

569 *Proof (Theorem 17).* The proof goes by induction on the structure of φ . Let
 570 (S, \rightsquigarrow) be a distributed signal and φ an STL formula. We sketch each case
 571 below.

572 **Atomic propositions.** For the base case, let $\varphi = p$ for some $p \in \text{AP}$. Note that
 573 $\llbracket (S, \rightsquigarrow), I \models p \rrbracket$ expands to $\gamma(x_p, I)$ where x_p is the signal encoding p . Then, it is
 574 easy to see that for every $u \in \Omega(S, \rightsquigarrow)$ the component corresponding to x_p starts
 575 with the value $\text{first}(\gamma(x_p, I))$ (which is a singleton). Moreover, by Lemma 18, the

576 same holds for $\pi(\text{Tr}^+(S, \rightsquigarrow))$, thus we conclude $[(S, \rightsquigarrow) \models p]_+ = \top$ (resp. \perp) iff
 577 $\text{first}(\llbracket (S, \rightsquigarrow) \models p \rrbracket) = \{1\}$ (resp. $\{0\}$).

578 **Negation.** Let $\varphi = \neg\psi$ for some formula ψ such that $[(S, \rightsquigarrow) \models \psi]_+$ and
 579 $\text{first}(\llbracket (S, \rightsquigarrow) \models \psi \rrbracket)$ agree. If $[(S, \rightsquigarrow) \models \psi]_+ = \perp$ (resp. \top , $?$) and $\text{first}(\llbracket (S, \rightsquigarrow) \models$
 580 $\psi \rrbracket) = \{0\}$ (resp. $\{1\}$, $\{0, 1\}$), then, thanks to their inductive definitions, it is easy
 581 to see that $[(S, \rightsquigarrow) \models \neg\psi]_+ = \top$ (resp. \perp , $?$) and $\text{first}(\llbracket (S, \rightsquigarrow) \models \neg\psi \rrbracket) = \{1\}$
 582 (resp. $\{0\}$, $\{1, 0\}$).

583 **Conjunction.** Let $\varphi = \psi_1 \wedge \psi_2$ for some formulas ψ_1, ψ_2 such that $[(S, \rightsquigarrow) \models$
 584 $\psi_i]_+$ and $\text{first}(\llbracket (S, \rightsquigarrow) \models \psi_i \rrbracket)$ agree for each $i \in \{1, 2\}$. Suppose $[(S, \rightsquigarrow) \models \psi_i]_+ =$
 585 \top and $\text{first}(\llbracket (S, \rightsquigarrow) \models \psi_i \rrbracket) = \{1\}$ agree for each $i \in \{1, 2\}$. Then, $[(S, \rightsquigarrow) \models$
 586 $\psi_1 \wedge \psi_2]_+ = \top$ because $\text{tr}^+(S, \rightsquigarrow)$ being contained in the languages of ψ_1 and
 587 ψ_2 implies that it is also contained in their intersection, and $\text{first}(\llbracket (S, \rightsquigarrow) \models$
 588 $\psi_i \rrbracket) = \{1\}$ because the conjunction of any asynchronous product where both
 589 components start with 1 also starts with 1. The remaining cases follow from
 590 similar arguments.

591 **Untimed until.** Let $\varphi = \psi_1 \wedge \psi_2$ for some formulas ψ_1, ψ_2 such that $[(S, \rightsquigarrow) \models$
 592 $\psi_i]_+$ and $\text{first}(\llbracket (S, \rightsquigarrow) \models \psi_i \rrbracket)$ agree for each $i \in \{1, 2\}$.

593 **Timed until.** TODO

Generalization to Real-Valued Signals The methods described in Section 4 and here can be generalized to handle arithmetic operations over real-valued signals and numerical predicates. The key observation enabling this is that such finite-length piecewise-constant signals take only finitely many values. Then, letting Σ be a finite alphabet of these values, we are able to compute the γ function as described. To handle arithmetic operations, we can compute the asynchronous product of the signals involved and apply the given operation letter-by-letter. Finally, to transform them into atomic propositions, we can compare the resulting value expressions letter-by-letter with a constant. We call this approach ORIG. For example, suppose the asynchronous product of two signals x_1 and x_2 contains $(2 \cdot 2 \cdot 3, 1 \cdot 0 \cdot 1)$. Applying addition to this pair letter by letter would give us $3 \cdot 2 \cdot 4$, and applying the comparison > 2 would result in the Boolean value expression 101. Repeating this for all pairs, we obtain the atomic proposition required.

As described in the previous subsection, we can efficiently represent and manipulate Boolean value expressions, avoiding explicit computation of asynchronous products. This method does not apply to real-valued signals in general, therefore formulas with numerical predicates cause an additional overhead. However, there are several ways how we can avoid explicit computation of asynchronous products for some classes of formulas and numerical predicates.

Since signals are completely asynchronous within segments, we can ignore the order and compute a set of potential *values* instead of a set of *sequences of values*. Continuing the example above, instead of repeating the process above for all the interleavings, we can consider the value sets $X_1 = \{2, 3\}$ for x_1 and $X_2 = \{0, 1\}$ for x_2 to conclude by pairwise addition that the values of $x_1 + x_2$ come from the set $\{2, 3, 4\}$. Then, we can simply assume that $x_1 + x_2$ is constant with one of these values within the segment. We call this approach FINE. Notice that ignoring the order may lead to additionally overapproximating the original set of traces. Therefore, FINE does not apply to formulas where the order of values matters, e.g., $x_1 > c_1 \mathcal{U} x_2 > c_2$. Nonetheless, when it applies, e.g. for $\Box((x_1 - x_2) \times x_3 > c)$, it saves us the time to explicitly compute the interleavings and the value expressions.

We can take FINE further and only consider the extreme values. Some arithmetic operations are *monotonic* in the sense that the extreme values of the output's value expressions can be computed from those of inputs. For example, since $\min X_1 = 2$, $\max X_1 = 3$, $\min X_2 = 0$, and $\max X_2 = 1$, we immediately obtain that the minimal and maximal values of $x_1 + x_2$ are 2 and 4. We call this approach COARSE.

Finally, instead of assuming that the monitor runs on a separate agent, we can assume it runs on one of the monitored agents. This helps reduce the computational overhead caused by asynchrony by taking the agent where the monitor runs as a reference point, removing any uncertainty associated with this agent's signal. We call this approach RELATIVE. We will demonstrate the benefits and downsides of each of these in Section 6.