

# Approximate Distributed Monitoring under Partial Synchrony: Balancing Speed with Accuracy

Author(s)

Affiliation(s)

**Abstract.** In distributed systems with processes that do not share a global clock, *partial synchrony* is achieved by clock synchronization that guarantees bounded clock skew among all applications. Existing solutions for distributed runtime verification under partial Synchrony against temporal logic specifications are exact but suffer from significant computational overhead. In this paper, we propose an *approximate* distributed monitoring algorithm for Signal Temporal Logic (STL) that mitigates this issue by abstracting away potential interleaving behaviors. This conservative abstraction enables a significant speedup of the distributed monitors, albeit with a trade-off in accuracy. We address this trade-off with a methodology that combines our approximate monitor with its exact counterpart, resulting in enhanced monitoring efficiency without sacrificing precision. We validate our approach with multiple experiments, showcasing its effectiveness and efficacy on both a real-world application and synthetic examples.

You can leave notes using the command `\firstName{note}`.  
Page limit: 16 + 4 (appendix)

## 1 Introduction

*Distributed systems* are networks of independent agents that work together to achieve a common objective. Distributed systems are everywhere around us and come in many different forms. For example, cloud computing uses distribution of resources and services over the internet to offer to their users a scalable infrastructure with transparent on-demand access to computing power and storage. Swarms of drones represent another family of distributed systems where individual drones collaborate to accomplish tasks like surveillance, search and rescue, or package delivery. While each drone operates independently, it also communicates and coordinates with others to successfully achieve their common objectives. The individual agents in a distributed system typically do not share a global clock. To coordinate actions across multiple agents, clock synchronization is often needed. While perfect clock synchronization is impractical due to network latency and node failures, algorithms such as the Network Time Protocol (NTP) allow agents

36 to maintain a *bounded skew* between the synchronized clocks. In that case, we  
 37 say that a distributed system has *partial synchrony*.

38 Formal verification of distributed system is a notoriously hard problem, due  
 39 to the combinatorial explosion of all possible interleavings in the behaviors col-  
 40 lected from individual agents. *Runtime verification (RV)* provides a more prag-  
 41 matic approach, in which a monitor observes a behavior of a distributed sys-  
 42 tem and checks its correctness against a formal specification. The problem of  
 43 distributed RV under partial synchrony assumption has been studied for Lin-  
 44 ear Temporal Logic (LTL) and Signal Temporal Logic (STL) specification lan-  
 45 guages. The proposed solutions use Satisfiability-Modulo-Theory (SMT) solving  
 46 to provide sound and complete distributed monitoring procedures. Although dis-  
 47 tributed RV monitors consume only a single distributed behavior at a time, this  
 48 behavior can nevertheless have an excessive number of possible interleavings.  
 49 Hence, the exact distributed monitors from the literature can still suffer from  
 50 significant computational overhead.

51 To mitigate this issue, we present an approach for *approximate* RV of STL  
 52 specifications under partial synchrony. In essence, we abstract away potential  
 53 interleavings in distributed behaviors in a conservative manner, resulting in an  
 54 effective over-approximation of global behaviors. This abstraction simplifies the  
 55 representation of distributed behaviors and the monitoring operations required  
 56 to evaluate temporal specifications. There is an inevitable trade-off in approxi-  
 57 mate RV – gains in the monitoring speed-up may result in reduced accuracy. For  
 58 some applications, reduced accuracy may not be acceptable. Therefore, we pro-  
 59 pose a methodology that combines our approximate monitors with their exact  
 60 counterparts, with the aim to benefit from the enhanced monitoring efficiency  
 61 without sacrificing precision. We implemented our approach in a prototype tool  
 62 and performed thorough evaluations on both synthetic and real-world case stud-  
 63 ies. We first demonstrated that our approximate monitors achieve speed-ups of  
 64 several orders of magnitudes compared to the exact SMT-based distributed RV  
 65 solution. We empirically characterized the classes of specifications and behaviors  
 66 for which our approximate monitoring approach achieves good precision. We fi-  
 67 nally showed that by combining exact and approximate distributed RV, there is  
 68 still a significant efficiency gain on average without the sacrifice of the precision,  
 69 even in cases where approximate monitors have low accuracy.

## 70 2 Preliminaries

71 We denote by  $\mathbb{B} = \{\top, \perp\}$  the set of Booleans,  $\mathbb{R}$  the set of reals,  $\mathbb{R}_{\geq 0}$  the set of  
 72 nonnegative reals, and  $\mathbb{R}_{> 0}$  the set of positive reals. An interval  $I \subseteq \mathbb{R}$  of reals  
 73 with the end points  $a < b$  has length  $|b - a|$ .

74 Let  $\Sigma$  be a finite *alphabet*. We denote by  $\Sigma^*$  the set of finite words over  
 75  $\Sigma$  and by  $\epsilon$  the empty word. For  $u \in \Sigma^*$ , we respectively write  $\text{prefix}(u)$  and  
 76  $\text{suffix}(u)$  for the sets of prefixes and suffixes of  $u$ . We also let  $\text{infix}(u) = \{v \in$   
 77  $\Sigma^* \mid \exists x, y \in \Sigma^* : u = xvy\}$ . For a nonempty word  $u \in \Sigma^*$  and  $1 \leq i \leq |u|$ ,  
 78 we denote by  $u[i]$  the  $i$ th letter of  $u$ , by  $u[..i]$  the prefix of  $u$  of length  $i$ , and by

79  $u[i..]$  the suffix of  $u$  of length  $|u| - i + 1$ . Given  $u \in \Sigma^*$  and  $\ell \geq 1$ , we denote by  
 80  $u^\ell$  the word obtained by concatenating  $u$  by itself  $\ell - 1$  times. Moreover, given  
 81  $L \subseteq \Sigma^*$ , we define  $\text{first}(L) = \{u[0] \mid u \in L\}$ . For sets  $L_1, L_2 \in \Sigma^*$  of words, we  
 82 let  $L_1 \cdot L_2 = \{u \cdot v \mid u \in L_1, v \in L_2\}$ . For tuples  $(u_1, \dots, u_m)$  and  $(v_1, \dots, v_m)$  of  
 83 words, we let  $(u_1, \dots, u_m) \cdot (v_1, \dots, v_m) = (u_1v_1, \dots, u_mv_m)$ .

84 We define the function  $\text{destutter} : \Sigma^* \rightarrow \Sigma^*$  inductively as follows. For all  
 85  $\sigma \in \Sigma \cup \{\epsilon\}$ , let  $\text{destutter}(\sigma) = \sigma$ . For all  $u \in \Sigma^*$  such that  $u = \sigma_1\sigma_2v$  for  
 86 some  $\sigma_1, \sigma_2 \in \Sigma$  and  $v \in \Sigma^*$ , let (i)  $\text{destutter}(u) = \text{destutter}(\sigma_2v)$  if  $\sigma_1 = \sigma_2$ ,  
 87 and (ii)  $\text{destutter}(u) = \sigma_1 \cdot \text{destutter}(\sigma_2v)$  otherwise. By extension, for a set  
 88  $L \subseteq \Sigma^*$  of finite words, we write  $\text{destutter}(L) = \{\text{destutter}(u) \mid u \in L\}$ . Given  
 89 a tuple  $(u_1, \dots, u_m) = (\sigma_{1,1}\sigma_{1,2}v_1, \dots, \sigma_{m,1}\sigma_{m,2}v_m)$  of finite words of the same  
 90 length, we define  $\text{destutter}(u_1, \dots, u_m)$  as expected: (i)  $\text{destutter}(u_1, \dots, u_m) =$   
 91  $\text{destutter}(\sigma_{1,2}v_1, \dots, \sigma_{m,2}v_m)$  if  $\sigma_{i,1} = \sigma_{i,2}$  for all  $1 \leq i \leq m$ , and (ii)  $\text{destutter}(u_1, \dots, u_m) =$   
 92  $(\sigma_{1,1}, \dots, \sigma_{m,1}) \cdot \text{destutter}(\sigma_{1,2}v_1, \dots, \sigma_{m,2}v_m)$  otherwise.

93 Moreover, given an integer  $k \geq 0$ , we define  $\text{stutter}_k : \Sigma^* \rightarrow \Sigma^*$  such  
 94 that  $\text{stutter}_k(u) = \{v \in \Sigma^* \mid |v| = k \wedge \text{destutter}(v) = \text{destutter}(u)\}$  if  $k \geq$   
 95  $|\text{destutter}(u)|$ , and  $\text{stutter}_k(u) = \emptyset$  otherwise.

96 **Signal Temporal Logic (STL)** Let  $A, B \subset \mathbb{R}$ . A function  $f : A \rightarrow B$  is  
 97 *right-continuous* iff  $\lim_{a \rightarrow c^+} f(a) = f(c)$  for all  $c \in A$ , and *non-Zeno* iff for  
 98 every bounded interval  $I \subseteq A$  there are finitely many  $a \in I$  such that  $f$  is not  
 99 continuous at  $a$ . A *signal* is a right-continuous, non-Zeno, piecewise-constant  
 100 function  $x : [0, d) \rightarrow \mathbb{R}$  where  $d \in \mathbb{R}_{>0}$  is the duration of  $x$  and  $[0, d)$  is its  
 101 temporal domain. Let  $x : [0, d) \rightarrow \mathbb{R}$  be a signal. An *event* of  $x$  is a pair  $(t, x(t))$   
 102 where  $t \in [0, d)$ . An *edge* of  $x$  is an event  $(t, x(t))$  such that  $\lim_{s \rightarrow t^-} x(s) \neq$   
 103  $\lim_{s \rightarrow t^+} x(s)$ . In particular, an edge is *rising* if  $\lim_{s \rightarrow t^-} x(s) < \lim_{s \rightarrow t^+} x(s)$ , and  
 104 it is *falling* otherwise. A signal  $x : [0, d) \rightarrow \mathbb{R}$  can be represented finitely by its  
 105 initial value and edges: if  $x$  has  $m$  edges, then  $x = (t_0, v_0)(t_1, v_1) \dots (t_m, v_m)$   
 106 such that  $t_0 = 0$ ,  $t_{i-1} < t_i$ , and  $(t_i, v_i)$  is an edge of  $x$  for all  $1 \leq i \leq m$ .

107 Let  $\text{AP}$  be a set of atomic propositions. The syntax is given by the following  
 108 grammar where  $p \in \text{AP}$  and  $I \subseteq \mathbb{R}_{\geq 0}$  is an interval.

$$\varphi := p \mid \neg\varphi \mid \varphi \wedge \varphi \mid \varphi \mathcal{U}_I \varphi$$

109 A *trace*  $w = (x_1, \dots, x_n)$  is a finite vector of signals. We express atomic  
 110 propositions as functions of trace values at a time point  $t$ , i.e., a proposition  
 111  $p \in \text{AP}$  over a trace  $w = (x_1, \dots, x_n)$  is defined as  $f_p(x_1(t), \dots, x_n(t)) > 0$   
 112 where  $f_p : \mathbb{R}^n \rightarrow \mathbb{R}$  is a function. Given intervals  $I, J \subseteq \mathbb{R}_{\geq 0}$ , we define  $I \oplus J =$   
 113  $\{i + j \mid i \in I \wedge j \in J\}$ , and we simply write  $t$  for the singleton set  $\{t\}$ .

114 Below we recall the finite-trace qualitative semantics of STL defined over  $\mathbb{B}$   
 115  $[1]$ . Let  $d \in \mathbb{R}_{>0}$  and  $w = (x_1, \dots, x_n)$  with  $x_i : [0, d) \rightarrow \mathbb{R}$  for all  $1 \leq i \leq n$ . Let  
 116  $\varphi_1, \varphi_2$  be STL formulas and let  $t \in [0, d)$ .

$$\begin{aligned}
(w, t) \models p &\iff f_p(x_1(t), \dots, x_n(t)) > 0 \\
(w, t) \models \neg \varphi_1 &\iff \overline{(w, t) \models \varphi_1} \\
(w, t) \models \varphi_1 \wedge \varphi_2 &\iff (w, t) \models \varphi_1 \wedge (w, t) \models \varphi_2 \\
(w, t) \models \varphi_1 \mathcal{U}_I \varphi_2 &\iff \exists t' \in (t \oplus I) \cap [0, d) : \\
&\quad (w, t') \models \varphi_2 \wedge \forall t'' \in (t, t') : (w, t'') \models \varphi_1
\end{aligned}$$

117 We simply write  $w \models \varphi$  for  $(w, 0) \models \varphi$ . We additionally use the following  
118 standard abbreviations: **false** =  $p \wedge \neg p$ , **true** =  $\neg \mathbf{false}$ ,  $\varphi_1 \vee \varphi_2 = \neg(\neg \varphi_1 \wedge$   
119  $\neg \varphi_2)$ ,  $\Diamond_I \varphi = \mathbf{true} \mathcal{U}_I \varphi$ , and  $\Box_I \varphi = \neg \Diamond_I \neg \varphi$ . Moreover, the untimed temporal  
120 operators are defined through their timed counterparts on the interval  $(0, \infty)$ ,  
121 e.g.,  $\varphi_1 \mathcal{U} \varphi_2 = \varphi_1 \mathcal{U}_{(0, \infty)} \varphi_2$ .

122 **Distributed Semantics of STL** We consider an asynchronous and loosely-  
123 coupled message-passing system of  $n \geq 2$  reliable agents producing a set of  
124 signals  $x_1, \dots, x_n$ , where for some  $d \in \mathbb{R}_{>0}$  we have  $x_i : [0, d) \rightarrow \mathbb{R}$  for all  
125  $1 \leq i \leq n$ . The agents do not share memory or a global clock. Only to formalize  
126 statements, we speak of a *hypothetical* global clock and denote its value by  $T$ .  
127 For local time values, we use the lowercase letters  $t$  and  $s$ .

128 For a signal  $x_i$ , we denote by  $V_i$  the set of its events, by  $E_i^\uparrow$  the set of its  
129 rising edges, and by  $E_i^\downarrow$  that of falling edges. Moreover, we let  $E_i = E_i^\uparrow \cup E_i^\downarrow$ . We  
130 represent the local clock of the  $i$ th agent as an increasing and divergent function  
131  $c_i : \mathbb{R}_{\geq 0} \rightarrow \mathbb{R}_{\geq 0}$  that maps a global time  $T$  to a local time  $c_i(T)$ . We denote by  
132  $c_i^{-1}$  the inverse of the local clock function  $c_i$ .

133 We assume that the system is *partially synchronous*: the agents use a clock  
134 synchronization algorithm that guarantees a bounded clock skew with respect  
135 to the global clock, i.e.,  $|c_i(T) - c_j(T)| < \varepsilon$  for all  $1 \leq i, j \leq N$  and  $T \in \mathbb{R}_{\geq 0}$ ,  
136 where  $\varepsilon \in \mathbb{R}_{>0}$  is the maximum clock skew.

137 **Definition 1.** A distributed signal is a pair  $(S, \rightsquigarrow)$ , where  $S = (x_1, \dots, x_n)$  is a  
138 vector of signals and  $\rightsquigarrow$  is the happened-before relation between events in signals  
139 extended with the partial synchrony assumption as follows.

- 140 – For every agent, the events of its signals are totally ordered, i.e., for all  $1 \leq$   
141  $i \leq n$  and all  $(t, x_i(t)), (t', x_i(t')) \in V_i$ , if  $t < t'$  then  $(t, x_i(t)) \rightsquigarrow (t', x_i(t'))$ .
- 142 – Every pair of events whose timestamps are at least  $\varepsilon$  apart is totally ordered,  
143 i.e., for all  $1 \leq i, j \leq n$  and all  $(t, x_i(t)) \in V_i$  and  $(t', x_j(t')) \in V_j$ , if  $t + \varepsilon \leq t'$   
144 then  $(t, x_i(t)) \rightsquigarrow (t', x_j(t'))$ .

145 *Example 2.* **TODO: distributed signal, happened-before relation**

146 **Definition 3.** Let  $(S, \rightsquigarrow)$  be a distributed signal of  $n$  signals, and  $V = \bigcup_{i=1}^n V_i$   
147 be the set of its events. A set  $C \subseteq V$  is a consistent cut iff for every event in  
148  $C$ , all events that happened before it also belong to  $C$ , i.e., for all  $e, e' \in V$ , if  
149  $e \in C$  and  $e' \rightsquigarrow e$ , then  $e' \in C$ .

150 We denote by  $\mathbb{C}(T)$  the (infinite) set of consistent cuts at global time  $T$ .  
 151 Given a consistent cut  $C$ , its *frontier*  $\text{front}(C) \subseteq C$  is the set consisting of the  
 152 last events in  $C$  of each signal, i.e.,  $\text{front}(C) = \bigcup_{i=1}^n \{(t, x_i(t)) \in V_i \cap C \mid \forall t' >$   
 153  $t : (t', x_i(t')) \notin V_i \cap C\}$ .

154 **Definition 4.** A consistent cut flow is a function  $\text{ccf} : \mathbb{R}_{\geq 0} \rightarrow 2^V$  that maps a  
 155 global clock value  $T$  to the frontier of a consistent cut at time  $T$ , i.e.,  $\text{ccf}(T) \in$   
 156  $\{\text{front}(C) \mid C \in \mathbb{C}(T)\}$ .

157 For all  $T, T' \in \mathbb{R}_{\geq 0}$  and  $1 \leq i \leq n$ , if  $T < T'$ , then for every pair of  
 158 events  $(c_i(T), x_i(c_i(T))) \in \text{ccf}(T)$  and  $(c_i(T'), x_i(c_i(T')))) \in \text{ccf}(T')$  we have  
 159  $(c_i(T), x_i(c_i(T))) \rightsquigarrow (c_i(T'), x_i(c_i(T')))$ . We denote by  $\text{CCF}(S, \rightsquigarrow)$  the set of all  
 160 consistent cut flows of the distributed signal  $(S, \rightsquigarrow)$ .

161 *Example 5. TODO: consistent cut, frontier, consistent cut flow*

162 Observe that a consistent cut flow of a distributed signal induces a vector  
 163 of synchronous signals which can be evaluated using the standard semantics  
 164 described in Section 2. Let  $(S, \rightsquigarrow)$  be a distributed signal of  $n$  signals  $x_1, \dots, x_n$ .  
 165 A consistent cut flow  $\text{ccf} \in \text{CCF}(S, \rightsquigarrow)$  yields a trace  $w_{\text{ccf}} = (x'_1, \dots, x'_n)$  on a  
 166 temporal domain  $[0, D]$  where  $D \in \mathbb{R}_{> 0}$  such that  $(c_i(T), x_i(c_i(T))) \in \text{ccf}(T)$   
 167 implies  $x'_i(T) = x_i(c_i(T))$  for all  $1 \leq i \leq n$  and  $T \in [0, D]$ . The set of traces of  
 168  $(S, \rightsquigarrow)$  is given by  $\text{Tr}(S, \rightsquigarrow) = \{w_{\text{ccf}} \mid \text{ccf} \in \text{CCF}(S, \rightsquigarrow)\}$ .

169 We define the satisfaction of an STL formula  $\varphi$  by a distributed signal  $(S, \rightsquigarrow)$   
 170 over a three-valued domain  $\{\top, \perp, ?\}$ . If the set of synchronous traces  $\text{Tr}(S, \rightsquigarrow)$   
 171 defined by a distributed signal  $(S, \rightsquigarrow)$  is contained in the set of traces allowed  
 172 by the formula  $\varphi$ , then  $(S, \rightsquigarrow)$  satisfies  $\varphi$ . Similarly, if  $\text{Tr}(S, \rightsquigarrow)$  has an empty  
 173 intersection with the set of traces  $\varphi$  defines, then  $(S, \rightsquigarrow)$  violates  $\varphi$ . Otherwise,  
 174 the evaluation is inconclusive since some traces satisfy the property and some  
 175 violate it.

$$[(S, \rightsquigarrow) \models \varphi] = \begin{cases} \top & \text{if } \forall w \in \text{Tr}(S, \rightsquigarrow) : w \models \varphi \\ \perp & \text{if } \forall w \in \text{Tr}(S, \rightsquigarrow) : w \models \neg \varphi \\ ? & \text{otherwise} \end{cases}$$

### 176 3 Overapproximation of the STL Distributed Semantics

177 To address the computational overhead in exact distributed monitoring, we de-  
 178 fine a new logic  $\text{STL}^+$  whose syntax is the same as STL but semantics provide  
 179 a sound approximation of the STL distributed semantics presented in Section 2.  
 180 In essence, given a distributed signal  $(S, \rightsquigarrow)$ ,  $\text{STL}^+$  considers an overapproxima-  
 181 tion  $\text{Tr}^+(S, \rightsquigarrow)$  of the set  $\text{Tr}(S, \rightsquigarrow)$  of synchronous traces. A signal  $(S, \rightsquigarrow)$  satisfies  
 182 (resp. violates) an  $\text{STL}^+$  formula  $\varphi$  iff all the traces in  $\text{Tr}^+(S, \rightsquigarrow)$  belong to the  
 183 language of  $\varphi$  (resp.  $\neg \varphi$ ).

$$[(S, \rightsquigarrow) \models \varphi]_+ = \begin{cases} \top & \text{if } \forall w \in \text{Tr}^+(S, \rightsquigarrow) : w \models \varphi \\ \perp & \text{if } \forall w \in \text{Tr}^+(S, \rightsquigarrow) : w \models \neg \varphi \\ ? & \text{otherwise} \end{cases}$$

184 In Sections 4 and 5, we respectively define  $\text{Tr}^+$  and present an algorithm  
 185 to compute the semantics of  $\text{STL}^+$ . We finally prove the correctness of our  
 186 approach.

187 **Theorem 6.** *For every STL formula  $\varphi$  and every distributed signal  $(S, \rightsquigarrow)$ , if*  
 188  *$[(S, \rightsquigarrow) \models \varphi]_+ = \top$  (resp.  $\perp$ ) then  $[(S, \rightsquigarrow) \models \varphi] = \top$  (resp.  $\perp$ ).*

## 189 4 Overapproximation of Synchronous Traces

190 In this section, given a distributed signal  $(S, \rightsquigarrow)$ , we describe an overapproxima-  
 191 tion  $\text{Tr}^+(S, \rightsquigarrow)$  of its set  $\text{Tr}(S, \rightsquigarrow)$  of synchronous traces. First, we present the  
 192 notion of *canonical segmentation*, a systematic way of partitioning the temporal  
 193 domain of a given distributed signal to keep track of the partial asynchrony.  
 194 Second, we introduce the notion of *value expressions*, sets of finite words repre-  
 195 senting how a signal behaves in a time interval. Finally, we define  $\text{Tr}^+$  based on  
 196 these notions, and show that it soundly approximates  $\text{Tr}$ .

197 *Remark 7.* We assume boolean signals in this section for convenience. The defi-  
 198 nitions and results presented here extend to real-valued signals.

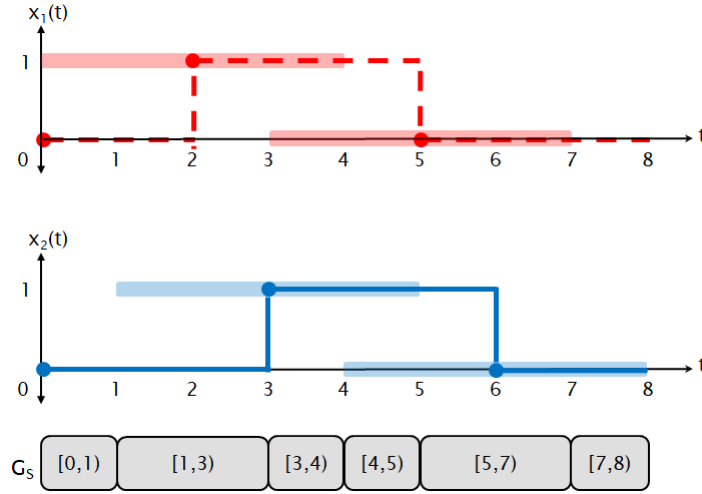
199 **Canonical Segmentation** Consider a boolean signal  $x$  with a rising edge at  
 200 time  $t > \varepsilon$ . Due to clock skew, this edge occurs in the range  $(t - \varepsilon, t + \varepsilon)$   
 201 from the monitor's point of view. This range is called an *uncertainty region*  
 202 because in  $(t - \varepsilon, t + \varepsilon)$  the monitor cannot tell the value of  $x$  precisely, but  
 203 only that it changes from 0 to 1. Formally, given an edge  $(t, x(t))$ , we define  
 204  $\theta_{\text{lo}}(x, t) = \max(0, t - \varepsilon)$  and  $\theta_{\text{hi}}(x, t) = t + \varepsilon$  as the end points of the edge's  
 205 uncertainty region.

206 Given a temporal domain  $I = [0, d) \subset \mathbb{R}_{\geq 0}$ , a *segmentation* of  $I$  is a partition  
 207 of  $I$  into finitely many intervals  $I_1, \dots, I_k$ , called *segments*, of the form  $I_j =$   
 208  $[t_j, t_{j+1})$  such that  $t_j < t_{j+1}$  for all  $1 \leq j \leq k$ . By extension, a segmentation of  
 209 a collection of signals with the same temporal domain  $I$  is a segmentation of  $I$ .

210 Let  $(S, \rightsquigarrow)$  be a distributed signal of  $n$  signals. The *canonical segmentation*  
 211  $G_S$  of  $(S, \rightsquigarrow)$  is the segmentation of  $S$  where the end points of the segments co-  
 212 incide with the end points of its temporal domain and uncertainty regions. For-  
 213 mally, we define  $G_S$  as follows. For each signal  $x_i$ , let  $F_i$  be the set of end points  
 214 of its uncertainty regions. Let  $d' = \max(d, \max(\bigcup_{i=1}^n F_i))$ , which corresponds to  
 215 the duration of the distributed signal with respect to the monitor's clock. Let  
 216  $F = \{0, d'\} \cup \bigcup_{i=1}^n F_i$  and let  $(s_j)_{1 \leq j \leq |F|}$  be a nondecreasing sequence of clock  
 217 values corresponding to the elements of  $F$ . Then, the canonical segmentation of  
 218  $(S, \rightsquigarrow)$  is  $G_S = \{I_1, \dots, I_{|F|-1}\}$  where  $I_j = [s_j, s_{j+1})$  for all  $1 \leq j < |F|$ .

219 *Example 8.* Let  $(S, \rightsquigarrow)$  be a distributed boolean signal with  $S = (x_1, x_2)$  and  
 220  $\varepsilon = 2$  over the temporal domain  $[0, 8)$  as given in Figure 1. Both signals are  
 221 initially 0. The signal  $x_1$  has a rising edge at time 2 and a falling edge at time 5,  
 222 while  $x_2$  has a rising edge at time 3 and a falling edge at time 6. The uncertainty  
 223 regions of  $x_1$  are  $(0, 4)$  and  $(3, 7)$ , while those of  $x_2$  are  $(1, 5)$  and  $(4, 8)$ . Then,

we have  $F = \{0, 8\} \cup \{0, 1, 3, 4, 5, 7, 8\}$ , and thus the canonical segmentation is  
 $G_S = \{[0, 1), [1, 3), [3, 4), [4, 5), [5, 7), [7, 8)\}$ .



**Fig. 1.** The signals  $x_1$  (top, red, dashed) and  $x_2$  (bottom, blue, solid) from Example 8. The edges are marked with solid balls and their uncertainty regions are given as semi-transparent boxes around the edges. The resulting canonical segmentation  $G_S$  is shown below the graphical representation of the signals.

**Value Expressions** Consider a boolean signal  $x$  with a rising edge with an uncertainty region of  $(t_1, t_2)$ . As discussed above, the monitor only knows that the value of  $x$  changes from 0 to 1 in this interval. We represent this knowledge as a finite word  $v = 0 \cdot 1$ . This representation is called a *value expression* and it encodes the uncertain behavior of an observed signal relative to the monitor. Formally, a value expression is an element of  $\Sigma^*$  where  $\Sigma$  is the finite alphabet of values the signal takes. Given a signal  $x$  and an edge  $(t, x(t))$ , the value expression corresponding to the uncertainty region  $(\theta_{lo}(x, t), \theta_{hi}(x, t))$  is given by  $v_{x,t} = v_- \cdot v_+$  where  $v_- = \lim_{s \rightarrow t^-} x(s)$  and  $v_+ = \lim_{s \rightarrow t^+} x(s)$ .

Notice that (i) uncertainty regions may overlap, and (ii) the canonical segmentation may split an uncertainty region into multiple segments. Consider a signal  $x$  with a rising edge in  $(1, 5)$  and a falling edge in  $(4, 8)$ . The corresponding value expressions are respectively  $v_1 = 0 \cdot 1$  and  $v_2 = 1 \cdot 0$ . Notice that the behavior of  $x$  in the interval  $[1, 4)$  can be expressed as  $\text{prefix}(v_1)$ , encoding whether the rising edge has happened yet or not. Similarly, the behavior in  $[4, 5)$  is given by  $\text{suffix}(v_1) \cdot \text{prefix}(v_2)$ , which captures whether the edges occur in this interval

(thanks to prefixing and suffixing) and the fact that the rising edge happens before the falling edge (thanks to concatenation).

Formally, given a distributed signal  $(S, \rightsquigarrow)$ , we define a function  $\gamma : S \times G_S \rightarrow 2^{\Sigma^*}$  that maps each signal and segment of the canonical segmentation to a set of value expressions, capturing the signal's potential behaviors in the given segment. Let  $x$  be a signal in  $S$ , and let  $R_1, \dots, R_m$  be its uncertainty regions where  $R_i = (t_i, t'_i)$  and the corresponding value expression is  $v_i$  for all  $1 \leq i \leq m$ . Now, let  $I \in G_S$  be a segment with  $I = [s, s']$  and for each  $1 \leq i \leq m$  define the set  $V_i$  of value expressions capturing how  $I$  relates with  $R_i$  as follows:

$$V_i = \begin{cases} \{v_i\} & \text{if } t_i = s \wedge s' = t'_i \\ \text{prefix}(v_i) & \text{if } t_i = s \wedge s' < t'_i \\ \text{suffix}(v_i) & \text{if } t_i > s \wedge s' = t'_i \\ \text{infix}(v_i) & \text{if } t_i > s \wedge s' < t'_i \\ \{\epsilon\} & \text{otherwise} \end{cases}$$

The last case happens only when  $I \cap R_i$  is empty. We finally define  $\gamma$  as follows:

$$\gamma(x, I) = \text{destutter}(V_1 \cdot V_2 \cdot \dots \cdot V_m) \setminus \{\epsilon\}$$

Observe that  $\gamma(x, I)$  contains all the potential behaviors of  $x$  in segment  $I$  by construction. However, it is potentially overapproximate. This is mainly because the sets  $V_1, \dots, V_m$  contain redundancy by definition and the concatenation does not guarantee that an edge is considered exactly once.

*Example 9.* Recall the distributed signal  $(S, \rightsquigarrow)$  in Example 8 and Figure 1. In Figure 2a, we show the value expressions corresponding to its uncertainty regions. For example, the falling edge of  $x_1$  has an uncertainty region of  $(3, 7)$ , represented by the value expression  $1 \cdot 0$ . In Figure 2b, we give the function  $\gamma$  for  $(S, \rightsquigarrow)$ . For example,  $\gamma(x_1, [3, 4])$  is obtained from  $\text{suffix}(0 \cdot 1) \cdot \text{prefix}(1 \cdot 0)$  and  $\gamma(x_2, [0, 1]) = \{0\}$ .


**Overapproximation of Tr** Consider a distributed signal  $(S, \rightsquigarrow)$  of  $n$  signals, and let  $G_S$  be its canonical segmentation. We describe how the function  $\gamma$  defines a set  $\text{Tr}^+(S, \rightsquigarrow)$  of synchronous traces that overapproximates the set  $\text{Tr}(S, \rightsquigarrow)$ .

Let  $x \in S$  and  $x'$  be two signals with the same temporal domain, and let  $I = [s, s']$  be a segment in  $G_S$ . Let  $(t_1, x'(t_1)), \dots, (t_\ell, x'(t_\ell))$  be the edges of  $x'$  in segment  $I$  with  $t_i < t_{i+1}$  for all  $1 \leq i < \ell$ . The signals  $x$  and  $x'$  are *consistent* in  $I$  iff the value expression  $x'(s) \cdot x'(t_1) \cdot \dots \cdot x'(t_\ell)$  belongs to  $\gamma(x, I)$ . Moreover,  $x$  and  $x'$  are *consistent* iff they are consistent in  $I$  for all  $I \in G_S$ . Now, let  $S = (x_1, \dots, x_n)$  and define  $\text{Tr}^+(S, \rightsquigarrow)$  as follows:

$$\text{Tr}^+(S, \rightsquigarrow) = \{(x'_1, \dots, x'_n) \mid x_i \text{ and } x'_i \text{ are consistent for all } 1 \leq i \leq n\}$$

*Example 10.* Recall the distributed signal  $(S, \rightsquigarrow)$  in Example 8 whose  $\gamma$  function is given in Figure 2b. Consider the synchronous trace  $w \in \text{Tr}(S, \rightsquigarrow)$  where the





**Fig. 2.** (a) The uncertainty regions of the distributed signal in Example 8 and the corresponding value expressions. (b) The tabular representation of the function  $\gamma$  for the given distributed signal.

273 rising edges of both signals occur at time 3 and the falling edges at time 5. One  
 274 can verify that  $w \in \text{Tr}^+(S, \rightsquigarrow)$  since for each  $i \in \{1, 2\}$  the value expression 1 is  
 275 contained in  $\gamma(x_i, [3, 4))$  and  $\gamma(x_i, [4, 5))$  while 0 is contained in the remaining  
 276 sets  $\gamma$  maps  $x_i$  to.

277 Now, consider a synchronous trace  $(x'_1, x'_2)$  where both signals are initially 0,  
 278 have rising edges at time 2 and 3.5, and falling edges at time 3 and 5. Evidently,  
 279 this trace does not belong to  $\text{Tr}(S, \rightsquigarrow)$  since  $x'_1$  and  $x'_2$  have more edges than  $x_1$   
 280 and  $x_2$ . Nonetheless, it belongs to  $\text{Tr}^+(S, \rightsquigarrow)$  since  $x'_1$  and  $x'_2$  are respectively  
 281 consistent with  $x_1$  and  $x_2$ . To witness, notice that for each  $i \in \{1, 2\}$  the value  
 282 expression  $0 \cdot 1$  is contained in  $\gamma(x_i, [1, 3))$  and  $\gamma(x_i, [3, 4))$ , the expression 1  
 283 is contained in  $\gamma(x_i, [4, 5))$ , and 0 is contained in the remaining sets  $\gamma$  maps  $x_i$  to.

284 Finally, we show the correctness of the canonical overapproximation.

285 **Lemma 11.** *For every distributed signal  $(S, \rightsquigarrow)$ , we have  $\text{Tr}(S, \rightsquigarrow) \subseteq \text{Tr}^+(S, \rightsquigarrow)$ .*

## 286 5 Monitoring Algorithm

287 In this section, given a distributed signal  $(S, \rightsquigarrow)$ , we describe an algorithm to  
 288 compute  $[(S, \rightsquigarrow) \models \varphi]_+$ . The algorithm makes use of the function  $\gamma$  defined in  
 289 Section 4 without explicitly computing  $\text{Tr}^+(S, \rightsquigarrow)$ . To achieve this, we first de-  
 290 scribe the notion of *asynchronous product* of value expressions to capture poten-  
 291 tial interleavings within segments. Then, we present an algorithm to compute  
 292 the *semantics*  $[(S, \rightsquigarrow) \models \varphi]_+$  using operations on value expressions and their  
 293 asynchronous products that encode  $\gamma$ . Finally, we discuss how *bit vectors* can  
 294 represent asynchronous products and allow us to manipulate them faster.

295 *Remark 12.* For the sake of convenience, we focus on boolean signals for the rest  
 296 of the section. Note that asynchronous products and the algorithm to compute  
 297  $[(S, \rightsquigarrow) \models \varphi]_+$  can be extended to value expressions over arbitrary finite alphabets,  
 298 e.g., encoding real-valued signals. This allows us to express more complex  
 299 properties where atomic propositions can be functions of real-valued signals.

300 **Asynchronous Product of Value Expressions** Consider the value expres-  
 301 sions  $u_1 = 0 \cdot 1$  and  $u_2 = 1 \cdot 0$  encoding the behaviors of two signals within a  
 302 segment. Due to partial asynchrony, the behaviors within segments can be seen  
 303 as completely asynchronous. To capture the potential interleavings of these be-  
 304 haviors, we consider how the values in  $u_1$  and  $u_2$  can align. In particular, there  
 305 are three potential alignments: (i) the rising edge of  $u_1$  happens before the falling  
 306 edge of  $u_2$ , (ii) the falling edge of  $u_2$  happens before the rising edge of  $u_1$ , and  
 307 (iii) the two edges happen simultaneously. We respectively represent these with  
 308 the tuples  $(011, 110)$ ,  $(001, 100)$ , and  $(01, 10)$  where the first component encodes  
 309  $u_1$  and the second  $u_2$ .

Formally, given two value expressions  $u_1$  and  $u_2$ , we define their *asynchronous product* as follows:

$$u_1 \otimes u_2 = \{\text{destutter}(v_1, v_2) \mid v_i \in \text{stutter}_k(u_i), k = |u_1| + |u_2| - 1, i \in \{1, 2\}\}$$

310 Moreover, given two sets  $L_1$  and  $L_2$  of value expressions, we write  $L_1 \otimes L_2 =$   
 311  $\{u_1 \otimes u_2 \mid u_i \in L_i, i \in \{1, 2\}\}$ .

312 Asynchronous products of value expressions allow us to lift value expressions  
 313 to satisfaction signals of formulas.

314 *Example 13.* Recall the distributed signal  $(S, \rightsquigarrow)$  in Example 8 and its  $\gamma$  function  
 315 given in Figure 2b. Suppose we want to compute the value expressions encoding  
 316 the satisfaction of  $x_1 \wedge x_2$  in the segment  $[1, 3)$ . We can achieve this by first com-  
 317 puting the asynchronous product  $\gamma(x_1, [3, 4)) \otimes \gamma(x_2, [3, 4))$ , and then computing  
 318 the bitwise conjunction of each pair in the set. For example, considering the ex-  
 319 pression  $0 \cdot 1 \cdot 0$  for  $x_1$  and  $0 \cdot 1$  for  $x_2$ , the product contains the pair  $(010, 011)$ .  
 320 Taking the bitwise conjunction of this pair gives us the expression  $0 \cdot 1 \cdot 0$  as a  
 321 potential behavior for the satisfaction of  $x_1 \wedge x_2$  in this segment.

322 **Computing the Semantics of STL<sup>+</sup>** As hinted in Example 13, to compute  
 323 the semantics, we apply bitwise operations on value expressions and their asyn-  
 324 chronous products to transform them into encodings of satisfaction signals of  
 325 formulas. Consider the distributed signal  $(S, \rightsquigarrow)$  in Example 8 and suppose we  
 326 want to compute  $[(S, \rightsquigarrow) \models \Diamond(x_1 \wedge x_2)]_+$ . To achieve this, we first compute for  
 327 each segment in  $G_S$  the set of value expressions for the satisfaction of  $x_1 \wedge x_2$ ,  
 328 and then from these compute that of  $\Diamond(x_1 \wedge x_2)$ . This compositional approach  
 329 allows us to evaluate arbitrary STL<sup>+</sup> formulas.

330 Let us start with the building blocks of our approach. First, we define bitwise  
 331 operations on boolean value expressions encoding atomic propositions. Then, we  
 332 present how these extend to asynchronous products of value expressions.

Let  $u$  and  $v$  be boolean value expressions of length  $\ell$ . We denote by  $u \& v$  the bitwise-and operation, by  $u \mid v$  the bitwise-or, and by  $\sim u$  the bitwise-negation. In addition, we define the *bitwise strong-until* operator as follows:

$$u \mathcal{U}^0 v = \left( \max_{i \leq j \leq \ell} \left( \min \left( v[j], \min_{i \leq k \leq j} u[k] \right) \right) \right)_{1 \leq i \leq \ell}$$

As usual, we derive *bitwise eventually* as  $\mathbf{E}u = 1^\ell \mathcal{U}^0 u$ , *bitwise always* as  $\mathbf{A}u = \sim(\mathbf{E}\sim u)$ , and *bitwise weak-until* as follows:

$$u \mathcal{U}^1 v = \left( \max \left( u[i..] \mathcal{U}^0 v[i..], \mathbf{A}u[i..] \right) \right)_{1 \leq i \leq \ell}$$

This distinction will be useful later when we evaluate a formula segment by segment. We remark that the definitions of these operators coincide with the robustness semantics of (discrete time) STL. Finally, note that the output of these operations is a value expression of length  $\ell$ . For example, if  $u = 010$ , we have  $\mathbf{E}u = 110$  and  $\mathbf{A}u = 000$ .

Now, let  $L_1$  and  $L_2$  be sets of boolean value expressions. We define the following untimed operations:

$$\begin{aligned} \neg L_1 &= \{\sim u \mid u \in L_1\} \\ L_1 \wedge L_2 &= \text{destutter}(\{u_1 \& u_2 \mid (u_1, u_2) \in L_1 \otimes L_2\}) \\ L_1 \mathcal{U}^x L_2 &= \text{destutter}(\{u_1 \mathcal{U}^x u_2 \mid (u_1, u_2) \in L_1 \otimes L_2\}) \text{ for } x \in \{0, 1\} \end{aligned}$$

Similarly as above, we derive  $L_1 \vee L_2 = \neg(L_1 \wedge L_2)$ ,  $\Diamond L_1 = \{1\} \mathcal{U}^0 L_1$ , and  $\Box L_1 = \neg \Diamond \neg L_1$ .

*Example 14.* demonstrate above operations

We additionally define the following timed operators:

TODO

Finally, given a distributed signal  $(S, \rightsquigarrow)$  and an STL<sup>+</sup> formula  $\varphi$ , we compute  $[(S, \rightsquigarrow) \models \varphi]_+$  as follows.

TODO

**Boolean Asynchronous Products as Bit Vectors** getting to atomic propositions represented as bitvectors intuition on shortcuts provide equivalent shortcuts for bitvectors.

## 6 Experimental Evaluation

TODO

## 7 Conclusion

TODO

## 360 References

- 361 1. Maler, O., Nickovic, D.: Monitoring properties of analog and mixed-signal circuits.  
362 Int. J. Softw. Tools Technol. Transf. **15**(3), 247–268 (2013). [https://doi.org/10.](https://doi.org/10.1007/s10009-012-0247-9)  
363 1007/s10009-012-0247-9

## 364 Appendix

365 *Proof (Proof of Theorem 6).* **TODO**

366 *Proof (Proof of Lemma 11).* **TODO**

367 *Proof (Proof of ??).* **TODO**