

Approximate Distributed Monitoring under Partial Synchrony: Balancing Speed with Accuracy

Author(s)

Affiliation(s)

Abstract. In distributed systems with processes that do not share a global clock, *partial synchrony* is achieved by clock synchronization that guarantees bounded clock skew among all applications. Existing solutions for distributed runtime verification under partial synchrony against temporal logic specifications are exact but suffer from significant computational overhead. In this paper, we propose an *approximate* distributed monitoring algorithm for Signal Temporal Logic (STL) that mitigates this issue by abstracting away potential interleaving behaviors. This conservative abstraction enables a significant speedup of the distributed monitors, albeit with a trade-off in accuracy. We address this trade-off with a methodology that combines our approximate monitor with its exact counterpart, resulting in enhanced monitoring efficiency without sacrificing precision. We validate our approach with multiple experiments, showcasing its effectiveness and efficacy on both a real-world application and synthetic examples.

You can leave notes using the command `\firstName{note}`.
Page limit: 16 + 4 (appendix)

1 Introduction

Distributed systems are networks of independent agents that work together to achieve a common objective. Distributed systems are everywhere around us and come in many different forms. For example, cloud computing uses distribution of resources and services over the internet to offer to their users a scalable infrastructure with transparent on-demand access to computing power and storage. Swarms of drones represent another family of distributed systems where individual drones collaborate to accomplish tasks like surveillance, search and rescue, or package delivery. While each drone operates independently, it also communicates and coordinates with others to successfully achieve their common objectives. The individual agents in a distributed system typically do not share a global clock. To coordinate actions across multiple agents, clock synchronization is often needed. While perfect clock synchronization is impractical due to network latency and node failures, algorithms such as the Network Time Protocol (NTP) allow agents

36 to maintain a *bounded skew* between the synchronized clocks. In that case, we
 37 say that a distributed system has *partial synchrony*.

38 Formal verification of distributed system is a notoriously hard problem, due
 39 to the combinatorial explosion of all possible interleavings in the behaviors col-
 40 lected from individual agents. *Runtime verification (RV)* provides a more prag-
 41 matic approach, in which a monitor observes a behavior of a distributed sys-
 42 tem and checks its correctness against a formal specification. The problem of
 43 distributed RV under partial synchrony assumption has been studied for Lin-
 44 ear Temporal Logic (LTL) and Signal Temporal Logic (STL) specification lan-
 45 guages. The proposed solutions use Satisfiability-Modulo-Theory (SMT) solving
 46 to provide sound and complete distributed monitoring procedures. Although dis-
 47 tributed RV monitors consume only a single distributed behavior at a time, this
 48 behavior can nevertheless have an excessive number of possible interleavings.
 49 Hence, the exact distributed monitors from the literature can still suffer from
 50 significant computational overhead.

51 To mitigate this issue, we present an approach for *approximate* RV of STL
 52 specifications under partial synchrony. In essence, we abstract away potential
 53 interleavings in distributed behaviors in a conservative manner, resulting in an
 54 effective over-approximation of global behaviors. This abstraction simplifies the
 55 representation of distributed behaviors and the monitoring operations required
 56 to evaluate temporal specifications. There is an inevitable trade-off in approxi-
 57 mate RV – gains in the monitoring speed-up may result in reduced accuracy. For
 58 some applications, reduced accuracy may not be acceptable. Therefore, we pro-
 59 pose a methodology that combines our approximate monitors with their exact
 60 counterparts, with the aim to benefit from the enhanced monitoring efficiency
 61 without sacrificing precision. We implemented our approach in a prototype tool
 62 and performed thorough evaluations on both synthetic and real-world case stud-
 63 ies. We first demonstrated that our approximate monitors achieve speed-ups of
 64 several orders of magnitudes compared to the exact SMT-based distributed RV
 65 solution. We empirically characterized the classes of specifications and behaviors
 66 for which our approximate monitoring approach achieves good precision. We fi-
 67 nally showed that by combining exact and approximate distributed RV, there is
 68 still a significant efficiency gain on average without the sacrifice of the precision,
 69 even in cases where approximate monitors have low accuracy.

70 2 Preliminaries

71 We denote by $\mathbb{B} = \{\top, \perp\}$ the set of Booleans, \mathbb{R} the set of reals, $\mathbb{R}_{\geq 0}$ the set of
 72 nonnegative reals, and $\mathbb{R}_{> 0}$ the set of positive reals. An interval $I \subseteq \mathbb{R}$ of reals
 73 with the end points $a < b$ has length $|b - a|$.

74 Let Σ be a finite *alphabet*. We denote by Σ^* the set of finite words over
 75 Σ and by ϵ the empty word. For $u \in \Sigma^*$, we respectively write $\text{prefix}(u)$ and
 76 $\text{suffix}(u)$ for the sets of prefixes and suffixes of u . We also let $\text{infix}(u) = \{v \in$
 77 $\Sigma^* \mid \exists x, y \in \Sigma^* : u = xvy\}$. For a nonempty word $u \in \Sigma^*$ and $1 \leq i \leq |u|$,
 78 we denote by $u[i]$ the i th letter of u , by $u[..i]$ the prefix of u of length i , and by

79 $u[i..]$ the suffix of u of length $|u| - i + 1$. Given $u \in \Sigma^*$ and $\ell \geq 1$, we denote by
 80 u^ℓ the word obtained by concatenating u by itself $\ell - 1$ times. Moreover, given
 81 $L \subseteq \Sigma^*$, we define $\text{first}(L) = \{u[0] \mid u \in L\}$. For sets $L_1, L_2 \in \Sigma^*$ of words, we
 82 let $L_1 \cdot L_2 = \{u \cdot v \mid u \in L_1, v \in L_2\}$. For tuples (u_1, \dots, u_m) and (v_1, \dots, v_m) of
 83 words, we let $(u_1, \dots, u_m) \cdot (v_1, \dots, v_m) = (u_1v_1, \dots, u_mv_m)$.

84 We define the function $\text{destutter} : \Sigma^* \rightarrow \Sigma^*$ inductively as follows. For all
 85 $\sigma \in \Sigma \cup \{\epsilon\}$, let $\text{destutter}(\sigma) = \sigma$. For all $u \in \Sigma^*$ such that $u = \sigma_1\sigma_2v$ for
 86 some $\sigma_1, \sigma_2 \in \Sigma$ and $v \in \Sigma^*$, let (i) $\text{destutter}(u) = \text{destutter}(\sigma_2v)$ if $\sigma_1 = \sigma_2$,
 87 and (ii) $\text{destutter}(u) = \sigma_1 \cdot \text{destutter}(\sigma_2v)$ otherwise. By extension, for a set
 88 $L \subseteq \Sigma^*$ of finite words, we write $\text{destutter}(L) = \{\text{destutter}(u) \mid u \in L\}$. Given
 89 a tuple $(u_1, \dots, u_m) = (\sigma_{1,1}\sigma_{1,2}v_1, \dots, \sigma_{m,1}\sigma_{m,2}v_m)$ of finite words of the same
 90 length, we define $\text{destutter}(u_1, \dots, u_m)$ as expected: (i) $\text{destutter}(u_1, \dots, u_m) =$
 91 $\text{destutter}(\sigma_{1,2}v_1, \dots, \sigma_{m,2}v_m)$ if $\sigma_{i,1} = \sigma_{i,2}$ for all $1 \leq i \leq m$, and (ii) $\text{destutter}(u_1, \dots, u_m) =$
 92 $(\sigma_{1,1}, \dots, \sigma_{m,1}) \cdot \text{destutter}(\sigma_{1,2}v_1, \dots, \sigma_{m,2}v_m)$ otherwise.

93 Moreover, given an integer $k \geq 0$, we define $\text{stutter}_k : \Sigma^* \rightarrow \Sigma^*$ such
 94 that $\text{stutter}_k(u) = \{v \in \Sigma^* \mid |v| = k \wedge \text{destutter}(v) = \text{destutter}(u)\}$ if $k \geq$
 95 $|\text{destutter}(u)|$, and $\text{stutter}_k(u) = \emptyset$ otherwise.

96 **Signal Temporal Logic (STL)** Let $A, B \subset \mathbb{R}$. A function $f : A \rightarrow B$ is
 97 *right-continuous* iff $\lim_{a \rightarrow c^+} f(a) = f(c)$ for all $c \in A$, and *non-Zeno* iff for
 98 every bounded interval $I \subseteq A$ there are finitely many $a \in I$ such that f is not
 99 continuous at a . A *signal* is a right-continuous, non-Zeno, piecewise-constant
 100 function $x : [0, d) \rightarrow \mathbb{R}$ where $d \in \mathbb{R}_{>0}$ is the duration of x and $[0, d)$ is its
 101 temporal domain. Let $x : [0, d) \rightarrow \mathbb{R}$ be a signal. An *event* of x is a pair $(t, x(t))$
 102 where $t \in [0, d)$. An *edge* of x is an event $(t, x(t))$ such that $\lim_{s \rightarrow t^-} x(s) \neq$
 103 $\lim_{s \rightarrow t^+} x(s)$. In particular, an edge is *rising* if $\lim_{s \rightarrow t^-} x(s) < \lim_{s \rightarrow t^+} x(s)$, and
 104 it is *falling* otherwise. A signal $x : [0, d) \rightarrow \mathbb{R}$ can be represented finitely by its
 105 initial value and edges: if x has m edges, then $x = (t_0, v_0)(t_1, v_1) \dots (t_m, v_m)$
 106 such that $t_0 = 0$, $t_{i-1} < t_i$, and (t_i, v_i) is an edge of x for all $1 \leq i \leq m$.

107 Let AP be a set of atomic propositions. The syntax is given by the following
 108 grammar where $p \in \text{AP}$ and $I \subseteq \mathbb{R}_{\geq 0}$ is an interval.

$$\varphi := p \mid \neg\varphi \mid \varphi \wedge \varphi \mid \varphi \mathcal{U}_I \varphi$$

109 A *trace* $w = (x_1, \dots, x_n)$ is a finite vector of signals. We express atomic
 110 propositions as functions of trace values at a time point t , i.e., a proposition
 111 $p \in \text{AP}$ over a trace $w = (x_1, \dots, x_n)$ is defined as $f_p(x_1(t), \dots, x_n(t)) > 0$
 112 where $f_p : \mathbb{R}^n \rightarrow \mathbb{R}$ is a function. Given intervals $I, J \subseteq \mathbb{R}_{\geq 0}$, we define $I \oplus J =$
 113 $\{i + j \mid i \in I \wedge j \in J\}$, and we simply write t for the singleton set $\{t\}$.

114 Below we recall the finite-trace qualitative semantics of STL defined over \mathbb{B}
 115 $[1]$. Let $d \in \mathbb{R}_{>0}$ and $w = (x_1, \dots, x_n)$ with $x_i : [0, d) \rightarrow \mathbb{R}$ for all $1 \leq i \leq n$. Let
 116 φ_1, φ_2 be STL formulas and let $t \in [0, d)$.

$$\begin{aligned}
(w, t) \models p &\iff f_p(x_1(t), \dots, x_n(t)) > 0 \\
(w, t) \models \neg \varphi_1 &\iff \overline{(w, t) \models \varphi_1} \\
(w, t) \models \varphi_1 \wedge \varphi_2 &\iff (w, t) \models \varphi_1 \wedge (w, t) \models \varphi_2 \\
(w, t) \models \varphi_1 \mathcal{U}_I \varphi_2 &\iff \exists t' \in (t \oplus I) \cap [0, d) : \\
&\quad (w, t') \models \varphi_2 \wedge \forall t'' \in (t, t') : (w, t'') \models \varphi_1
\end{aligned}$$

117 We simply write $w \models \varphi$ for $(w, 0) \models \varphi$. We additionally use the following
118 standard abbreviations: **false** = $p \wedge \neg p$, **true** = $\neg \mathbf{false}$, $\varphi_1 \vee \varphi_2 = \neg(\neg \varphi_1 \wedge$
119 $\neg \varphi_2)$, $\Diamond_I \varphi = \mathbf{true} \mathcal{U}_I \varphi$, and $\Box_I \varphi = \neg \Diamond_I \neg \varphi$. Moreover, the untimed temporal
120 operators are defined through their timed counterparts on the interval $(0, \infty)$,
121 e.g., $\varphi_1 \mathcal{U} \varphi_2 = \varphi_1 \mathcal{U}_{(0, \infty)} \varphi_2$.

122 **Distributed Semantics of STL** We consider an asynchronous and loosely-
123 coupled message-passing system of $n \geq 2$ reliable agents producing a set of
124 signals x_1, \dots, x_n , where for some $d \in \mathbb{R}_{>0}$ we have $x_i : [0, d) \rightarrow \mathbb{R}$ for all
125 $1 \leq i \leq n$. The agents do not share memory or a global clock. Only to formalize
126 statements, we speak of a *hypothetical* global clock and denote its value by T .
127 For local time values, we use the lowercase letters t and s .

128 For a signal x_i , we denote by V_i the set of its events, by E_i^\uparrow the set of its
129 rising edges, and by E_i^\downarrow that of falling edges. Moreover, we let $E_i = E_i^\uparrow \cup E_i^\downarrow$. We
130 represent the local clock of the i th agent as an increasing and divergent function
131 $c_i : \mathbb{R}_{\geq 0} \rightarrow \mathbb{R}_{\geq 0}$ that maps a global time T to a local time $c_i(T)$.

132 We assume that the system is *partially synchronous*: the agents use a clock
133 synchronization algorithm that guarantees a bounded clock skew with respect
134 to the global clock, i.e., $|c_i(T) - c_j(T)| < \varepsilon$ for all $1 \leq i, j \leq N$ and $T \in \mathbb{R}_{\geq 0}$,
135 where $\varepsilon \in \mathbb{R}_{>0}$ is the maximum clock skew.

136 **Definition 1.** A distributed signal is a pair (S, \rightsquigarrow) , where $S = (x_1, \dots, x_n)$ is a
137 vector of signals and \rightsquigarrow is the happened-before relation between events in signals
138 extended with the partial synchrony assumption as follows.

- 139 – For every agent, the events of its signals are totally ordered, i.e., for all $1 \leq$
140 $i \leq n$ and all $(t, x_i(t)), (t', x_i(t')) \in V_i$, if $t < t'$ then $(t, x_i(t)) \rightsquigarrow (t', x_i(t'))$.
- 141 – Every pair of events whose timestamps are at least ε apart is totally ordered,
142 i.e., for all $1 \leq i, j \leq n$ and all $(t, x_i(t)) \in V_i$ and $(t', x_j(t')) \in V_j$, if $t + \varepsilon \leq t'$
143 then $(t, x_i(t)) \rightsquigarrow (t', x_j(t'))$.

144 *Example 2.* **TODO: distributed signal, happened-before relation**

145 **Definition 3.** Let (S, \rightsquigarrow) be a distributed signal of n signals, and $V = \bigcup_{i=1}^n V_i$
146 be the set of its events. A set $C \subseteq V$ is a consistent cut iff for every event in
147 C , all events that happened before it also belong to C , i.e., for all $e, e' \in V$, if
148 $e \in C$ and $e' \rightsquigarrow e$, then $e' \in C$.

149 We denote by $\mathbb{C}(T)$ the (infinite) set of consistent cuts at global time T .
 150 Given a consistent cut C , its *frontier* $\text{front}(C) \subseteq C$ is the set consisting of the
 151 last events in C of each signal, i.e., $\text{front}(C) = \bigcup_{i=1}^n \{(t, x_i(t)) \in V_i \cap C \mid \forall t' >$
 152 $t : (t', x_i(t')) \notin V_i \cap C\}$.

153 **Definition 4.** A consistent cut flow is a function $\text{ccf} : \mathbb{R}_{\geq 0} \rightarrow 2^V$ that maps a
 154 global clock value T to the frontier of a consistent cut at time T , i.e., $\text{ccf}(T) \in$
 155 $\{\text{front}(C) \mid C \in \mathbb{C}(T)\}$.

156 For all $T, T' \in \mathbb{R}_{\geq 0}$ and $1 \leq i \leq n$, if $T < T'$, then for every pair of
 157 events $(c_i(T), x_i(c_i(T))) \in \text{ccf}(T)$ and $(c_i(T'), x_i(c_i(T')))) \in \text{ccf}(T')$ we have
 158 $(c_i(T), x_i(c_i(T))) \rightsquigarrow (c_i(T'), x_i(c_i(T')))$. We denote by $\text{CCF}(S, \rightsquigarrow)$ the set of all
 159 consistent cut flows of the distributed signal (S, \rightsquigarrow) .

160 *Example 5. TODO: consistent cut, frontier, consistent cut flow*

161 Observe that a consistent cut flow of a distributed signal induces a vector
 162 of synchronous signals which can be evaluated using the standard semantics
 163 described in Section 2. Let (S, \rightsquigarrow) be a distributed signal of n signals x_1, \dots, x_n .
 164 A consistent cut flow $\text{ccf} \in \text{CCF}(S, \rightsquigarrow)$ yields a trace $w_{\text{ccf}} = (x'_1, \dots, x'_n)$ on a
 165 temporal domain $[0, D]$ where $D \in \mathbb{R}_{> 0}$ such that $(c_i(T), x_i(c_i(T))) \in \text{ccf}(T)$
 166 implies $x'_i(T) = x_i(c_i(T))$ for all $1 \leq i \leq n$ and $T \in [0, D]$. The set of traces of
 167 (S, \rightsquigarrow) is given by $\text{Tr}(S, \rightsquigarrow) = \{w_{\text{ccf}} \mid \text{ccf} \in \text{CCF}(S, \rightsquigarrow)\}$.

168 We define the satisfaction of an STL formula φ by a distributed signal (S, \rightsquigarrow)
 169 over a three-valued domain $\{\top, \perp, ?\}$. If the set of synchronous traces $\text{Tr}(S, \rightsquigarrow)$
 170 defined by a distributed signal (S, \rightsquigarrow) is contained in the set of traces allowed
 171 by the formula φ , then (S, \rightsquigarrow) satisfies φ . Similarly, if $\text{Tr}(S, \rightsquigarrow)$ has an empty
 172 intersection with the set of traces φ defines, then (S, \rightsquigarrow) violates φ . Otherwise,
 173 the evaluation is inconclusive since some traces satisfy the property and some
 174 violate it.

$$[(S, \rightsquigarrow) \models \varphi] = \begin{cases} \top & \text{if } \forall w \in \text{Tr}(S, \rightsquigarrow) : w \models \varphi \\ \perp & \text{if } \forall w \in \text{Tr}(S, \rightsquigarrow) : w \models \neg \varphi \\ ? & \text{otherwise} \end{cases}$$

175 3 Overapproximation of the STL Distributed Semantics

176 To address the computational overhead in exact distributed monitoring, we de-
 177 fine a new logic STL^+ whose syntax is the same as STL but semantics provide
 178 a sound approximation of the STL distributed semantics presented in Section 2.
 179 In essence, given a distributed signal (S, \rightsquigarrow) , STL^+ considers an overapproxima-
 180 tion $\text{Tr}^+(S, \rightsquigarrow)$ of the set $\text{Tr}(S, \rightsquigarrow)$ of synchronous traces. A signal (S, \rightsquigarrow) satisfies
 181 (resp. violates) an STL^+ formula φ iff all the traces in $\text{Tr}^+(S, \rightsquigarrow)$ belong to the
 182 language of φ (resp. $\neg \varphi$).

$$[(S, \rightsquigarrow) \models \varphi]_+ = \begin{cases} \top & \text{if } \forall w \in \text{Tr}^+(S, \rightsquigarrow) : w \models \varphi \\ \perp & \text{if } \forall w \in \text{Tr}^+(S, \rightsquigarrow) : w \models \neg \varphi \\ ? & \text{otherwise} \end{cases}$$

183 In Sections 4 and 5, we respectively define Tr^+ and present an algorithm
 184 to compute the semantics of STL^+ . We finally prove the correctness of our
 185 approach.

186 **Theorem 6.** *For every STL formula φ and every distributed signal (S, \rightsquigarrow) , if*
 187 *$[(S, \rightsquigarrow) \models \varphi]_+ = \top$ (resp. \perp) then $[(S, \rightsquigarrow) \models \varphi] = \top$ (resp. \perp).*

188 4 Overapproximation of Synchronous Traces

189 In this section, given a distributed signal (S, \rightsquigarrow) , we describe an overapproxima-
 190 tion $\text{Tr}^+(S, \rightsquigarrow)$ of its set $\text{Tr}(S, \rightsquigarrow)$ of synchronous traces. First, we present the
 191 notion of *canonical segmentation*, a systematic way of partitioning the temporal
 192 domain of a given distributed signal to keep track of the partial asynchrony.
 193 Second, we introduce the notion of *value expressions*, sets of finite words repre-
 194 senting how a signal behaves in a time interval. Finally, we define Tr^+ based on
 195 these notions, and show that it soundly approximates Tr .

196 *Remark 7.* We assume boolean signals in this section for convenience. The def-
 197 initions and results presented here extend to real-valued signals because finite-
 198 length piecewise-constant signals will only use a finite number of values.

199 **Canonical Segmentation** Consider a boolean signal x with a rising edge at
 200 time $t > \varepsilon$. Due to clock skew, this edge occurs in the range $(t - \varepsilon, t + \varepsilon)$
 201 from the monitor's point of view. This range is called an *uncertainty region*
 202 because in $(t - \varepsilon, t + \varepsilon)$ the monitor cannot tell the value of x precisely, but
 203 only that it changes from 0 to 1. Formally, given an edge $(t, x(t))$, we define
 204 $\theta_{\text{lo}}(x, t) = \max(0, t - \varepsilon)$ and $\theta_{\text{hi}}(x, t) = \min(d, t + \varepsilon)$ as the end points of the
 205 edge's uncertainty region.

206 Given a temporal domain $I = [0, d] \subset \mathbb{R}_{\geq 0}$, a *segmentation* of I is a partition
 207 of I into finitely many intervals I_1, \dots, I_k , called *segments*, of the form $I_j =$
 208 $[t_j, t_{j+1})$ such that $t_j < t_{j+1}$ for all $1 \leq j \leq k$. By extension, a segmentation of
 209 a collection of signals with the same temporal domain I is a segmentation of I .

210 Let (S, \rightsquigarrow) be a distributed signal of n signals. The *canonical segmentation*
 211 G_S of (S, \rightsquigarrow) is the segmentation of S where the end points of the segments
 212 coincide with the end points of its temporal domain and uncertainty regions.
 213 Formally, we define G_S as follows. For each signal x_i , let F_i be the set of end
 214 points of its uncertainty regions. Let $F = \{0, d\} \cup \bigcup_{i=1}^n F_i$ and let $(s_j)_{1 \leq j \leq |F|}$
 215 be a nondecreasing sequence of clock values corresponding to the elements of
 216 F . Then, the canonical segmentation of (S, \rightsquigarrow) is $G_S = \{I_1, \dots, I_{|F|-1}\}$ where
 217 $I_j = [s_j, s_{j+1})$ for all $1 \leq j < |F|$.

218 *Example 8.* Let (S, \rightsquigarrow) be a distributed boolean signal with $S = (x_1, x_2)$ and
 219 $\varepsilon = 2$ over the temporal domain $[0, 8)$ as given in Figure 1. Both signals are
 220 initially 0. The signal x_1 has a rising edge at time 2 and a falling edge at time 5,
 221 while x_2 has a rising edge at time 3 and a falling edge at time 6. The uncertainty
 222 regions of x_1 are $(0, 4)$ and $(3, 7)$, while those of x_2 are $(1, 5)$ and $(4, 8)$. Then,

223 we have $F = \{0, 8\} \cup \{0, 1, 3, 4, 5, 7, 8\}$, and thus the canonical segmentation is
 224 $G_S = \{[0, 1), [1, 3), [3, 4), [4, 5), [5, 7), [7, 8)\}$.




Fig. 1. The signals x_1 (top, red, dashed) and x_2 (bottom, blue, solid) from Example 8. The edges are marked with solid balls and their uncertainty regions are given as semi-transparent boxes around the edges. The resulting canonical segmentation G_S is shown below the graphical representation of the signals.

225 **Value Expressions** Consider a boolean signal x with a rising edge with an
 226 uncertainty region of (t_1, t_2) . As discussed above, the monitor only knows that
 227 the value of x changes from 0 to 1 in this interval. We represent this knowledge
 228 as a finite word $v = 0 \cdot 1$. This representation is called a *value expression* and
 229 it encodes the uncertain behavior of an observed signal relative to the monitor.
 230 Formally, a value expression is an element of Σ^* where Σ is the finite alphabet
 231 of values the signal takes. Given a signal x and an edge $(t, x(t))$, the value
 232 expression corresponding to the uncertainty region $(\theta_{lo}(x, t), \theta_{hi}(x, t))$ is given
 233 by $v_{x,t} = v_- \cdot v_+$ where $v_- = \lim_{s \rightarrow t^-} x(s)$ and $v_+ = \lim_{s \rightarrow t^+} x(s)$. Let us
 234 remark that this definition is general because finite-length piecewise-constant
 235 real-valued signals will only have a finite number of values, making Σ finite.

236 Notice that (i) uncertainty regions may overlap, and (ii) the canonical seg-
 237 mentation may split an uncertainty region into multiple segments. Consider a
 238 signal x with a rising edge in $(1, 5)$ and a falling edge in $(4, 8)$. The corresponding
 239 value expressions are respectively $v_1 = 0 \cdot 1$ and $v_2 = 1 \cdot 0$. Notice that the be-
 240 havior of x in the interval $[1, 4)$ can be expressed as $\text{prefix}(v_1)$, encoding whether
 241 the rising edge has happened yet or not. Similarly, the behavior in $[4, 5)$ is given

by $\text{suffix}(v_1) \cdot \text{prefix}(v_2)$, which captures whether the edges occur in this interval (thanks to prefixing and suffixing) and the fact that the rising edge happens before the falling edge (thanks to concatenation).

Formally, given a distributed signal (S, \rightsquigarrow) , we define a function $\gamma : S \times G_S \rightarrow 2^{\Sigma^*}$ that maps each signal and segment of the canonical segmentation to a set of value expressions, capturing the signal's potential behaviors in the given segment. Let x be a signal in S , and let R_1, \dots, R_m be its uncertainty regions where $R_i = (t_i, t'_i)$ and the corresponding value expression is v_i for all $1 \leq i \leq m$. Now, let $I \in G_S$ be a segment with $I = [s, s']$ and for each $1 \leq i \leq m$ define the set V_i of value expressions capturing how I relates with R_i as follows:

$$V_i = \begin{cases} \{v_i\} & \text{if } t_i = s \wedge s' = t'_i \\ \text{prefix}(v_i) & \text{if } t_i = s \wedge s' < t'_i \\ \text{suffix}(v_i) & \text{if } t_i > s \wedge s' = t'_i \\ \text{infix}(v_i) & \text{if } t_i > s \wedge s' < t'_i \\ \{\epsilon\} & \text{otherwise} \end{cases}$$

The last case happens only when $I \cap R_i$ is empty. We finally define γ as follows:

$$\gamma(x, I) = \text{destutter}(V_1 \cdot V_2 \cdot \dots \cdot V_m) \setminus \{\epsilon\}$$

Observe that $\gamma(x, I)$ contains all the potential behaviors of x in segment I by construction. However, it is potentially overapproximate. This is mainly because the sets V_1, \dots, V_m contain redundancy by definition and the concatenation does not guarantee that an edge is considered exactly once.

Example 9. Recall the distributed signal (S, \rightsquigarrow) in Example 8 and Figure 1. In Figure 2a, we show the value expressions corresponding to its uncertainty regions. For example, the falling edge of x_1 has an uncertainty region of $(3, 7)$, represented by the value expression $1 \cdot 0$. In Figure 2b, we give the function γ for (S, \rightsquigarrow) . For example, $\gamma(x_1, [3, 4])$ is obtained from $\text{suffix}(0 \cdot 1) \cdot \text{prefix}(1 \cdot 0)$ and $\gamma(x_2, [0, 1]) = \{0\}$.

Overapproximation of Tr Consider a distributed signal (S, \rightsquigarrow) of n signals, and let G_S be its canonical segmentation. We describe how the function γ defines a set $\text{Tr}^+(S, \rightsquigarrow)$ of synchronous traces that overapproximates the set $\text{Tr}(S, \rightsquigarrow)$.

Let $x \in S$ and x' be two signals with the same temporal domain, and let $I = [s, s']$ be a segment in G_S . Let $(t_1, x'(t_1)), \dots, (t_\ell, x'(t_\ell))$ be the edges of x' in segment I with $t_i < t_{i+1}$ for all $1 \leq i < \ell$. The signal x' is *I-consistent with* x iff the value expression $x'(s) \cdot x'(t_1) \cdot \dots \cdot x'(t_\ell)$ belongs to $\gamma(x, I)$. Moreover, x' is *consistent with* x iff it is *I-consistent with* x for all $I \in G_S$.

Now, let $S = (x_1, \dots, x_n)$ and define $\text{Tr}^+(S, \rightsquigarrow)$ as follows:

$$\text{Tr}^+(S, \rightsquigarrow) = \{(x'_1, \dots, x'_n) \mid x'_i \text{ is consistent with } x_i \text{ for all } 1 \leq i \leq n\}$$

Example 10. Recall the distributed signal (S, \rightsquigarrow) in Example 8 whose γ function is given in Figure 2b. Consider the synchronous trace $w \in \text{Tr}(S, \rightsquigarrow)$ where the




Fig. 2. (a) The uncertainty regions of the distributed signal in Example 8 and the corresponding value expressions. (b) The tabular representation of the function γ for the given distributed signal.

274 rising edges of both signals occur at time 3 and the falling edges at time 5. One
 275 can verify that $w \in \text{Tr}^+(S, \rightsquigarrow)$ since for each $i \in \{1, 2\}$ the value expression 1 is
 276 contained in $\gamma(x_i, [3, 4))$ and $\gamma(x_i, [4, 5))$ while 0 is contained in the remaining
 277 sets γ maps x_i to.

278 Now, consider a synchronous trace (x'_1, x'_2) where both signals are initially 0,
 279 have rising edges at time 2 and 3.5, and falling edges at time 3 and 5. Evidently,
 280 this trace does not belong to $\text{Tr}(S, \rightsquigarrow)$ since x'_1 and x'_2 have more edges than x_1
 281 and x_2 . Nonetheless, it belongs to $\text{Tr}^+(S, \rightsquigarrow)$ since x'_1 and x'_2 are respectively
 282 consistent with x_1 and x_2 . To witness, notice that for each $i \in \{1, 2\}$ the value
 283 expression $0 \cdot 1$ is contained in $\gamma(x_i, [1, 3))$ and $\gamma(x_i, [3, 4))$, the expression 1 is
 284 contained in $\gamma(x_i, [4, 5))$, and 0 is contained in the remaining sets γ maps x_i to.

285 Finally, we prove Tr^+ overapproximates Tr .

286 **Lemma 11.** *For every distributed signal (S, \rightsquigarrow) , we have $\text{Tr}(S, \rightsquigarrow) \subseteq \text{Tr}^+(S, \rightsquigarrow)$.*

287 5 Monitoring Algorithm

288 In this section, given a distributed signal (S, \rightsquigarrow) , we describe an algorithm to
 289 compute $[(S, \rightsquigarrow) \models \varphi]_+$. The algorithm makes use of the function γ defined
 290 in Section 4 without explicitly computing $\text{Tr}^+(S, \rightsquigarrow)$. To achieve this, we first
 291 describe the notion of *asynchronous product* of value expressions to capture po-
 292 tential interleavings within segments. We continue with the evaluation of *untimed*
 293 *operators* and then *timed operators*. Finally, we conclude with putting all these
 294 together to compute the *semantics* of STL^+ and discuss an efficient implemen-
 295 tation of the monitoring algorithm using *bit vectors* to represent and manipulate
 296 sets of boolean value expressions.

297 *Remark 12.* For the sake of convenience, we focus on boolean signals for the rest
 298 of the section. Note that asynchronous products and the algorithm to compute
 299 $[(S, \rightsquigarrow) \models \varphi]_+$ can be extended to value expressions over arbitrary finite alphabets, e.g., encoding real-valued signals. This allows us to express more complex
 300 properties where atomic propositions can be functions of real-valued signals.
 301

302 **Asynchronous Products** Consider the value expressions $u_1 = 0 \cdot 1$ and
 303 $u_2 = 1 \cdot 0$ encoding the behaviors of two signals within a segment. Due to partial
 304 asynchrony, the behaviors within segments can be seen as completely asyn-
 305 chronous. To capture the potential interleavings of these behaviors, we consider
 306 how the values in u_1 and u_2 can align. In particular, there are three potential
 307 alignments: (i) the rising edge of u_1 happens before the falling edge of u_2 , (ii)
 308 the falling edge of u_2 happens before the rising edge of u_1 , and (iii) the two
 309 edges happen simultaneously. We respectively represent these with the tuples
 310 $(011, 110)$, $(001, 100)$, and $(01, 10)$ where the first component encodes u_1 and
 311 the second u_2 . Formally, given two value expressions u_1 and u_2 , we define their
 312 *asynchronous product* as follows:

$$u_1 \otimes u_2 = \{\text{destutter}(v_1, v_2) \mid v_i \in \text{stutter}_k(u_i), k = |u_1| + |u_2| - 1, i \in \{1, 2\}\}$$

313 Moreover, given two sets L_1 and L_2 of value expressions, we define the following:

$$L_1 \otimes L_2 = \{u_1 \otimes u_2 \mid u_1 \in L_1, u_2 \in L_2\}$$

314 Asynchronous products of value expressions allow us to lift value expressions
 315 to satisfaction signals of formulas.

316 *Example 13.* Recall the distributed signal (S, \rightsquigarrow) in Example 8 and its γ function
 317 given in Figure 2b. Suppose we want to compute the value expressions encoding
 318 the satisfaction of $x_1 \wedge x_2$ in the segment $[1, 3)$. We can achieve this by first com-
 319 puting the asynchronous product $\gamma(x_1, [3, 4)) \otimes \gamma(x_2, [3, 4))$, and then computing
 320 the bitwise conjunction of each pair in the set. For example, considering the ex-
 321 pression $0 \cdot 1 \cdot 0$ for x_1 and $0 \cdot 1$ for x_2 , the product contains the pair $(010, 011)$.
 322 Taking the bitwise conjunction of this pair gives us the expression $0 \cdot 1 \cdot 0$ as a
 323 potential behavior for the satisfaction of $x_1 \wedge x_2$ in this segment.

324 **Untimed Operations** As hinted in Example 13, to compute the semantics, we
 325 apply bitwise operations on value expressions and their asynchronous products
 326 to transform them into encodings of satisfaction signals of formulas. Consider
 327 the distributed signal (S, \rightsquigarrow) in Example 8 and suppose we want to compute
 328 $[(S, \rightsquigarrow) \models \Diamond(x_1 \wedge x_2)]_+$. To achieve this, we first compute for each segment in
 329 G_S the set of value expressions for the satisfaction of $x_1 \wedge x_2$, and then from
 330 these compute that of $\Diamond(x_1 \wedge x_2)$. This compositional approach allows us to
 331 evaluate arbitrary STL⁺ formulas.

332 First, we define bitwise operations on boolean value expressions encoding
 333 atomic propositions. Then, we use these to evaluate (untimed) STL formulas
 334 over sets of value expressions.

Let u and v be boolean value expressions of length ℓ . We denote by $u \& v$ the bitwise-and operation, by $u \mid v$ the bitwise-or, and by $\sim u$ the bitwise-negation. In addition, we define the *bitwise strong until* operator as follows:

$$u \mathbf{U}^0 v = \left(\max_{i \leq j \leq \ell} \left(\min \left(v[j], \min_{i \leq k \leq j} u[k] \right) \right) \right)_{1 \leq i \leq \ell}$$

As usual, we derive *bitwise eventually* as $\mathbf{E}u = 1^\ell \mathbf{U}^0 u$, *bitwise always* as $\mathbf{A}u = \sim(\mathbf{E}\sim u)$, and *bitwise weak until* as $u \mathbf{U}^1 v = (u \mathbf{U}^0 v) \mid (\mathbf{A}u)$. The distinction between \mathbf{U}^0 and \mathbf{U}^1 will be useful later when we evaluate a formula segment by segment. We remark that the definitions of these operators coincide with the robustness semantics of (discrete time) STL. Finally, note that the output of these operations is a value expression of length ℓ . For example, if $u = 010$, we have $\mathbf{E}u = 110$ and $\mathbf{A}u = 000$.

Let (S, \rightsquigarrow) be a distributed signal. Consider an atomic proposition $p \in \mathbf{AP}$ encoded as $x_p \in S$ and let φ_1, φ_2 be two STL formulas. We define the evaluation of untimed formulas with respect to (S, \rightsquigarrow) and a segment $I \in G_S$ inductively:

$$\begin{aligned} \llbracket (S, \rightsquigarrow), I \models p \rrbracket &= \gamma(x_p, I) \\ \llbracket (S, \rightsquigarrow), I \models \neg \varphi_1 \rrbracket &= \{\sim u \mid u \in \llbracket (S, \rightsquigarrow), I \models \varphi_1 \rrbracket\} \\ \llbracket (S, \rightsquigarrow), I \models \varphi_1 \wedge \varphi_2 \rrbracket &= \{u_1 \& u_2 \mid (u_1, u_2) \in \llbracket (S, \rightsquigarrow), I \models \varphi_1 \rrbracket \otimes \llbracket (S, \rightsquigarrow), I \models \varphi_2 \rrbracket\} \\ \llbracket (S, \rightsquigarrow), I \models \varphi_1 \mathbf{U} \varphi_2 \rrbracket &= \{u_1 \mathbf{U}^a u_2 \mid (u_1, u_2) \in \llbracket (S, \rightsquigarrow), I \models \varphi_1 \rrbracket \otimes \llbracket (S, \rightsquigarrow), I \models \varphi_2 \rrbracket, \\ &\quad a \in \text{first}(\llbracket (S, \rightsquigarrow), I' \models \varphi_1 \mathbf{U} \varphi_2 \rrbracket)\} \end{aligned}$$

where I' is the segment that follows I in G_S , if it exists. For completeness, for every formula φ we define $\llbracket (S, \rightsquigarrow), I' \models \varphi \rrbracket = \{0\}$ when $I' \notin G_S$. When I is the first segment in G_S , we simply write $\llbracket (S, \rightsquigarrow) \models \varphi \rrbracket$. Similarly as above, we can use the standard derived operators to compute the corresponding sets of value expressions. Intuitively, for a given formula and a segment, the evaluation above produces a set of value expressions encoding the formula's satisfaction within the segment.

Example 14. Recall the distributed signal (S, \rightsquigarrow) in Example 8 and its γ function given in Figure 2b. Suppose we want to compute $\llbracket (S, \rightsquigarrow), [5, 7] \models \Diamond(x_1 \wedge x_2) \rrbracket$. First, we compute $\llbracket (S, \rightsquigarrow), [5, 7] \models x_1 \wedge x_2 \rrbracket$ by computing the bitwise conjunction over the asynchronous product $\gamma(x_1, [5, 7]) \otimes \gamma(x_2, [5, 7])$ and destuttering. For example, since $010 \in \gamma(x_1, [5, 7])$ and $01 \in \gamma(x_2, [5, 7])$, the pair $(0010, 0111)$ is in the product, whose conjunction gives us 010 after destuttering. Repeating this for the rest, we obtain $\llbracket (S, \rightsquigarrow), [5, 7] \models x_1 \wedge x_2 \rrbracket = \{0, 01, 010, 1, 10\}$. Finally, we compute $\llbracket (S, \rightsquigarrow), [5, 7] \models \Diamond(x_1 \wedge x_2) \rrbracket$ by applying each expression in $\llbracket (S, \rightsquigarrow), [5, 7] \models x_1 \wedge x_2 \rrbracket$ the bitwise eventually operator and destuttering. The resulting set $\{0, 1, 10\}$ encodes the satisfaction signal of $\Diamond(x_1 \wedge x_2)$ in $[5, 7]$. Note that we do not need to consider the evaluation of the next segment for the eventually operator since $\llbracket (S, \rightsquigarrow), [7, 8] \models x_1 \wedge x_2 \rrbracket = \{0\}$.

Timed Operations Handling timed operations requires a closer inspection as value expressions are untimed by definition. We address this issue by considering how a given evaluation interval relates with a given segmentation. For example, take a segmentation $G_S = \{[0, 4), [4, 6), [6, 10)\}$ and an evaluation interval $J = [0, 5)$. Suppose we are interested in how a signal $x \in S$ behaves with respect to J over the first segment $I = [0, 4)$. First, to see how J relates with G_S with respect to $I = [0, 4)$, we “slide” the interval J over $I \oplus J = [0, 9)$ and consider the different ways it intersects the segments in G_S . Initially, J covers the entire segment $[0, 4)$ and the beginning of $[4, 6)$, for which the potential behaviors of x are captured by the set $\gamma(x, [0, 4)) \cdot \text{prefix}(\gamma(x, [4, 6)))$. Now, if we slide the window and take $J' = [3, 7)$, the window covers the ending of $[0, 4)$, the entire $[4, 6)$, and the beginning of $[6, 10)$, for which the potential behaviors are captured by the set $\text{suffix}(\gamma(x, [0, 4))) \cdot \gamma(x, [4, 6)) \cdot \text{prefix}(\gamma(x, [6, 9)))$. We call these sets the *profiles* of J and J' with respect to (S, \rightsquigarrow) , x , and I .

Let (S, \rightsquigarrow) be a distributed signal, $I \in G_S$ be a segment, and φ be an STL formula. Let us introduce the notation we use in the definition below. First, we abbreviate the set $\llbracket (S, \rightsquigarrow), I \models \varphi \rrbracket$ of value expressions as $\tau_{\varphi, I}$. Second, given an interval K , we respectively denote by l_K and r_K its left and right end points. Third, recall that we denote by F the set of end points of G_S (see Section 4). Given an interval J , we define the *profile* of J with respect to (S, \rightsquigarrow) , φ , and I as follows.

$$\text{profile}((S, \rightsquigarrow), \varphi, I, J) = \begin{cases} \text{prefix}(\tau_{\varphi, I}) & \text{if } l_I = l_J \wedge r_I > r_J \\ \text{infix}(\tau_{\varphi, I}) & \text{if } l_I < l_J \wedge r_I > r_J \\ \tau_{\varphi, I} \cdot \kappa(\varphi, I, J) & \text{if } l_I = l_J \wedge r_I \leq r_J \wedge r_J \in F \setminus J \\ \tau_{\varphi, I} \cdot \kappa(\varphi, I, J) \cdot \text{first}(\tau_{\varphi, I'}) & \text{if } l_I = l_J \wedge r_I \leq r_J \wedge r_J \in F \cap J \\ \tau_{\varphi, I} \cdot \kappa(\varphi, I, J) \cdot \text{prefix}(\tau_{\varphi, I'}) & \text{if } l_I = l_J \wedge r_I \leq r_J \wedge r_J \notin F \\ \text{suffix}(\tau_{\varphi, I}) \cdot \kappa(\varphi, I, J) & \text{if } l_I < l_J < r_I \leq r_J \wedge r_J \in F \setminus J \\ \text{suffix}(\tau_{\varphi, I}) \cdot \kappa(\varphi, I, J) \cdot \text{first}(\tau_{\varphi, I'}) & \text{if } l_I < l_J < r_I \leq r_J \wedge r_J \in F \cap J \\ \text{suffix}(\tau_{\varphi, I}) \cdot \kappa(\varphi, I, J) \cdot \text{prefix}(\tau_{\varphi, I'}) & \text{if } l_I < l_J < r_I \leq r_J \wedge r_J \notin F \\ \{\epsilon\} & \text{otherwise} \end{cases}$$

where we assume J is trimmed to fit the temporal domain of S and $I' \in G_S$ is such that $r_J \in I'$. Moreover, $\kappa(\varphi, I, J)$ is the concatenation $\tau_{\varphi, I_1} \cdot \dots \cdot \tau_{\varphi, I_m}$ such that I, I_1, \dots, I_m, I' are consecutive segments in G_S . If I_1, \dots, I_m do not exist, we let $\kappa(\varphi, I, J) = \{\epsilon\}$. Note that the last case happens when $I \cap J$ is empty. We now formalize the intuitive approach of “sliding” J over the segmentation to obtain the various profiles it produces as follows.

$$\text{pfs}((S, \rightsquigarrow), \varphi, I, J) = \{\text{destutter}(\text{profile}((S, \rightsquigarrow), \varphi, I, J')) \mid J' \subseteq I \oplus J, J' \sim J\}$$

where $J' \sim J$ holds when $|J'| = |J|$ and J' contains an end point (left or right) iff J does so. Note that although infinitely many intervals J' satisfy the conditions given above (due to denseness of time), the set defined by **pfs** is finite. We demonstrate this and the computation of **pfs** in Example 15 and Figure 3.

Example 15. Recall the distributed signal (S, \rightsquigarrow) in Example 8 and its γ function given in Figure 2b. We demonstrate the computation of **pfs** $((S, \rightsquigarrow), x_1, [1, 3), [0, 1))$.

Intuitively, sliding the interval $[0, 1)$ over the window $[1, 3) \oplus [0, 1)$ (as shown in Figure 3) gives us the following sets:

$$\begin{aligned}
 P_1 &= \text{destutter}(\text{prefix}(\gamma(x_1, [1, 3)))) = \{0, 01, 1\} \\
 P_2 &= \text{destutter}(\text{infix}(\gamma(x_1, [1, 3)))) = \{0, 01, 1\} \\
 P_3 &= \text{destutter}(\text{suffix}(\gamma(x_1, [1, 3)))) = \{0, 01, 1\} \\
 P_4 &= \text{destutter}(\text{suffix}(\gamma(x_1, [1, 3))) \cdot \text{prefix}(\gamma(x_1, [3, 4)))) \\
 &= \{0, 01, 010, 0101, 01010, 1, 10, 101, 1010\}
 \end{aligned}$$

Therefore, we obtain $\text{pfs}((S, \rightsquigarrow), x_1, [1, 3), [0, 1)) = \{P_1, P_2, P_3, P_4\}$. This set overapproximates the potential behaviors of x_1 , for all $t \in [1, 3)$, in the interval $t \oplus [0, 1)$.




Fig. 3. The profiles of $J = [0, 1)$ with respect to $x_1 \in S$ of Example 8. The four representative intervals of each profile is shown with solid black lines below the tabular representation of γ for x_1 .

Let φ_1 and φ_2 be two STL formulas. Intuitively, once we have the profiles of a given interval J with respect to φ_1 and φ_2 , we can evaluate the corresponding untimed formulas on the product of these profiles and concatenate them. Formally, we handle the evaluation of timed formulas inductively as follows.

$$\llbracket (S, \rightsquigarrow), I \models \varphi_1 \mathcal{U}_J \varphi_2 \rrbracket = \text{destutter}(\{u_1 \mathcal{U}^0 u_2 \mid (u_1, u_2) \in P_1 \otimes Q_1\} \cdot \dots \cdot \{u_1 \mathcal{U}^0 u_2 \mid (u_1, u_2) \in P_k \otimes Q_k\})$$

where $\text{pfs}((S, \rightsquigarrow), \varphi_1, I, J) = \{P_1, \dots, P_k\}$ and $\text{pfs}((S, \rightsquigarrow), \varphi_2, I, J) = \{Q_1, \dots, Q_k\}$ such that the intervals producing P_i and Q_i respectively start before those producing P_{i+1} and Q_{i+1} for all $1 \leq i < k$.

Example 16. Let (S, \rightsquigarrow) be as in Example 8 and its γ function as given in Figure 2b. We demonstrate the evaluation of the timed formula $\Diamond_{[0,1)} x_1$ over the segment $[1, 3)$. Recall from Example 15 the set $\text{pfs}((S, \rightsquigarrow), x_1, [1, 3), [0, 1)) = \{P_1, P_2, P_3, P_4\}$ of profiles. First, we apply the bitwise eventually operator to each value expression in each of these profiles separately: $\{\text{Eu} \mid u \in P_1\} = \{0, 1\}$, $\{\text{Eu} \mid u \in P_2\} = \{0, 1\}$, $\{\text{Eu} \mid u \in P_3\} = \{0, 1\}$, and $\{\text{Eu} \mid u \in P_4\} = \{0, 10, 1\}$. Then, we concatenate these sets and destutter to obtain the following:

$$\llbracket (S, \rightsquigarrow), [1, 3) \models \Diamond_{[0,1)} x_1 \rrbracket = \{0, 01, 010, 0101, 01010, 1, 10, 101, 1010\}$$

419 **Computing the Semantics of STL⁺** Putting it all together, given a dis-
 420 tributed signal (S, \rightsquigarrow) and an STL⁺ formula φ , we can compute $[(S, \rightsquigarrow) \models \varphi]_+$
 421 thanks to the following theorem.

422 **Theorem 17.** *For every distributed signal (S, \rightsquigarrow) , we have $[(S, \rightsquigarrow) \models \varphi]_+ = \top$*
 423 *(resp. \perp , $?$) iff $\text{first}(\llbracket (S, \rightsquigarrow) \models \varphi \rrbracket) = \{1\}$ (resp. $\{0\}, \{0, 1\})$.*

424 **Sets of Boolean Value Expressions as Bit Vectors** Evidently, asynchronous
 425 products are expensive to compute. Our implementation of the algorithm we de-
 426 scribe in this section relies on the following observation: Sets of boolean value
 427 expressions and their operations can be efficiently implemented through bit vec-
 428 tors. Intuitively, to represent such a set, we can encode each element using its first
 429 bit and its length since value expressions are boolean and always destuttered.
 430 Moreover, to evaluate untimed operations on such sets, we only need to know
 431 the maximal lengths of the four possible types of expressions ($0 \dots 0$, $0 \dots 1$,
 432 $1 \dots 0$, and $1 \dots 1$) and whether the set contains 0 or 1 (to handle some edge
 433 cases). This is because value expressions corresponding to same segments can be
 434 seen as completely asynchronous and the possible interleavings obtained from
 435 shorter expressions can be obtained from longer ones. This approach enables,
 436 for example, an algorithm for conjunction of sets of value expressions that runs
 437 in $O(|u| + |v|)$ time where u and v are the longest expressions in the two sets.
 438 Note that the same idea also applies to untimed temporal operators.

439 6 Experimental Evaluation

440 **TODO**

441 7 Conclusion

442 **TODO**

443 References

- 444 1. Maler, O., Nickovic, D.: Monitoring properties of analog and mixed-signal circuits.
 445 Int. J. Softw. Tools Technol. Transf. **15**(3), 247–268 (2013). [https://doi.org/10.](https://doi.org/10.1007/s10009-012-0247-9)
 446 [1007/s10009-012-0247-9](https://doi.org/10.1007/s10009-012-0247-9)
- 447 2. Momtaz, A., Abbas, H., Bonakdarpour, B.: Monitoring signal temporal logic in
 448 distributed cyber-physical systems. In: Mitra, S., Venkatasubramanian, N., Dubey,
 449 A., Feng, L., Ghasemi, M., Sprinkle, J. (eds.) Proceedings of the ACM/IEEE 14th
 450 International Conference on Cyber-Physical Systems, ICCPS 2023, (with CPS-IoT
 451 Week 2023), San Antonio, TX, USA, May 9–12, 2023. pp. 154–165. ACM (2023).
 452 <https://doi.org/10.1145/3576841.3585937>

453 Appendix

454 *Proof (Proof of Theorem 6).* Let φ be an STL formula and (S, \rightsquigarrow) be a dis-
 455 tributed signal. Assume $[(S, \rightsquigarrow) \models \varphi]_+ = \top$. We want to show that $[(S, \rightsquigarrow) \models$
 456 $\varphi] = \top$. Expanding the definition of $[(S, \rightsquigarrow) \models \varphi]_+ = \top$, we have $w \models \varphi$ for all
 457 $w \in \text{Tr}^+(S, \rightsquigarrow)$. By Lemma 11, we have $\text{Tr}(S, \rightsquigarrow) \subseteq \text{Tr}^+(S, \rightsquigarrow)$. Then, it holds
 458 that $w \models \varphi$ for all $w \in \text{Tr}(S, \rightsquigarrow)$. Therefore, $[(S, \rightsquigarrow) \models \varphi] = \top$ by definition.
 459 The case of $[(S, \rightsquigarrow) \models \varphi]_+ = \perp$ follows from the same arguments.

460 *Proof (Proof of Lemma 11).* Let (S, \rightsquigarrow) be a distributed signal where $S =$
 461 (x_1, \dots, x_n) . Let $w = (y_1, \dots, y_n) \in \text{Tr}(S, \rightsquigarrow)$ be a trace. We want to show
 462 that $w \in \text{Tr}^+(S, \rightsquigarrow)$. First, recall the following:

$$\text{Tr}^+(S, \rightsquigarrow) = \{(x'_1, \dots, x'_n) \mid x'_i \text{ is consistent with } x_i \text{ for all } 1 \leq i \leq n\}$$

463 To achieve our goal, we need to show that x_i and y_i are consistent for all $1 \leq i \leq$
 464 n . Let $1 \leq i \leq n$ be arbitrary, and let $I = [t_0, s)$ be a segment in G_S . Moreover,
 465 let $(t_1, y_i(t_1)), \dots, (t_\ell, y_i(t_\ell))$ be the edges of y_i in segment I with $t_j < t_{j+1}$ for
 466 all $1 \leq j < \ell$. We want to show that the expression $y_i(t_0) \cdot y_i(t_1) \cdot \dots \cdot y_i(t_\ell)$
 467 belongs to $\gamma(x_i, I)$. We sketch the proof below.

468 Note that w can be seen as a trace obtained through an ε -retiming of S (see
 469 [2, Section 4.2]). It is then clear that the timestamps of the events (and thus
 470 edges) of x_i are mapped to clock values that are less than ε away. In particular,
 471 $|t - c_i^{-1}(t)| < \varepsilon$ for all $t \in \{t_0, t_1, \dots, t_\ell\}$, where $c_i^{-1}(t)$ is the local clock value of
 472 x_i that is mapped to t . Since y_i has ℓ edges in I , it holds that x_i has at least ℓ
 473 edges in $[t_0 - \varepsilon, s + \varepsilon)$. Since I is a segment in G_S , there are ℓ of these that are
 474 consecutive such that the intersection of their uncertainty regions contain (t_0, s) ,
 475 i.e., $(t_0, s) \subseteq \bigcap_{1 \leq j \leq \ell} (\theta_{\text{lo}}(t'_j), \theta_{\text{hi}}(t'_j))$ where $t'_j = c_i^{-1}(t_j)$ is the corresponding
 476 timestamp in x_i for all $0 \leq j \leq \ell$. In particular, note that $y_i(t_j) = x_i(t'_j)$ for all
 477 $0 \leq j \leq \ell$.

478 Now, notice that, by definition, $\gamma(x_i, I)$ takes into account every edge of x_i
 479 whose uncertainty region has a nonempty intersection with I . Moreover, thanks
 480 to concatenation in its definition, it contains a value expression in which all ℓ
 481 edges occur and their order is preserved. We conclude that $x_i(t'_0) \cdot x_i(t'_1) \cdot \dots \cdot$
 482 $x_i(t'_\ell) \in \gamma(x_i, I)$.

483 *Proof (Proof of Theorem 17).* **TODO**