

# Approximate Distributed Monitoring under Partial Synchrony: Balancing Speed with Accuracy

Author(s)

Affiliation(s)

**Abstract.** In distributed systems with processes that do not share a global clock, *partial synchrony* is achieved by clock synchronization that guarantees bounded clock skew among all applications. Existing solutions for distributed runtime verification under partial synchrony against temporal logic specifications are exact but suffer from significant computational overhead. In this paper, we propose an *approximate* distributed monitoring algorithm for Signal Temporal Logic (STL) that mitigates this issue by abstracting away potential interleaving behaviors. This conservative abstraction enables a significant speedup of the distributed monitors, albeit with a trade-off in accuracy. We address this trade-off with a methodology that combines our approximate monitor with its exact counterpart, resulting in enhanced monitoring efficiency without sacrificing precision. We validate our approach with multiple experiments, showcasing its effectiveness and efficacy on both a real-world application and synthetic examples.

You can leave notes using the command `\firstName{note}`.  
Page limit: 16 + 4 (appendix)

## 1 Introduction

*Distributed systems* are networks of independent agents that work together to achieve a common objective. Distributed systems are everywhere around us and come in many different forms. For example, cloud computing uses distribution of resources and services over the internet to offer to their users a scalable infrastructure with transparent on-demand access to computing power and storage. Swarms of drones represent another family of distributed systems where individual drones collaborate to accomplish tasks like surveillance, search and rescue, or package delivery. While each drone operates independently, it also communicates and coordinates with others to successfully achieve their common objectives. The individual agents in a distributed system typically do not share a global clock. To coordinate actions across multiple agents, clock synchronization is often needed. While perfect clock synchronization is impractical due to network latency and node failures, algorithms such as the Network Time Protocol (NTP) allow agents

36 to maintain a *bounded skew* between the synchronized clocks. In that case, we  
 37 say that a distributed system has *partial synchrony*.

38 Formal verification of distributed system is a notoriously hard problem, due  
 39 to the combinatorial explosion of all possible interleavings in the behaviors col-  
 40 lected from individual agents. *Runtime verification (RV)* provides a more prag-  
 41 matic approach, in which a monitor observes a behavior of a distributed sys-  
 42 tem and checks its correctness against a formal specification. The problem of  
 43 distributed RV under partial synchrony assumption has been studied for Lin-  
 44 ear Temporal Logic (LTL) and Signal Temporal Logic (STL) specification lan-  
 45 guages. The proposed solutions use Satisfiability-Modulo-Theory (SMT) solving  
 46 to provide sound and complete distributed monitoring procedures. Although dis-  
 47 tributed RV monitors consume only a single distributed behavior at a time, this  
 48 behavior can nevertheless have an excessive number of possible interleavings.  
 49 Hence, the exact distributed monitors from the literature can still suffer from  
 50 significant computational overhead.

51 To mitigate this issue, we present an approach for *approximate* RV of STL  
 52 specifications under partial synchrony. In essence, we abstract away potential  
 53 interleavings in distributed behaviors in a conservative manner, resulting in an  
 54 effective over-approximation of global behaviors. This abstraction simplifies the  
 55 representation of distributed behaviors and the monitoring operations required  
 56 to evaluate temporal specifications. There is an inevitable trade-off in approxi-  
 57 mate RV – gains in the monitoring speed-up may result in reduced accuracy. For  
 58 some applications, reduced accuracy may not be acceptable. Therefore, we pro-  
 59 pose a methodology that combines our approximate monitors with their exact  
 60 counterparts, with the aim to benefit from the enhanced monitoring efficiency  
 61 without sacrificing precision. We implemented our approach in a prototype tool  
 62 and performed thorough evaluations on both synthetic and real-world case stud-  
 63 ies. We first demonstrated that our approximate monitors achieve speed-ups of  
 64 several orders of magnitudes compared to the exact SMT-based distributed RV  
 65 solution. We empirically characterized the classes of specifications and behaviors  
 66 for which our approximate monitoring approach achieves good precision. We fi-  
 67 nally showed that by combining exact and approximate distributed RV, there is  
 68 still a significant efficiency gain on average without the sacrifice of the precision,  
 69 even in cases where approximate monitors have low accuracy.

## 70 2 Preliminaries

71 We denote by  $\mathbb{B} = \{\top, \perp\}$  the set of Booleans,  $\mathbb{R}$  the set of reals,  $\mathbb{R}_{\geq 0}$  the set of  
 72 nonnegative reals, and  $\mathbb{R}_{> 0}$  the set of positive reals. An interval  $I \subseteq \mathbb{R}$  of reals  
 73 with the end points  $a < b$  has length  $|b - a|$ .

74 Let  $\Sigma$  be a finite *alphabet*. We denote by  $\Sigma^*$  the set of finite words over  
 75  $\Sigma$  and by  $\epsilon$  the empty word. For  $u \in \Sigma^*$ , we respectively write  $\text{prefix}(u)$  and  
 76  $\text{suffix}(u)$  for the sets of prefixes and suffixes of  $u$ . We also let  $\text{infix}(u) = \{v \in$   
 77  $\Sigma^* \mid \exists x, y \in \Sigma^* : u = xvy\}$ . For a nonempty word  $u \in \Sigma^*$  and  $1 \leq i \leq |u|$ ,  
 78 we denote by  $u[i]$  the  $i$ th letter of  $u$ , by  $u[..i]$  the prefix of  $u$  of length  $i$ , and by

79  $u[i..]$  the suffix of  $u$  of length  $|u| - i + 1$ . Given  $u \in \Sigma^*$  and  $\ell \geq 1$ , we denote by  
 80  $u^\ell$  the word obtained by concatenating  $u$  by itself  $\ell - 1$  times. Moreover, given  
 81  $L \subseteq \Sigma^*$ , we define  $\text{first}(L) = \{u[0] \mid u \in L\}$ . For sets  $L_1, L_2 \subseteq \Sigma^*$  of words, we  
 82 let  $L_1 \cdot L_2 = \{u \cdot v \mid u \in L_1, v \in L_2\}$ . For tuples  $(u_1, \dots, u_m)$  and  $(v_1, \dots, v_m)$  of  
 83 words, we let  $(u_1, \dots, u_m) \cdot (v_1, \dots, v_m) = (u_1v_1, \dots, u_mv_m)$ .

84 We define the function  $\text{destutter} : \Sigma^* \rightarrow \Sigma^*$  inductively as follows. For all  
 85  $\sigma \in \Sigma \cup \{\epsilon\}$ , let  $\text{destutter}(\sigma) = \sigma$ . For all  $u \in \Sigma^*$  such that  $u = \sigma_1\sigma_2v$  for  
 86 some  $\sigma_1, \sigma_2 \in \Sigma$  and  $v \in \Sigma^*$ , let (i)  $\text{destutter}(u) = \text{destutter}(\sigma_2v)$  if  $\sigma_1 = \sigma_2$ ,  
 87 and (ii)  $\text{destutter}(u) = \sigma_1 \cdot \text{destutter}(\sigma_2v)$  otherwise. By extension, for a set  
 88  $L \subseteq \Sigma^*$  of finite words, we write  $\text{destutter}(L) = \{\text{destutter}(u) \mid u \in L\}$ . Given  
 89 a tuple  $(u_1, \dots, u_m) = (\sigma_{1,1}\sigma_{1,2}v_1, \dots, \sigma_{m,1}\sigma_{m,2}v_m)$  of finite words of the same  
 90 length, we define  $\text{destutter}(u_1, \dots, u_m)$  as expected: (i)  $\text{destutter}(u_1, \dots, u_m) =$   
 91  $\text{destutter}(\sigma_{1,2}v_1, \dots, \sigma_{m,2}v_m)$  if  $\sigma_{i,1} = \sigma_{i,2}$  for all  $1 \leq i \leq m$ , and (ii)  $\text{destutter}(u_1, \dots, u_m) =$   
 92  $(\sigma_{1,1}, \dots, \sigma_{m,1}) \cdot \text{destutter}(\sigma_{1,2}v_1, \dots, \sigma_{m,2}v_m)$  otherwise.

93 Moreover, given an integer  $k \geq 0$ , we define  $\text{stutter}_k : \Sigma^* \rightarrow \Sigma^*$  such  
 94 that  $\text{stutter}_k(u) = \{v \in \Sigma^* \mid |v| = k \wedge \text{destutter}(v) = \text{destutter}(u)\}$  if  $k \geq$   
 95  $|\text{destutter}(u)|$ , and  $\text{stutter}_k(u) = \emptyset$  otherwise.

96 **Signal Temporal Logic (STL) [1].** Let  $A, B \subset \mathbb{R}$ . A function  $f : A \rightarrow B$   
 97 is *right-continuous* iff  $\lim_{a \rightarrow c^+} f(a) = f(c)$  for all  $c \in A$ , and *non-Zeno* iff for  
 98 every bounded interval  $I \subseteq A$  there are finitely many  $a \in I$  such that  $f$  is not  
 99 continuous at  $a$ . A *signal* is a right-continuous, non-Zeno, piecewise-constant  
 100 function  $x : [0, d) \rightarrow \mathbb{R}$  where  $d \in \mathbb{R}_{>0}$  is the duration of  $x$  and  $[0, d)$  is its  
 101 temporal domain. Let  $x : [0, d) \rightarrow \mathbb{R}$  be a signal. An *event* of  $x$  is a pair  $(t, x(t))$   
 102 where  $t \in [0, d)$ . An *edge* of  $x$  is an event  $(t, x(t))$  such that  $\lim_{s \rightarrow t^-} x(s) \neq$   
 103  $\lim_{s \rightarrow t^+} x(s)$ . In particular, an edge is *rising* if  $\lim_{s \rightarrow t^-} x(s) < \lim_{s \rightarrow t^+} x(s)$ , and  
 104 it is *falling* otherwise. A signal  $x : [0, d) \rightarrow \mathbb{R}$  can be represented finitely by its  
 105 initial value and edges: if  $x$  has  $m$  edges, then  $x = (t_0, v_0)(t_1, v_1) \dots (t_m, v_m)$   
 106 such that  $t_0 = 0$ ,  $t_{i-1} < t_i$ , and  $(t_i, v_i)$  is an edge of  $x$  for all  $1 \leq i \leq m$ .

107 Let  $\text{AP}$  be a set of *atomic propositions*. The syntax is given by the following  
 108 grammar where  $p \in \text{AP}$  and  $I \subseteq \mathbb{R}_{\geq 0}$  is an interval.

$$\varphi := p \mid \neg\varphi \mid \varphi \wedge \varphi \mid \varphi \mathcal{U}_I \varphi$$

109 A *trace*  $w = (x_1, \dots, x_n)$  is a finite vector of signals. We express atomic  
 110 propositions as functions of trace values at a time point  $t$ , i.e., a proposition  
 111  $p \in \text{AP}$  over a trace  $w = (x_1, \dots, x_n)$  is defined as  $f_p(x_1(t), \dots, x_n(t)) > 0$   
 112 where  $f_p : \mathbb{R}^n \rightarrow \mathbb{R}$  is a function. Given intervals  $I, J \subseteq \mathbb{R}_{\geq 0}$ , we define  $I \oplus J =$   
 113  $\{i + j \mid i \in I \wedge j \in J\}$ , and we simply write  $t$  for the singleton set  $\{t\}$ .

114 Below, we recall the finite-trace qualitative semantics of STL defined over  $\mathbb{B}$   
 115 [1]. Let  $d \in \mathbb{R}_{>0}$  and  $w = (x_1, \dots, x_n)$  with  $x_i : [0, d) \rightarrow \mathbb{R}$  for all  $1 \leq i \leq n$ . Let  
 116  $\varphi_1, \varphi_2$  be STL formulas and let  $t \in [0, d)$ .

$$\begin{aligned}
(w, t) \models p &\iff f_p(x_1(t), \dots, x_n(t)) > 0 \\
(w, t) \models \neg \varphi_1 &\iff \overline{(w, t) \models \varphi_1} \\
(w, t) \models \varphi_1 \wedge \varphi_2 &\iff (w, t) \models \varphi_1 \wedge (w, t) \models \varphi_2 \\
(w, t) \models \varphi_1 \mathcal{U}_I \varphi_2 &\iff \exists t' \in (t \oplus I) \cap [0, d) : \\
&\quad (w, t') \models \varphi_2 \wedge \forall t'' \in (t, t') : (w, t'') \models \varphi_1
\end{aligned}$$

117 We simply write  $w \models \varphi$  for  $(w, 0) \models \varphi$ . We additionally use the following  
118 standard abbreviations: **false** =  $p \wedge \neg p$ , **true** =  $\neg \mathbf{false}$ ,  $\varphi_1 \vee \varphi_2 = \neg(\neg \varphi_1 \wedge$   
119  $\neg \varphi_2)$ ,  $\Diamond_I \varphi = \mathbf{true} \mathcal{U}_I \varphi$ , and  $\Box_I \varphi = \neg \Diamond_I \neg \varphi$ . Moreover, the untimed temporal  
120 operators are defined through their timed counterparts on the interval  $[0, \infty)$ ,  
121 e.g.,  $\varphi_1 \mathcal{U} \varphi_2 = \varphi_1 \mathcal{U}_{[0, \infty)} \varphi_2$ .

122 **Distributed Semantics of STL [2].** We consider an asynchronous and loosely-  
123 coupled message-passing system of  $n \geq 2$  reliable agents producing a set of  
124 signals  $x_1, \dots, x_n$ , where for some  $d \in \mathbb{R}_{>0}$  we have  $x_i : [0, d) \rightarrow \mathbb{R}$  for all  
125  $1 \leq i \leq n$ . The agents do not share memory or a global clock. Only to formalize  
126 statements, we speak of a *hypothetical* global clock and denote its value by  $T$ .  
127 For local time values, we use the lowercase letters  $t$  and  $s$ .

128 For a signal  $x_i$ , we denote by  $V_i$  the set of its events, by  $E_i^\uparrow$  the set of its  
129 rising edges, and by  $E_i^\downarrow$  that of falling edges. Moreover, we let  $E_i = E_i^\uparrow \cup E_i^\downarrow$ . We  
130 represent the local clock of the  $i$ th agent as an increasing and divergent function  
131  $c_i : \mathbb{R}_{\geq 0} \rightarrow \mathbb{R}_{\geq 0}$  that maps a global time  $T$  to a local time  $c_i(T)$ .

132 We assume that the system is *partially synchronous*: the agents use a clock  
133 synchronization algorithm that guarantees a bounded clock skew with respect  
134 to the global clock, i.e.,  $|c_i(T) - c_j(T)| < \varepsilon$  for all  $1 \leq i, j \leq N$  and  $T \in \mathbb{R}_{\geq 0}$ ,  
135 where  $\varepsilon \in \mathbb{R}_{>0}$  is the maximum clock skew.

136 **Definition 1.** A distributed signal is a pair  $(S, \rightsquigarrow)$ , where  $S = (x_1, \dots, x_n)$  is a  
137 vector of signals and  $\rightsquigarrow$  is the happened-before relation between events in signals  
138 extended with the partial synchrony assumption as follows.

- 139 – For every agent, the events of its signals are totally ordered, i.e., for all  $1 \leq$   
140  $i \leq n$  and all  $(t, x_i(t)), (t', x_i(t')) \in V_i$ , if  $t < t'$  then  $(t, x_i(t)) \rightsquigarrow (t', x_i(t'))$ .
- 141 – Every pair of events whose timestamps are at least  $\varepsilon$  apart is totally ordered,  
142 i.e., for all  $1 \leq i, j \leq n$  and all  $(t, x_i(t)) \in V_i$  and  $(t', x_j(t')) \in V_j$ , if  $t + \varepsilon \leq t'$   
143 then  $(t, x_i(t)) \rightsquigarrow (t', x_j(t'))$ .

144 *Example 2.* **TODO: distributed signal, happened-before relation**

145 **Definition 3.** Let  $(S, \rightsquigarrow)$  be a distributed signal of  $n$  signals, and  $V = \bigcup_{i=1}^n V_i$   
146 be the set of its events. A set  $C \subseteq V$  is a consistent cut iff for every event in  
147  $C$ , all events that happened before it also belong to  $C$ , i.e., for all  $e, e' \in V$ , if  
148  $e \in C$  and  $e' \rightsquigarrow e$ , then  $e' \in C$ .

We denote by  $\mathbb{C}(T)$  the (infinite) set of consistent cuts at global time  $T$ . Given a consistent cut  $C$ , its *frontier*  $\text{front}(C) \subseteq C$  is the set consisting of the last events in  $C$  of each signal, i.e.,  $\text{front}(C) = \bigcup_{i=1}^n \{(t, x_i(t)) \in V_i \cap C \mid \forall t' > t : (t', x_i(t')) \notin V_i \cap C\}$ .

**Definition 4.** A consistent cut flow is a function  $\text{ccf} : \mathbb{R}_{\geq 0} \rightarrow 2^V$  that maps a global clock value  $T$  to the frontier of a consistent cut at time  $T$ , i.e.,  $\text{ccf}(T) \in \{\text{front}(C) \mid C \in \mathbb{C}(T)\}$ .

For all  $T, T' \in \mathbb{R}_{\geq 0}$  and  $1 \leq i \leq n$ , if  $T < T'$ , then for every pair of events  $(c_i(T), x_i(c_i(T))) \in \text{ccf}(T)$  and  $(c_i(T'), x_i(c_i(T')))) \in \text{ccf}(T')$  we have  $(c_i(T), x_i(c_i(T))) \rightsquigarrow (c_i(T'), x_i(c_i(T')))$ . We denote by  $\text{CCF}(S, \rightsquigarrow)$  the set of all consistent cut flows of the distributed signal  $(S, \rightsquigarrow)$ .

*Example 5.* **TODO: consistent cut, frontier, consistent cut flow**

Observe that a consistent cut flow of a distributed signal induces a vector of synchronous signals which can be evaluated using the standard semantics described in Section 2. Let  $(S, \rightsquigarrow)$  be a distributed signal of  $n$  signals  $x_1, \dots, x_n$ . A consistent cut flow  $\text{ccf} \in \text{CCF}(S, \rightsquigarrow)$  yields a trace  $w_{\text{ccf}} = (x'_1, \dots, x'_n)$  on the temporal domain  $[0, d]$  such that  $(c_i(T), x_i(c_i(T))) \in \text{ccf}(T)$  implies  $x'_i(T) = x_i(c_i(T))$  for all  $1 \leq i \leq n$  and  $T \in [0, d]$ . The set of traces of  $(S, \rightsquigarrow)$  is given by  $\text{Tr}(S, \rightsquigarrow) = \{w_{\text{ccf}} \mid \text{ccf} \in \text{CCF}(S, \rightsquigarrow)\}$ .

We define the satisfaction of an STL formula  $\varphi$  by a distributed signal  $(S, \rightsquigarrow)$  over a three-valued domain  $\{\top, \perp, ?\}$ . If the set of synchronous traces  $\text{Tr}(S, \rightsquigarrow)$  defined by a distributed signal  $(S, \rightsquigarrow)$  is contained in the set of traces allowed by the formula  $\varphi$ , then  $(S, \rightsquigarrow)$  satisfies  $\varphi$ . Similarly, if  $\text{Tr}(S, \rightsquigarrow)$  has an empty intersection with the set of traces  $\varphi$  defines, then  $(S, \rightsquigarrow)$  violates  $\varphi$ . Otherwise, the evaluation is inconclusive since some traces satisfy the property and some violate it. Notice that we quantify universally over traces for both satisfaction and violation.

$$[(S, \rightsquigarrow) \models \varphi] = \begin{cases} \top & \text{if } \forall w \in \text{Tr}(S, \rightsquigarrow) : w \models \varphi \\ \perp & \text{if } \forall w \in \text{Tr}(S, \rightsquigarrow) : w \models \neg \varphi \\ ? & \text{otherwise} \end{cases}$$

### 3 Overapproximation of the STL Distributed Semantics

To address the computational overhead in exact distributed monitoring, we define  $\text{STL}^+$ , a variant of STL whose syntax is the same as STL but semantics provide a sound approximation of the STL distributed semantics.  $\text{STL}^+$  is better adapted to distributed monitoring as it overapproximates the set of traces. In particular, given a distributed signal  $(S, \rightsquigarrow)$ ,  $\text{STL}^+$  considers an approximation  $\text{Tr}^+(S, \rightsquigarrow)$  of the set  $\text{Tr}(S, \rightsquigarrow)$  of synchronous traces where  $\text{Tr}(S, \rightsquigarrow) \subseteq \text{Tr}^+(S, \rightsquigarrow)$ . A signal  $(S, \rightsquigarrow)$  satisfies (resp. violates) an  $\text{STL}^+$  formula  $\varphi$  iff all the traces in  $\text{Tr}^+(S, \rightsquigarrow)$  belong to the language of  $\varphi$  (resp.  $\neg \varphi$ ).

$$[(S, \rightsquigarrow) \models \varphi]_+ = \begin{cases} \top & \text{if } \forall w \in \text{Tr}^+(S, \rightsquigarrow) : w \models \varphi \\ \perp & \text{if } \forall w \in \text{Tr}^+(S, \rightsquigarrow) : w \models \neg \varphi \\ ? & \text{otherwise} \end{cases}$$

185 In Sections 4 and 5, we respectively define  $\text{Tr}^+$  and present an algorithm  
 186 to compute the semantics of  $\text{STL}^+$ . We finally prove the correctness of our  
 187 approach.

188 **Theorem 6.** *For every STL formula  $\varphi$  and every distributed signal  $(S, \rightsquigarrow)$ , if*  
 189  *$[(S, \rightsquigarrow) \models \varphi]_+ = \top$  (resp.  $\perp$ ) then  $[(S, \rightsquigarrow) \models \varphi] = \top$  (resp.  $\perp$ ).*

190 Notice that both the distributed semantics of STL and the semantics of  $\text{STL}^+$   
 191 quantify universally over the set of traces for the verdicts  $\top$  and  $\perp$ . Therefore,  
 192 the Theorem 6 holds for the verdicts  $\top$  and  $\perp$ , but not for  $?$ .

## 193 4 Overapproximation of Synchronous Traces

194 In this section, given a distributed signal  $(S, \rightsquigarrow)$ , we describe an overapproxima-  
 195 tion  $\text{Tr}^+(S, \rightsquigarrow)$  of its set  $\text{Tr}(S, \rightsquigarrow)$  of synchronous traces. First, we present the  
 196 notion of *canonical segmentation*, a systematic way of partitioning the temporal  
 197 domain of a given distributed signal to keep track of the partial asynchrony.  
 198 Second, we introduce the notion of *value expressions*, sets of finite words repre-  
 199 senting how a signal behaves in a time interval. Finally, we define  $\text{Tr}^+$  based on  
 200 these notions, and show that it soundly approximates  $\text{Tr}$ .

201 *Remark 7.* We assume boolean signals in this section for convenience. The def-  
 202 initions and results presented here extend to real-valued signals because finite-  
 203 length piecewise-constant signals will only use a finite number of values.


204 **Canonical Segmentation** Consider a boolean signal  $x$  with a rising edge at  
 205 time  $t > \varepsilon$ . Due to clock skew, this edge occurs in the range  $(t - \varepsilon, t + \varepsilon)$   
 206 from the monitor's point of view. This range is called an *uncertainty region*  
 207 because in  $(t - \varepsilon, t + \varepsilon)$  the monitor cannot tell the value of  $x$  precisely, but  
 208 only that it changes from 0 to 1. Formally, given an edge  $(t, x(t))$ , we define  
 209  $\theta_{\text{lo}}(x, t) = \max(0, t - \varepsilon)$  and  $\theta_{\text{hi}}(x, t) = \min(d, t + \varepsilon)$  as the end points of the  
 210 edge's uncertainty region.

211 Given a temporal domain  $I = [0, d] \subset \mathbb{R}_{\geq 0}$ , a *segmentation* of  $I$  is a partition  
 212 of  $I$  into finitely many intervals  $I_1, \dots, I_k$ , called *segments*, of the form  $I_j =$   
 213  $[t_j, t_{j+1})$  such that  $t_j < t_{j+1}$  for all  $1 \leq j \leq k$ . By extension, a segmentation of  
 214 a collection of signals with the same temporal domain  $I$  is a segmentation of  $I$ .

215 Let  $(S, \rightsquigarrow)$  be a distributed signal of  $n$  signals. The *canonical segmentation*  
 216  $G_S$  of  $(S, \rightsquigarrow)$  is the segmentation of  $S$  where the end points of the segments  
 217 coincide with the end points of its temporal domain and uncertainty regions.  
 218 Formally, we define  $G_S$  as follows. For each signal  $x_i$ , let  $F_i$  be the set of end  
 219 points of its uncertainty regions. Let  $F = \{0, d\} \cup \bigcup_{i=1}^n F_i$  and let  $(s_j)_{1 \leq j \leq |F|}$

be a nondecreasing sequence of clock values corresponding to the elements of  $F$ . Then, the canonical segmentation of  $(S, \rightsquigarrow)$  is  $G_S = \{I_1, \dots, I_{|F|-1}\}$  where  $I_j = [s_j, s_{j+1})$  for all  $1 \leq j < |F|$ .

*Example 8.* Let  $(S, \rightsquigarrow)$  be a distributed boolean signal with  $S = (x_1, x_2)$  and  $\varepsilon = 2$  over the temporal domain  $[0, 8)$  as given in Figure 1. Both signals are initially 0. The signal  $x_1$  has a rising edge at time 2 and a falling edge at time 5, while  $x_2$  has a rising edge at time 3 and a falling edge at time 6. The uncertainty regions of  $x_1$  are  $(0, 4)$  and  $(3, 7)$ , while those of  $x_2$  are  $(1, 5)$  and  $(4, 8)$ . Then, we have  $F = \{0, 8\} \cup \{0, 1, 3, 4, 5, 7, 8\}$ , and thus the canonical segmentation is  $G_S = \{[0, 1), [1, 3), [3, 4), [4, 5), [5, 7), [7, 8)\}$ .



**Fig. 1.** The signals  $x_1$  (top, red, dashed) and  $x_2$  (bottom, blue, solid) from Example 8. The edges are marked with solid balls and their uncertainty regions are given as semi-transparent boxes around the edges. The resulting canonical segmentation  $G_S$  is shown below the graphical representation of the signals.

**Value Expressions** Consider a boolean signal  $x$  with a rising edge with an uncertainty region of  $(t_1, t_2)$ . As discussed above, the monitor only knows that the value of  $x$  changes from 0 to 1 in this interval. We represent this knowledge as a finite word  $v = 0 \cdot 1$ . This representation is called a *value expression* and it encodes the uncertain behavior of an observed signal relative to the monitor. Formally, a value expression is an element of  $\Sigma^*$  where  $\Sigma$  is the finite alphabet of values the signal takes. Given a signal  $x$  and an edge  $(t, x(t))$ , the value expression corresponding to the uncertainty region  $(\theta_{lo}(x, t), \theta_{hi}(x, t))$  is given

by  $v_{x,t} = v_- \cdot v_+$  where  $v_- = \lim_{s \rightarrow t^-} x(s)$  and  $v_+ = \lim_{s \rightarrow t^+} x(s)$ . Let us remark that this definition is general because finite-length piecewise-constant real-valued signals will only have a finite number of values, making  $\Sigma$  finite.

Notice that (i) uncertainty regions may overlap, and (ii) the canonical segmentation may split an uncertainty region into multiple segments. Consider a signal  $x$  with a rising edge in  $(1, 5)$  and a falling edge in  $(4, 8)$ . The corresponding value expressions are respectively  $v_1 = 0 \cdot 1$  and  $v_2 = 1 \cdot 0$ . Notice that the behavior of  $x$  in the interval  $[1, 4)$  can be expressed as  $\text{prefix}(v_1)$ , encoding whether the rising edge has happened yet or not. Similarly, the behavior in  $[4, 5)$  is given by  $\text{suffix}(v_1) \cdot \text{prefix}(v_2)$ , which captures whether the edges occur in this interval (thanks to prefixing and suffixing) and the fact that the rising edge happens before the falling edge (thanks to concatenation).

Formally, given a distributed signal  $(S, \rightsquigarrow)$ , we define a function  $\gamma : S \times G_S \rightarrow 2^{\Sigma^*}$  that maps each signal and segment of the canonical segmentation to a set of value expressions, capturing the signal's potential behaviors in the given segment. Let  $x$  be a signal in  $S$ , and let  $R_1, \dots, R_m$  be its uncertainty regions where  $R_i = (t_i, t'_i)$  and the corresponding value expression is  $v_i$  for all  $1 \leq i \leq m$ . Now, let  $I \in G_S$  be a segment with  $I = [s, s')$  and for each  $1 \leq i \leq m$  define the set  $V_i$  of value expressions capturing how  $I$  relates with  $R_i$  as follows:

$$V_i = \begin{cases} \{v_i\} & \text{if } t_i = s \wedge s' = t'_i \\ \text{prefix}(v_i) & \text{if } t_i = s \wedge s' < t'_i \\ \text{suffix}(v_i) & \text{if } t_i > s \wedge s' = t'_i \\ \text{infix}(v_i) & \text{if } t_i > s \wedge s' < t'_i \\ \{\epsilon\} & \text{otherwise} \end{cases} \quad (1)$$

The last case happens only when  $I \cap R_i$  is empty. We finally define  $\gamma$  as follows:


$$\gamma(x, I) = \text{destutter}(V_1 \cdot V_2 \cdot \dots \cdot V_m) \setminus \{\epsilon\}$$

Observe that  $\gamma(x, I)$  contains all the potential behaviors of  $x$  in segment  $I$  by construction. However, it is potentially overapproximate. This is mainly because the sets  $V_1, \dots, V_m$  contain redundancy by definition and the concatenation does not guarantee that an edge is considered exactly once.

*Example 9.* Recall the distributed signal  $(S, \rightsquigarrow)$  in Example 8 and Figure 1. In Figure 2a, we show the value expressions corresponding to its uncertainty regions. For example, the falling edge of  $x_1$  has an uncertainty region of  $(3, 7)$ , represented by the value expression  $1 \cdot 0$ . In Figure 2b, we give the function  $\gamma$  for  $(S, \rightsquigarrow)$ . For example,  $\gamma(x_1, [3, 4))$  is obtained from  $\text{suffix}(0 \cdot 1) \cdot \text{prefix}(1 \cdot 0)$  and  $\gamma(x_2, [0, 1)) = \{0\}$ .

**Overapproximation of Tr** Consider a distributed signal  $(S, \rightsquigarrow)$  of  $n$  signals, and let  $G_S$  be its canonical segmentation. We describe how the function  $\gamma$  defines a set  $\text{Tr}^+(S, \rightsquigarrow)$  of synchronous traces that overapproximates the set  $\text{Tr}(S, \rightsquigarrow)$ .





**Fig. 2.** (a) The uncertainty regions of the distributed signal in Example 8 and the corresponding value expressions. (b) The tabular representation of the function  $\gamma$  for the given distributed signal.

Let  $x \in S$  and  $x'$  be two signals with the same temporal domain, and let  $I = [s, s']$  be a segment in  $G_S$ . Let  $(t_1, x'(t_1)), \dots, (t_\ell, x'(t_\ell))$  be the edges of  $x'$  in segment  $I$  with  $t_i < t_{i+1}$  for all  $1 \leq i < \ell$ . The signal  $x'$  is  $I$ -consistent with  $x$  iff the value expression  $x'(s) \cdot x'(t_1) \cdot \dots \cdot x'(t_\ell)$  belongs to  $\gamma(x, I)$ . Moreover,  $x'$  is consistent with  $x$  iff it is  $I$ -consistent with  $x$  for all  $I \in G_S$ .

Now, let  $S = (x_1, \dots, x_n)$  and define  $\text{Tr}^+(S, \rightsquigarrow)$  as follows:

$$\text{Tr}^+(S, \rightsquigarrow) = \{(x'_1, \dots, x'_n) \mid x'_i \text{ is consistent with } x_i \text{ for all } 1 \leq i \leq n\}$$

*Example 10.* Recall the distributed signal  $(S, \rightsquigarrow)$  in Example 8 whose  $\gamma$  function is given in Figure 2b. Consider the synchronous trace  $w \in \text{Tr}(S, \rightsquigarrow)$  where the rising edges of both signals occur at time 3 and the falling edges at time 5. One can verify that  $w \in \text{Tr}^+(S, \rightsquigarrow)$  since for each  $i \in \{1, 2\}$  the value expression 1 is contained in  $\gamma(x_i, [3, 4])$  and  $\gamma(x_i, [4, 5])$  while 0 is contained in the remaining sets  $\gamma$  maps  $x_i$  to.

Now, consider a synchronous trace  $(x'_1, x'_2)$  where both signals are initially 0, have rising edges at time 2 and 3.5, and falling edges at time 3 and 5. Evidently, this trace does not belong to  $\text{Tr}(S, \rightsquigarrow)$  since  $x'_1$  and  $x'_2$  have more edges than  $x_1$  and  $x_2$ . Nonetheless, it belongs to  $\text{Tr}^+(S, \rightsquigarrow)$  since  $x'_1$  and  $x'_2$  are respectively consistent with  $x_1$  and  $x_2$ . To witness, notice that for each  $i \in \{1, 2\}$  the value expression  $0 \cdot 1$  is contained in  $\gamma(x_i, [1, 3])$  and  $\gamma(x_i, [3, 4])$ , the expression 1 is contained in  $\gamma(x_i, [4, 5])$ , and 0 is contained in the remaining sets  $\gamma$  maps  $x_i$  to.

Finally, we prove  $\text{Tr}^+$  overapproximates  $\text{Tr}$ .

**Lemma 11.** For every distributed signal  $(S, \rightsquigarrow)$ , we have  $\text{Tr}(S, \rightsquigarrow) \subseteq \text{Tr}^+(S, \rightsquigarrow)$ .

## 5 Monitoring Algorithm

In this section, given a distributed signal  $(S, \rightsquigarrow)$ , we describe an algorithm to compute  $[(S, \rightsquigarrow) \models \varphi]_+$ . The algorithm makes use of the function  $\gamma$  defined in Section 4 without explicitly computing  $\text{Tr}^+(S, \rightsquigarrow)$ . To achieve this, we first describe the notion of *asynchronous product* of value expressions to capture potential interleavings within segments. We continue with the evaluation of *untimed operators* and then *timed operators*. Finally, we conclude with putting all these together to compute the *semantics* of  $\text{STL}^+$  and discuss an efficient implementation of the monitoring algorithm using *bit vectors* to represent and manipulate sets of boolean value expressions.

*Remark 12.* For the sake of convenience, we focus on boolean signals for the rest of the section. Note that asynchronous products and the algorithm to compute  $[(S, \rightsquigarrow) \models \varphi]_+$  can be extended to value expressions over arbitrary finite alphabets, e.g., encoding real-valued signals. This allows us to express more complex properties where atomic propositions can be functions of real-valued signals.

**Asynchronous Products** Consider the value expressions  $u_1 = 0 \cdot 1$  and  $u_2 = 1 \cdot 0$  encoding the behaviors of two signals within a segment. Due to partial asynchrony, the behaviors within segments can be seen as completely asynchronous. To capture the potential interleavings of these behaviors, we consider how the values in  $u_1$  and  $u_2$  can align. In particular, there are three potential alignments: (i) the rising edge of  $u_1$  happens before the falling edge of  $u_2$ , (ii) the falling edge of  $u_2$  happens before the rising edge of  $u_1$ , and (iii) the two edges happen simultaneously. We respectively represent these with the tuples  $(011, 110)$ ,  $(001, 100)$ , and  $(01, 10)$  where the first component encodes  $u_1$  and the second  $u_2$ . Formally, given two value expressions  $u_1$  and  $u_2$ , we define their *asynchronous product* as follows:

$$u_1 \otimes u_2 = \{\text{destutter}(v_1, v_2) \mid v_i \in \text{stutter}_k(u_i), k = |u_1| + |u_2| - 1, i \in \{1, 2\}\}$$

Moreover, given two sets  $L_1$  and  $L_2$  of value expressions, we define the following:

$$L_1 \otimes L_2 = \{u_1 \otimes u_2 \mid u_1 \in L_1, u_2 \in L_2\}$$

Asynchronous products of value expressions allow us to lift value expressions to satisfaction signals of formulas.

*Example 13.* Recall the distributed signal  $(S, \rightsquigarrow)$  in Example 8 and its  $\gamma$  function given in Figure 2b. Suppose we want to compute the value expressions encoding the satisfaction of  $x_1 \wedge x_2$  in the segment  $[1, 3)$ . We can achieve this by first computing the asynchronous product  $\gamma(x_1, [3, 4)) \otimes \gamma(x_2, [3, 4))$ , and then computing the bitwise conjunction of each pair in the set. For example, considering the expression  $0 \cdot 1 \cdot 0$  for  $x_1$  and  $0 \cdot 1$  for  $x_2$ , the product contains the pair  $(010, 011)$ . Taking the bitwise conjunction of this pair gives us the expression  $0 \cdot 1 \cdot 0$  as a potential behavior for the satisfaction of  $x_1 \wedge x_2$  in this segment.

**Untimed Operations** As hinted in Example 13, to compute the semantics, we apply bitwise operations on value expressions and their asynchronous products to transform them into encodings of satisfaction signals of formulas. Consider the distributed signal  $(S, \rightsquigarrow)$  in Example 8 and suppose we want to compute  $[(S, \rightsquigarrow) \models \Diamond(x_1 \wedge x_2)]_+$ . To achieve this, we first compute for each segment in  $G_S$  the set of value expressions for the satisfaction of  $x_1 \wedge x_2$ , and then from these compute that of  $\Diamond(x_1 \wedge x_2)$ . This compositional approach allows us to evaluate arbitrary STL<sup>+</sup> formulas.

First, we define bitwise operations on boolean value expressions encoding atomic propositions. Then, we use these to evaluate (untimed) STL formulas over sets of value expressions.

Let  $u$  and  $v$  be boolean value expressions of length  $\ell$ . We denote by  $u \& v$  the bitwise-and operation, by  $u \mid v$  the bitwise-or, and by  $\sim u$  the bitwise-negation. In addition, we define the *bitwise strong until* operator as follows:

$$u \mathbf{U}^0 v = \left( \max_{i \leq j \leq \ell} \left( \min \left( v[j], \min_{i \leq k \leq j} u[k] \right) \right) \right)_{1 \leq i \leq \ell}$$

As usual, we derive *bitwise eventually* as  $\mathbf{E}u = 1^\ell \mathbf{U}^0 u$ , *bitwise always* as  $\mathbf{A}u = \sim(\mathbf{E}\sim u)$ , and *bitwise weak until* as  $u \mathbf{U}^1 v = (u \mathbf{U}^0 v) \mid (\mathbf{A}u)$ . The distinction between  $\mathbf{U}^0$  and  $\mathbf{U}^1$  will be useful later when we evaluate a formula segment by segment. We remark that the definitions of these operators coincide with the robustness semantics of (discrete time) STL. Finally, note that the output of these operations is a value expression of length  $\ell$ . For example, if  $u = 010$ , we have  $\mathbf{E}u = 110$  and  $\mathbf{A}u = 000$ .

Let  $(S, \rightsquigarrow)$  be a distributed signal. Consider an atomic proposition  $p \in \mathbf{AP}$  encoded as  $x_p \in S$  and let  $\varphi_1, \varphi_2$  be two STL formulas. We define the evaluation of untimed formulas with respect to  $(S, \rightsquigarrow)$  and a segment  $I \in G_S$  inductively:

$$\begin{aligned} \llbracket (S, \rightsquigarrow), I \models p \rrbracket &= \gamma(x_p, I) \\ \llbracket (S, \rightsquigarrow), I \models \neg \varphi_1 \rrbracket &= \{\sim u \mid u \in \llbracket (S, \rightsquigarrow), I \models \varphi_1 \rrbracket\} \\ \llbracket (S, \rightsquigarrow), I \models \varphi_1 \wedge \varphi_2 \rrbracket &= \{u_1 \& u_2 \mid (u_1, u_2) \in \llbracket (S, \rightsquigarrow), I \models \varphi_1 \rrbracket \otimes \llbracket (S, \rightsquigarrow), I \models \varphi_2 \rrbracket\} \\ \llbracket (S, \rightsquigarrow), I \models \varphi_1 \mathcal{U} \varphi_2 \rrbracket &= \{u_1 \mathbf{U}^a u_2 \mid (u_1, u_2) \in \llbracket (S, \rightsquigarrow), I \models \varphi_1 \rrbracket \otimes \llbracket (S, \rightsquigarrow), I \models \varphi_2 \rrbracket, \\ &\quad a \in \mathbf{first}(\llbracket (S, \rightsquigarrow), I' \models \varphi_1 \mathcal{U} \varphi_2 \rrbracket)\} \end{aligned}$$

where  $I'$  is the segment that follows  $I$  in  $G_S$ , if it exists. For completeness, for every formula  $\varphi$  we define  $\llbracket (S, \rightsquigarrow), I' \models \varphi \rrbracket = \{0\}$  when  $I' \notin G_S$ . When  $I$  is the first segment in  $G_S$ , we simply write  $\llbracket (S, \rightsquigarrow) \models \varphi \rrbracket$ . Similarly as above, we can use the standard derived operators to compute the corresponding sets of value expressions. Intuitively, for a given formula and a segment, the evaluation above produces a set of value expressions encoding the formula's satisfaction within the segment.

*Example 14.* Recall the distributed signal  $(S, \rightsquigarrow)$  in Example 8 and its  $\gamma$  function given in Figure 2b. Suppose we want to compute  $\llbracket (S, \rightsquigarrow), [5, 7) \models \Diamond(x_1 \wedge x_2) \rrbracket$ .

First, we compute  $\llbracket (S, \rightsquigarrow), [5, 7) \models x_1 \wedge x_2 \rrbracket$  by computing the bitwise conjunction over the asynchronous product  $\gamma(x_1, [5, 7)) \otimes \gamma(x_2, [5, 7))$  and destuttering. For example, since  $010 \in \gamma(x_1, [5, 7))$  and  $01 \in \gamma(x_2, [5, 7))$ , the pair  $(0010, 0111)$  is in the product, whose conjunction gives us  $010$  after destuttering. Repeating this for the rest, we obtain  $\llbracket (S, \rightsquigarrow), [5, 7) \models x_1 \wedge x_2 \rrbracket = \{0, 01, 010, 1, 10\}$ . Finally, we compute  $\llbracket (S, \rightsquigarrow), [5, 7) \models \Diamond(x_1 \wedge x_2) \rrbracket$  by applying each expression in  $\llbracket (S, \rightsquigarrow), [5, 7) \models x_1 \wedge x_2 \rrbracket$  the bitwise eventually operator and destuttering. The resulting set  $\{0, 1, 10\}$  encodes the satisfaction signal of  $\Diamond(x_1 \wedge x_2)$  in  $[5, 7)$ . Note that we do not need to consider the evaluation of the next segment for the eventually operator since  $\llbracket (S, \rightsquigarrow), [7, 8) \models x_1 \wedge x_2 \rrbracket = \{0\}$ .

**Timed Operations** Handling timed operations requires a closer inspection as value expressions are untimed by definition. We address this issue by considering how a given evaluation interval relates with a given segmentation. For example, take a segmentation  $G_S = \{[0, 4), [4, 6), [6, 10)\}$  and an evaluation interval  $J = [0, 5)$ . Suppose we are interested in how a signal  $x \in S$  behaves with respect to  $J$  over the first segment  $I = [0, 4)$ . First, to see how  $J$  relates with  $G_S$  with respect to  $I = [0, 4)$ , we “slide” the interval  $J$  over  $I \oplus J = [0, 9)$  and consider the different ways it intersects the segments in  $G_S$ . Initially,  $J$  covers the entire segment  $[0, 4)$  and the beginning of  $[4, 6)$ , for which the potential behaviors of  $x$  are captured by the set  $\gamma(x, [0, 4)) \cdot \text{prefix}(\gamma(x, [4, 6)))$ . Now, if we slide the window and take  $J' = [3, 7)$ , the window covers the ending of  $[0, 4)$ , the entire  $[4, 6)$ , and the beginning of  $[6, 10)$ , for which the potential behaviors are captured by the set  $\text{suffix}(\gamma(x, [0, 4))) \cdot \gamma(x, [4, 6)) \cdot \text{prefix}(\gamma(x, [6, 9)))$ . We call these sets the *profiles* of  $J$  and  $J'$  with respect to  $(S, \rightsquigarrow)$ ,  $x$ , and  $I$ .

Let  $(S, \rightsquigarrow)$  be a distributed signal,  $I \in G_S$  be a segment, and  $\varphi$  be an STL formula. Let us introduce the notation we use in the definition below. First, we abbreviate the set  $\llbracket (S, \rightsquigarrow), I \models \varphi \rrbracket$  of value expressions as  $\tau_{\varphi, I}$ . Second, given an interval  $K$ , we respectively denote by  $l_K$  and  $r_K$  its left and right end points. Third, recall that we denote by  $F$  the set of end points of  $G_S$  (see Section 4). Given an interval  $J$ , we define the *profile* of  $J$  with respect to  $(S, \rightsquigarrow)$ ,  $\varphi$ , and  $I$  as follows.

$$\text{profile}((S, \rightsquigarrow), \varphi, I, J) = \begin{cases} \text{prefix}(\tau_{\varphi, I}) & \text{if } l_I = l_J \wedge r_I > r_J \\ \text{infix}(\tau_{\varphi, I}) & \text{if } l_I < l_J \wedge r_I > r_J \\ \tau_{\varphi, I} \cdot \kappa(\varphi, I, J) & \text{if } l_I = l_J \wedge r_I \leq r_J \wedge r_J \in F \setminus J \\ \tau_{\varphi, I} \cdot \kappa(\varphi, I, J) \cdot \text{first}(\tau_{\varphi, I'}) & \text{if } l_I = l_J \wedge r_I \leq r_J \wedge r_J \in F \cap J \\ \tau_{\varphi, I} \cdot \kappa(\varphi, I, J) \cdot \text{prefix}(\tau_{\varphi, I'}) & \text{if } l_I = l_J \wedge r_I \leq r_J \wedge r_J \notin F \\ \text{suffix}(\tau_{\varphi, I}) \cdot \kappa(\varphi, I, J) & \text{if } l_I < l_J < r_I \leq r_J \wedge r_J \in F \setminus J \\ \text{suffix}(\tau_{\varphi, I}) \cdot \kappa(\varphi, I, J) \cdot \text{first}(\tau_{\varphi, I'}) & \text{if } l_I < l_J < r_I \leq r_J \wedge r_J \in F \cap J \\ \text{suffix}(\tau_{\varphi, I}) \cdot \kappa(\varphi, I, J) \cdot \text{prefix}(\tau_{\varphi, I'}) & \text{if } l_I < l_J < r_I \leq r_J \wedge r_J \notin F \\ \{\epsilon\} & \text{otherwise} \end{cases}$$

where we assume  $J$  is trimmed to fit the temporal domain of  $S$  and  $I' \in G_S$  is such that  $r_J \in I'$ . Moreover,  $\kappa(\varphi, I, J)$  is the concatenation  $\tau_{\varphi, I_1} \cdot \dots \cdot \tau_{\varphi, I_m}$  such that  $I, I_1, \dots, I_m, I'$  are consecutive segments in  $G_S$ . If  $I_1, \dots, I_m$  do not exist, we let  $\kappa(\varphi, I, J) = \{\epsilon\}$ . Note that the last case happens when  $I \cap J$  is empty.

397 We now formalize the intuitive approach of “sliding”  $J$  over the segmentation to  
 398 obtain the various profiles it produces as follows.


$$\mathbf{pfs}((S, \rightsquigarrow), \varphi, I, J) = \{\mathbf{destutter}(\mathbf{profile}((S, \rightsquigarrow), \varphi, I, J')) \mid J' \subseteq I \oplus J, J' \sim J\}$$

399 where  $J' \sim J$  holds when  $|J'| = |J|$  and  $J'$  contains an end point (left or  
 400 right) iff  $J$  does so. Note that although infinitely many intervals  $J'$  satisfy the  
 401 conditions given above (due to denseness of time), the set defined by  $\mathbf{pfs}$  is finite.  
 402 We demonstrate this and the computation of  $\mathbf{pfs}$  in Example 15 and Figure 3.

403 *Example 15.* Recall the distributed signal  $(S, \rightsquigarrow)$  in Example 8 and its  $\gamma$  function  
 404 given in Figure 2b. We demonstrate the computation of  $\mathbf{pfs}((S, \rightsquigarrow), x_1, [1, 3], [0, 1])$ .  
 405 Intuitively, sliding the interval  $[0, 1]$  over the window  $[1, 3] \oplus [0, 1]$  (as shown in  
 406 Figure 3) gives us the following sets:

$$\begin{aligned} P_1 &= \mathbf{destutter}(\mathbf{prefix}(\gamma(x_1, [1, 3]))) = \{0, 01, 1\} \\ P_2 &= \mathbf{destutter}(\mathbf{infix}(\gamma(x_1, [1, 3]))) = \{0, 01, 1\} \\ P_3 &= \mathbf{destutter}(\mathbf{suffix}(\gamma(x_1, [1, 3]))) = \{0, 01, 1\} \\ P_4 &= \mathbf{destutter}(\mathbf{suffix}(\gamma(x_1, [1, 3])) \cdot \mathbf{prefix}(\gamma(x_1, [3, 4]))) \\ &= \{0, 01, 010, 0101, 01010, 1, 10, 101, 1010\} \end{aligned}$$

407 Therefore, we obtain  $\mathbf{pfs}((S, \rightsquigarrow), x_1, [1, 3], [0, 1]) = \{P_1, P_2, P_3, P_4\}$ . This set over-  
 408 approximates the potential behaviors of  $x_1$ , for all  $t \in [1, 3]$ , in the interval  
 409  $t \oplus [0, 1]$ .



**Fig. 3.** The profiles of  $J = [0, 1]$  with respect to  $x_1 \in S$  of Example 8. The four representative intervals of each profile is shown with solid black lines below the tabular representation of  $\gamma$  for  $x_1$ .

410 Let  $\varphi_1$  and  $\varphi_2$  be two STL formulas. Intuitively, once we have the profiles of  
 411 a given interval  $J$  with respect to  $\varphi_1$  and  $\varphi_2$ , we can evaluate the correspond-  
 412 ing untimed formulas on the product of these profiles and concatenate them.  
 413 Formally, we handle the evaluation of timed formulas inductively as follows.

$$\llbracket (S, \rightsquigarrow), I \models \varphi_1 \mathcal{U}_J \varphi_2 \rrbracket = \mathbf{destutter}(\{u_1 \mathbf{U}^0 u_2 \mid (u_1, u_2) \in P_1 \otimes Q_1\} \cdot \dots \cdot \{u_1 \mathbf{U}^0 u_2 \mid (u_1, u_2) \in P_k \otimes Q_k\})$$

where  $\mathbf{pfs}((S, \rightsquigarrow), \varphi_1, I, J) = \{P_1, \dots, P_k\}$  and  $\mathbf{pfs}((S, \rightsquigarrow), \varphi_2, I, J) = \{Q_1, \dots, Q_k\}$  such that the intervals producing  $P_i$  and  $Q_i$  respectively start before those producing  $P_{i+1}$  and  $Q_{i+1}$  for all  $1 \leq i < k$ .

*Example 16.* Let  $(S, \rightsquigarrow)$  be as in Example 8 and its  $\gamma$  function as given in Figure 2b. We demonstrate the evaluation of the timed formula  $\Diamond_{[0,1)} x_1$  over the segment  $[1, 3)$ . Recall from Example 15 the set  $\mathbf{pfs}((S, \rightsquigarrow), x_1, [1, 3), [0, 1)) = \{P_1, P_2, P_3, P_4\}$  of profiles. First, we apply the bitwise eventually operator to each value expression in each of these profiles separately:  $\{Eu \mid u \in P_1\} = \{0, 1\}$ ,  $\{Eu \mid u \in P_2\} = \{0, 1\}$ ,  $\{Eu \mid u \in P_3\} = \{0, 1\}$ , and  $\{Eu \mid u \in P_4\} = \{0, 10, 1\}$ . Then, we concatenate these sets and destutter to obtain the following:

$$\llbracket (S, \rightsquigarrow), [1, 3) \models \Diamond_{[0,1)} x_1 \rrbracket = \{0, 01, 010, 0101, 01010, 1, 10, 101, 1010\}$$

**Computing the Semantics of STL<sup>+</sup>** Putting it all together, given a distributed signal  $(S, \rightsquigarrow)$  and an STL<sup>+</sup> formula  $\varphi$ , we can compute  $\llbracket (S, \rightsquigarrow) \models \varphi \rrbracket_+$  thanks to the following theorem.

**Theorem 17.** *For every distributed signal  $(S, \rightsquigarrow)$ , we have  $\llbracket (S, \rightsquigarrow) \models \varphi \rrbracket_+ = \top$  (resp.  $\perp$ ,  $?$ ) iff  $\mathbf{first}(\llbracket (S, \rightsquigarrow) \models \varphi \rrbracket) = \{1\}$  (resp.  $\{0\}$ ,  $\{0, 1\}$ ).*

**Sets of Boolean Value Expressions as Bit Vectors** Evidently, asynchronous products are expensive to compute. Our implementation of the algorithm we describe in this section relies on the following observation: Sets of boolean value expressions and their operations can be efficiently implemented through bit vectors. Intuitively, to represent such a set, we can encode each element using its first bit and its length since value expressions are boolean and always destuttered. Moreover, to evaluate untimed operations on such sets, we only need to know the maximal lengths of the four possible types of expressions ( $0 \dots 0$ ,  $0 \dots 1$ ,  $1 \dots 0$ , and  $1 \dots 1$ ) and whether the set contains 0 or 1 (to handle some edge cases). This is because value expressions corresponding to same segments can be seen as completely asynchronous and the possible interleavings obtained from shorter expressions can be obtained from longer ones. This approach enables, for example, an algorithm for conjunction of sets of value expressions that runs in  $O(|u| + |v|)$  time where  $u$  and  $v$  are the longest expressions in the two sets. Note that the same idea also applies to untimed temporal operators.

## 6 Experimental Evaluation

TODO

## 7 Conclusion

TODO

## References

1. Maler, O., Nickovic, D.: Monitoring properties of analog and mixed-signal circuits. *Int. J. Softw. Tools Technol. Transf.* **15**(3), 247–268 (2013). <https://doi.org/10.1007/s10009-012-0247-9>
2. Momtaz, A., Abbas, H., Bonakdarpour, B.: Monitoring signal temporal logic in distributed cyber-physical systems. In: Mitra, S., Venkatasubramanian, N., Dubey, A., Feng, L., Ghasemi, M., Sprinkle, J. (eds.) *Proceedings of the ACM/IEEE 14th International Conference on Cyber-Physical Systems, ICCPS 2023, (with CPS-IoT Week 2023), San Antonio, TX, USA, May 9-12, 2023.* pp. 154–165. ACM (2023). <https://doi.org/10.1145/3576841.3585937>

## Appendix

### Proof of Theorem 6

*Proof.* Let  $\varphi$  be an STL formula and  $(S, \rightsquigarrow)$  be a distributed signal. Assume  $[(S, \rightsquigarrow) \models \varphi]_+ = \top$ . We want to show that  $[(S, \rightsquigarrow) \models \varphi] = \top$ . Expanding the definition of  $[(S, \rightsquigarrow) \models \varphi]_+ = \top$ , we have  $w \models \varphi$  for all  $w \in \text{Tr}^+(S, \rightsquigarrow)$ . By Lemma 11, we have  $\text{Tr}(S, \rightsquigarrow) \subseteq \text{Tr}^+(S, \rightsquigarrow)$ . Then, it holds that  $w \models \varphi$  for all  $w \in \text{Tr}(S, \rightsquigarrow)$ . Therefore,  $[(S, \rightsquigarrow) \models \varphi] = \top$  by definition. The case of  $[(S, \rightsquigarrow) \models \varphi]_+ = \perp$  follows from the same arguments.

### Proof of Lemma 11

*Proof.* Let  $(S, \rightsquigarrow)$  be a distributed signal where  $S = (x_1, \dots, x_n)$ . Let  $w = (y_1, \dots, y_n) \in \text{Tr}(S, \rightsquigarrow)$  be a trace. We want to show that  $w \in \text{Tr}^+(S, \rightsquigarrow)$ . First, let us recall the definition of  $\text{Tr}^+$ .

$$\text{Tr}^+(S, \rightsquigarrow) = \{(x'_1, \dots, x'_n) \mid x'_i \text{ is consistent with } x_i \text{ for all } 1 \leq i \leq n\}$$

Let  $1 \leq i \leq n$  be arbitrary. To show that  $y_i$  is consistent with  $x_i$ , we need to show that  $y_i$  is  $I$ -consistent with  $x_i$  for all  $I \in G_S$ . Let  $I = [t_0, s)$  be an arbitrary segment in  $G_S$ , let  $(t_1, y_i(t_1)), \dots, (t_\ell, y_i(t_\ell))$  be the edges of  $y_i$  in segment  $I$  with  $t_j < t_{j+1}$  for all  $1 \leq j < \ell$ . To show that  $y_i$  is  $I$ -consistent with  $x_i$ , we need to show that the expression  $y_i(t_0) \cdot y_i(t_1) \cdot \dots \cdot y_i(t_\ell)$  belongs to  $\gamma(x_i, I)$ . We sketch the proof idea below.

Note that  $w$  can be seen as a trace obtained through an  $\varepsilon$ -retiming of  $S$  (see [2, Section 4.2]). Then, the timestamp  $t$  of any edge of  $x_i$  is mapped to some clock value in the range  $(\theta_{\text{lo}}(t), \theta_{\text{hi}}(t))$ . In particular,  $|t - c_i^{-1}(t)| < \varepsilon$  for all  $t \in \{t_0, t_1, \dots, t_\ell\}$ , where  $c_i^{-1}(t)$  is the local clock value of  $x_i$  that is mapped to  $t$ .

Since  $y_i$  has  $\ell$  edges in  $I$ , it holds that  $x_i$  has at least  $\ell$  edges in  $(t_0 - \varepsilon, s + \varepsilon)$ . Since  $I$  is a segment in  $G_S$ , there are  $\ell$  of these that are consecutive such that the intersection of their uncertainty regions contain  $(t_0, s)$ , i.e.,  $(t_0, s) \subseteq \bigcap_{1 \leq j \leq \ell} (\theta_{\text{lo}}(t'_j), \theta_{\text{hi}}(t'_j))$  where  $t'_j = c_i^{-1}(t_j)$  is the corresponding timestamp in  $x_i$  for all  $0 \leq j \leq \ell$ . In particular, note that  $y_i(t_j) = x_i(t'_j)$  for all  $0 \leq j \leq \ell$ .

Now, notice that, by definition,  $\gamma(x_i, I)$  takes into account every edge of  $x_i$  whose uncertainty region has a nonempty intersection with  $I$ , and preserves their order. Let  $V_j$  be the set of value expressions capturing how  $I$  relates with the uncertainty intervals of the edge  $(t'_j, x_i(t'_j))$  for all  $1 \leq j \leq \ell$  (as defined in Equation (1)). Then,  $\text{destutter}(\{x_i(t'_0)\} \cdot V_1 \cdot \dots \cdot V_\ell) \subseteq \gamma(x_i, I)$ . One can verify that for all  $1 \leq j \leq \ell$ , either  $x_i(t'_j)$  or  $x_i(t'_{j-1}) \cdot x_i(t'_j)$  belongs to  $V_j$ . This allows us to choose a value expression  $v_j$  from each  $V_j$  such that  $\text{destutter}(\{x_i(t'_0)\} \cdot v_1 \cdot \dots \cdot v_\ell) = x_i(t'_0) \cdot x_i(t'_1) \cdot \dots \cdot x_i(t'_\ell)$ , which concludes the proof.

Note that if there are more edges of  $x_i$  with a timestamp smaller than  $t'_0$  or larger than  $t'_\ell$  whose uncertainty intervals intersect with  $I$ , then the corresponding set of value expressions is obtained either by prefixing or suffixing. In either case, we can choose  $\epsilon$  from these sets for concatenation with the remaining edges' value expressions and obtain the desired result.

#### Proof of Theorem 17

*Proof.* **TODO**